

CPSC 340: Machine Learning and Data Mining

Least Squares



Fall 2019

Admin

- **Assignment 2:**
 - 2 late days to hand in tonight, 1 for Wednesday.
- **Assignment 3** is up:
 - Start early, this is usually the longest assignment.
 - Has details about the project.
 - Only 1 late day allowed.
- We're going to start using **calculus** and **linear algebra** a lot.
 - You should **start reviewing these ASAP** if you are rusty.
 - A review of relevant calculus concepts is [here](#).
 - A review of relevant linear algebra concepts is [here](#).

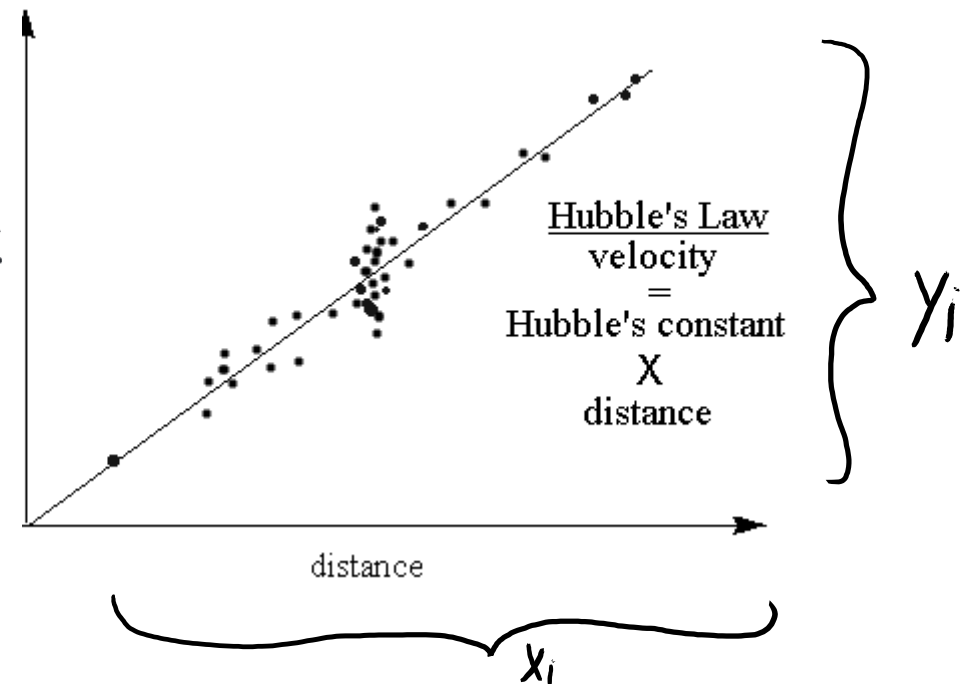

Example: Dependent vs. Explanatory Variables

- We want to discover relationship between numerical variables:
 - Does number of lung cancer deaths change with number of cigarettes?
 - Does number of skin cancer deaths change with latitude?





Example: Dependent vs. Explanatory Variables

- We want to discover relationship between numerical variables:
 - Do people in big cities walk faster?
 - Is the universe expanding or shrinking or staying the same size?




Example: Dependent vs. Explanatory Variables

- We want to discover relationship between numerical variables:
 - Does number of gun deaths change with gun ownership?
 - Does number violent crimes change with violent video games?



Example: Dependent vs. Explanatory Variables

- We want to discover relationship between numerical variables:
 - Does higher gender equality index lead to more women STEM grads?
- Not that we're doing supervised learning:
 - Trying to predict value of 1 variable (the ' y_i ' values). (instead of measuring correlation between 2).
- Supervised learning **does not give causality**:
 - OK: "Higher index **is correlated** with lower grad %".
 - OK: "Higher index **helps predict** lower grad %".
 - BAD: "Higher index **leads to** lower grads %".
 - People/media get these confused all the time, be careful!
 - There **are lots of potential reasons for this correlation**.



Handling Numerical Labels

- One way to handle numerical y_i : **discretize**.
 - E.g., for 'age' could we use {'age ≤ 20 ', ' $20 < \text{age} \leq 30$ ', 'age > 30 '}.
 - Now we can apply methods for classification to do regression.
 - But **coarse discretization loses resolution**.
 - And **fine discretization requires lots of data**.
- There exist regression versions of classification methods:
 - Regression trees, probabilistic models, non-parametric models.
- Today: one of oldest, but still most popular/important methods:
 - **Linear regression based on squared error**.
 - Interpretable and the building block for more-complex methods.


Linear Regression in 1 Dimension

- Assume we only have 1 feature ($d = 1$):
 - E.g., x_i is number of cigarettes and y_i is number of lung cancer deaths.
- Linear regression makes predictions \hat{y}_i using a linear function of x_i :

$$\hat{y}_i = w x_i$$

- The parameter 'w' is the weight or regression coefficient of x_i .
 - We're temporarily ignoring the y-intercept.
- As x_i changes, slope 'w' affects the rate that \hat{y}_i increases/decreases:
 - Positive 'w': \hat{y}_i increase as x_i increases.
 - Negative 'w': \hat{y}_i decreases as x_i increases.

Linear Regression in 1 Dimension



Aside: terminology woes

- Different fields use different terminology and symbols.
 - Data points = **objects** = **examples** = rows = observations.
 - **Inputs** = predictors = **features** = explanatory variables = regressors = independent variables = covariates = columns.
 - **Outputs** = outcomes = targets = response variables = dependent variables (also called a “label” if it’s categorical).
 - Regression coefficients = **weights** = parameters = betas.
- With linear regression, the symbols are inconsistent too:
 - In ML, the data is X and the weights are w .
 - In statistics, the data is X and the weights are β .
 - In optimization, the data is A and the weights are x .

Least Squares Objective

- Our **linear model** is given by:

$$\hat{y}_i = w x_i$$

- So we make **predictions** for a new example by using:

$$\hat{y}_i = w \tilde{x}_i$$

- But we **can't use the same error** as before:

- Even if data comes from a linear model but has noise,
we can have $\hat{y}_i \neq y_i$ for all training examples 'i' for the "best" model
- Unlikely that **line goes exactly** through the training examples.

Least Squares Objective

- We need a way to evaluate **numerical error**.
- Classic way is setting slope 'w' to minimize **sum of squared errors**:

$$f(w) = \sum_{i=1}^n (w x_i - y_i)^2$$

Annotations for the equation:

- A blue arrow points from y_i to the text "True value of y_i ".
- A blue arrow points from $w x_i$ to the text "Our prediction \hat{y}_i ".

Sum up the squared differences over all training examples.


Difference between prediction and true value for example 'i'.

- There are some justifications for this choice.
 - A probabilistic interpretation is coming later in the course.
- But usually, it is done because **it is easy to minimize**.

Least Squares Objective

- Classic way to set slope 'w' is minimizing **sum of squared errors**:

$$f(w) = \sum_{i=1}^n (wx_i - y_i)^2$$




"Error" is the sum of the squared values of these vertical distances between the line ($w x_i$) and the targets (y_i)

↓
If this error is small, then our predictions are close to the targets.

Least Squares Objective

- Classic way to set slope 'w' is minimizing **sum of squared errors**:

$$f(w) = \sum_{i=1}^n (wx_i - y_i)^2$$




"Error" is the sum of the squared values of these vertical distances between the line ($w x_i$) and the targets (y_i)

↓
If this error is **large**, then our predictions are **far from** the targets.

Minimizing a Differential Function

- Math 101 approach to minimizing a differentiable function 'f':
 1. Take the derivative of 'f'.
 2. Find points 'w' where the derivative $f'(w)$ is equal to 0.
 3. Choose the smallest one (and check that $f''(w)$ is positive).



Digression: Multiplying by a Positive Constant

- Note that this problem:

$$f(w) = \sum_{i=1}^n (w x_i - y_i)^2$$

- Has the **same set of minimizers** as this problem:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w x_i - y_i)^2$$

- And these also have the same minimizers:

$$f(w) = \frac{1}{n} \sum_{i=1}^n (w x_i - y_i)^2 \qquad f(w) = \frac{1}{2n} \sum_{i=1}^n (w x_i - y_i)^2 + 1000$$

- I can **multiply 'f' by any positive constant and not change solution.**
 - Gradient will still be zero at the same locations.
 - We'll use this trick a lot!

Finding Least Squares Solution

- Finding 'w' that minimizes **sum of squared errors**:

$$\begin{aligned} f(w) &= \frac{1}{2} \sum_{i=1}^n (w x_i - y_i)^2 = \frac{1}{2} \sum_{i=1}^n [w^2 x_i^2 - 2w x_i y_i + y_i^2] && \text{(expand square)} \\ &= \frac{w^2}{2} \underbrace{\sum_{i=1}^n x_i^2}_{\text{constant 'a'}} - w \underbrace{\sum_{i=1}^n x_i y_i}_{\text{constant 'b'}} + \frac{1}{2} \underbrace{\sum_{i=1}^n y_i^2}_{\text{constant 'c'}} && \text{(split sums, take 'w' outside)} \\ &= \frac{w^2}{2} a - wb + c \end{aligned}$$

Take derivative: $f'(w) = wa - b + 0$

Setting $f'(w)=0$ and solving gives $w = \frac{b}{a} = \frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2}$ (exists if we have a non-zero feature)

Finding Least Squares Solution

- Finding 'w' that minimizes **sum of squared errors**:

Setting $f'(w)=0$ and solving gives $w = \frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2}$ (exists if we have one non-zero x_i)

- Let's check that this is a **minimizer** by checking second derivative:

$$f'(w) = w \sum_{i=1}^n x_i^2 - \sum_{i=1}^n x_i y_i$$


$$f''(w) = \sum_{i=1}^n x_i^2$$

- Since (anything)² is non-negative, $f''(w) \geq 0$ and this is a minimizer.

Least Squares Objective/Solution (Another View)

- Least squares **minimizes a quadratic that is a sum of quadratics:**

$$f(w) = \underbrace{(wx_1 - y_1)^2}_{\downarrow} + \underbrace{(wx_2 - y_2)^2}_{\downarrow} + (wx_3 - y_3)^2 + \dots + \underbrace{(wx_n - y_n)^2}_{\downarrow}$$



(pause)

Motivation: Combining Explanatory Variables

- Smoking is **not the only contributor** to lung cancer.
 - For example, there environmental factors like exposure to asbestos.
- How can we model the **combined effect** of smoking and asbestos?
- A simple way is with a **2-dimensional linear function**:

$$\hat{y}_i = w_1 x_{i1} + w_2 x_{i2}$$

Handwritten annotations for the equation above:

- Blue arrows point from w_1 and x_{i1} to the text: "weight" of feature 1
- Blue arrows point from x_{i1} and x_{i2} to the text: Value of feature 1 in example 'i'
- Green arrows point from w_2 and x_{i2} to the text: "weight" on feature 2.
- Green arrows point from x_{i2} and the entire right side of the equation to the text: Value of feature 2 in example 'i'

- We have a weight w_1 for feature '1' and w_2 for feature '2':


$$\hat{y}_i = 10(\# \text{ cigarettes}) + 25(\# \text{ asbestos})$$

Least Squares in 2-Dimensions

- Linear model:

$$\hat{y}_i = w_1 x_{i1} + w_2 x_{i2}$$

- This defines a
**two-dimensional
hyper-plane.**




Least Squares in 2-Dimensions

- Linear model:


$$\hat{y}_i = w_1 x_{i1} + w_2 x_{i2}$$

- This defines a
two-dimensional
hyper-plane.
- Not just a line!



Why don't we have a y-intercept?


- Linear model is $\hat{y}_i = wx_i$ instead of $\hat{y}_i = wx_i + w_0$ with y-intercept w_0 .
- Without an intercept, if $x_i = 0$ then we **must predict $\hat{y}_i = 0$** .



Why don't we have a y-intercept?

- Linear model is $\hat{y}_i = wx_i$ instead of $\hat{y}_i = wx_i + w_0$ with y-intercept w_0 .
- Without an intercept, if $x_i = 0$ then we **must predict $\hat{y}_i = 0$** .

Adding
y-intercept
fixes this.



Adding a Bias Variable

- Simple trick to add a y-intercept (“bias”) variable:
 - Make a new matrix “Z” with an extra feature that is always “1”.

$$X = \begin{bmatrix} -0.1 \\ 0.3 \\ 0.2 \end{bmatrix}$$

$$Z = \begin{bmatrix} 1 & -0.1 \\ 1 & 0.3 \\ 1 & 0.2 \end{bmatrix}$$

"always 1" X

- Now use “Z” as your features in linear regression.
 - We’ll use ‘v’ instead of ‘w’ as regression weights when we use features ‘Z’.

$$\hat{y}_i = v_1 z_{i1} + v_2 z_{i2} = w_0 + w_1 x_{i1}$$

\downarrow \downarrow \downarrow \downarrow
 w_0 1 w_1 x_{i1}

- So we can have a non-zero y-intercept by changing features.
 - This means we can ignore the y-intercept in our derivations, which is cleaner.

Different Notations for Least Squares

- If we have 'd' features, the **d-dimensional linear model** is:

$$\hat{y}_i = w_1 x_{i1} + w_2 x_{i2} + w_3 x_{i3} + \dots + w_d x_{id}$$

– In words, our model is that the **output is a weighted sum of the inputs**.

- We can re-write this in **summation notation**:

$$\hat{y}_i = \sum_{j=1}^d w_j x_{ij}$$

- We can also re-write this in **vector notation**:

$$\hat{y}_i = \underbrace{w^T x_i}_{\text{"inner product" between vectors}} \quad (\text{assuming 'w' and } x_i \text{ are column-vectors})$$

Notation Alert (again)

- In this course, all **vectors are assumed to be column-vectors**:

$$w = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_d \end{bmatrix}$$

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{id} \end{bmatrix}$$

- So **$w^T x_i$ is a scalar**:
$$w^T x_i = \begin{bmatrix} w_1 & w_2 & \dots & w_d \end{bmatrix} \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{id} \end{bmatrix} = w_1 x_{i1} + w_2 x_{i2} + \dots + w_d x_{id} = \sum_{j=1}^d w_j x_{ij}$$

- So **rows of 'X' are actually transpose of column-vector x_i** :

$$X = \begin{bmatrix} \text{---} x_1^T \text{---} \\ \text{---} x_2^T \text{---} \\ \vdots \\ \text{---} x_n^T \text{---} \end{bmatrix}$$

Least Squares in d-Dimensions

- The **linear least squares** model in d-dimensions minimizes:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2$$


'w' is now a vector

*prediction is inner product of 'w' and 'x_i'
(linear combination of features)*


"Error" is still the sum of squared differences between "true" y_i and our "prediction" $w^T x_i$

- Dates back to 1801: Gauss used it to predict location of Ceres.
- How do we find the **best vector 'w'** in 'd' dimensions?
 - Set the derivative of each variable ("**partial derivative**") to 0?

Partial Derivatives



Partial Derivatives




"Partial" derivative of 'f' with respect to w_2 is the derivative with respect to w when all other variables are held fixed.

Denoted by $\frac{\partial}{\partial w_2}$ for variable w_2


Gradient and Critical Points in d-Dimensions

- Generalizing “set the derivative to 0 and solve” in d-dimensions:
 - Find ‘w’ where the **gradient** vector **equals the zero vector**.
- **Gradient** is vector with partial derivative ‘j’ in position ‘j’:

$$\nabla f(w) = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \frac{\partial f}{\partial w_2} \\ \vdots \\ \frac{\partial f}{\partial w_d} \end{bmatrix}$$



Tangent slope is 0 in every direction at minimizer.



Least Squares Partial Derivatives (1 Example)

- The **linear least squares** model in d-dimensions for 1 example:

$$f(w_1, w_2, \dots, w_d) = \frac{1}{2} (\hat{y}_i - y_i)^2 = \frac{1}{2} \hat{y}_i^2 - \hat{y}_i y_i + \frac{1}{2} y_i^2$$
$$\hat{y}_i = w_1 x_{i1} + w_2 x_{i2} + \dots + w_d x_{id} = \frac{1}{2} \left(\sum_{j=1}^d w_j x_{ij} \right)^2 + \left(\sum_{j=1}^d w_j x_{ij} \right) y_i + \frac{1}{2} y_i^2$$

- Computing the **partial derivative** for variable '1':

$$\begin{aligned} \frac{\partial}{\partial w_1} f(w_1, w_2, \dots, w_d) &= \left(\sum_{j=1}^d w_j x_{ij} \right) x_{i1} - y_i x_{i1} + 0 \\ &= \left(\sum_{j=1}^d w_j x_{ij} - y_i \right) x_{i1} \\ &= (w^T x_i - y_i) x_{i1} \end{aligned}$$

Least Squares Partial Derivatives ('n' Examples)

- Linear least squares partial derivative for variable 1 on example 'i':

$$\frac{\partial}{\partial w_1} f(w_1, w_2, \dots, w_d) = (w^T x_i - y_i) x_{i1}$$

- For a generic variable 'j' we would have:

$$\frac{\partial}{\partial w_j} f(w_1, w_2, \dots, w_d) = (w^T x_i - y_i) x_{ij}$$

- And if 'f' is summed over all 'n' examples we would have:

$$\frac{\partial}{\partial w_j} f(w_1, w_2, \dots, w_d) = \sum_{i=1}^n (w^T x_i - y_i) x_{ij}$$

- Unfortunately, the partial derivative for w_j depends on all $\{w_1, w_2, \dots, w_d\}$
 - I can't just "set equal to 0 and solve for w_j ".

Gradient and Critical Points in d-Dimensions

- Generalizing “set the derivative to 0 and solve” in d-dimensions:
 - Find ‘w’ where the **gradient** vector **equals the zero vector**.
- **Gradient** is vector with partial derivative ‘j’ in position ‘j’:

$$\nabla f(w) = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \frac{\partial f}{\partial w_2} \\ \vdots \\ \frac{\partial f}{\partial w_d} \end{bmatrix}$$

For linear least squares:

$$\nabla f(w) = \begin{bmatrix} \sum_{i=1}^n (w^T x_i - y_i) x_{i1} \\ \sum_{i=1}^n (w^T x_i - y_i) x_{i2} \\ \vdots \\ \sum_{i=1}^n (w^T x_i - y_i) x_{id} \end{bmatrix}$$

Claims for linear least square:

1. Finding a ‘w’ where $\nabla f(w) = 0$ can be done by solving a system of linear equations.
2. All ‘w’ where $\nabla f(w) = 0$ are minimizers.

Matrix/Norm Notation (MEMORIZE/STUDY THIS)

- To solve the d-dimensional least squares, we use **matrix notation**:
 - We use ' y ' as an "n times 1" vector containing target ' y_i ' in position 'i'.
 - We use ' x_i ' as a "d times 1" vector containing features 'j' of example 'i'.
 - We're now going to be careful to make sure these are **column vectors**.
 - So ' X ' is a matrix with the x_i^T in row 'i'.

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{id} \end{bmatrix}$$

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1d} \\ x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{nd} \end{bmatrix} = \begin{bmatrix} \text{---} x_1^T \text{---} \\ \text{---} x_2^T \text{---} \\ \vdots \\ \text{---} x_n^T \text{---} \end{bmatrix}$$

Matrix/Norm Notation (MEMORIZE/STUDY THIS)

- To solve the d-dimensional least squares, we use **matrix notation**:
 - Our **prediction for example 'i'** is given by the **scalar** $w^T x_i$.
 - Our **predictions for all 'i'** (n times 1 vector) is the **matrix-vector product** Xw .

$$\hat{y}_i = w^T x_i$$

Also, because $w^T x_i$ is a scalar,
we have $w^T x_i = x_i^T w$.
(e.g., $[5]^T = [5]$)

$$\hat{\mathbf{y}} = Xw = \underbrace{\begin{bmatrix} - & x_1^T & - \\ - & x_2^T & - \\ & \vdots & \\ - & x_n^T & - \end{bmatrix}}_X \underbrace{\begin{bmatrix} 1 \\ w \\ 1 \end{bmatrix}}_w = \begin{bmatrix} x_1^T w \\ x_2^T w \\ \vdots \\ x_n^T w \end{bmatrix}$$

Prediction for example 'i' is in row 'i'.

Matrix/Norm Notation (MEMORIZE/STUDY THIS)

- To solve the d-dimensional least squares, we use **matrix notation**:
 - Our **prediction for example 'i'** is given by the **scalar** $w^T x_i$.
 - Our **predictions for all 'i'** (n times 1 vector) is the **matrix-vector product** Xw .
 - **Residual vector 'r'** gives difference between prediction and y_i (n times 1).
 - **Least squares can be written as the squared L2-norm of the residual.**

$$f(w) = \sum_{i=1}^n (\underbrace{w^T x_i}_{r_i} - y_i)^2 = \sum_{i=1}^n (r_i)^2$$

$$r = \hat{y} - y = \underbrace{Xw}_{\hat{y}} - y = \begin{bmatrix} w^T x_1 \\ w^T x_2 \\ \vdots \\ w^T x_n \end{bmatrix} - \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} w^T x_1 - y_1 \\ w^T x_2 - y_2 \\ \vdots \\ w^T x_n - y_n \end{bmatrix}$$

r_2 is difference for example 2.

$$= \sum_{i=1}^n r_i r_i$$

$$= r^T r$$

$$= \|r\|^2 = \|Xw - y\|^2$$

Summary

- **Regression** considers the case of a numerical y_i .
- **Least squares** is a classic method for fitting linear models.
 - With 1 feature, it has a simple closed-form solution.
 - Can be generalized to 'd' features.
- **Gradient** is vector containing partial derivatives of all variables.
- **Matrix notation** for expressing least squares problem.
- Next time: minimizing $\frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2$ in terms of 'w' is:

$$w = (X^T X)^{-1} (X^T y)$$

(in Julia)