

CPSC 340: Machine Learning and Data Mining

Convolutions

Fall 2019

Last Time: “Global” and “Local” Features

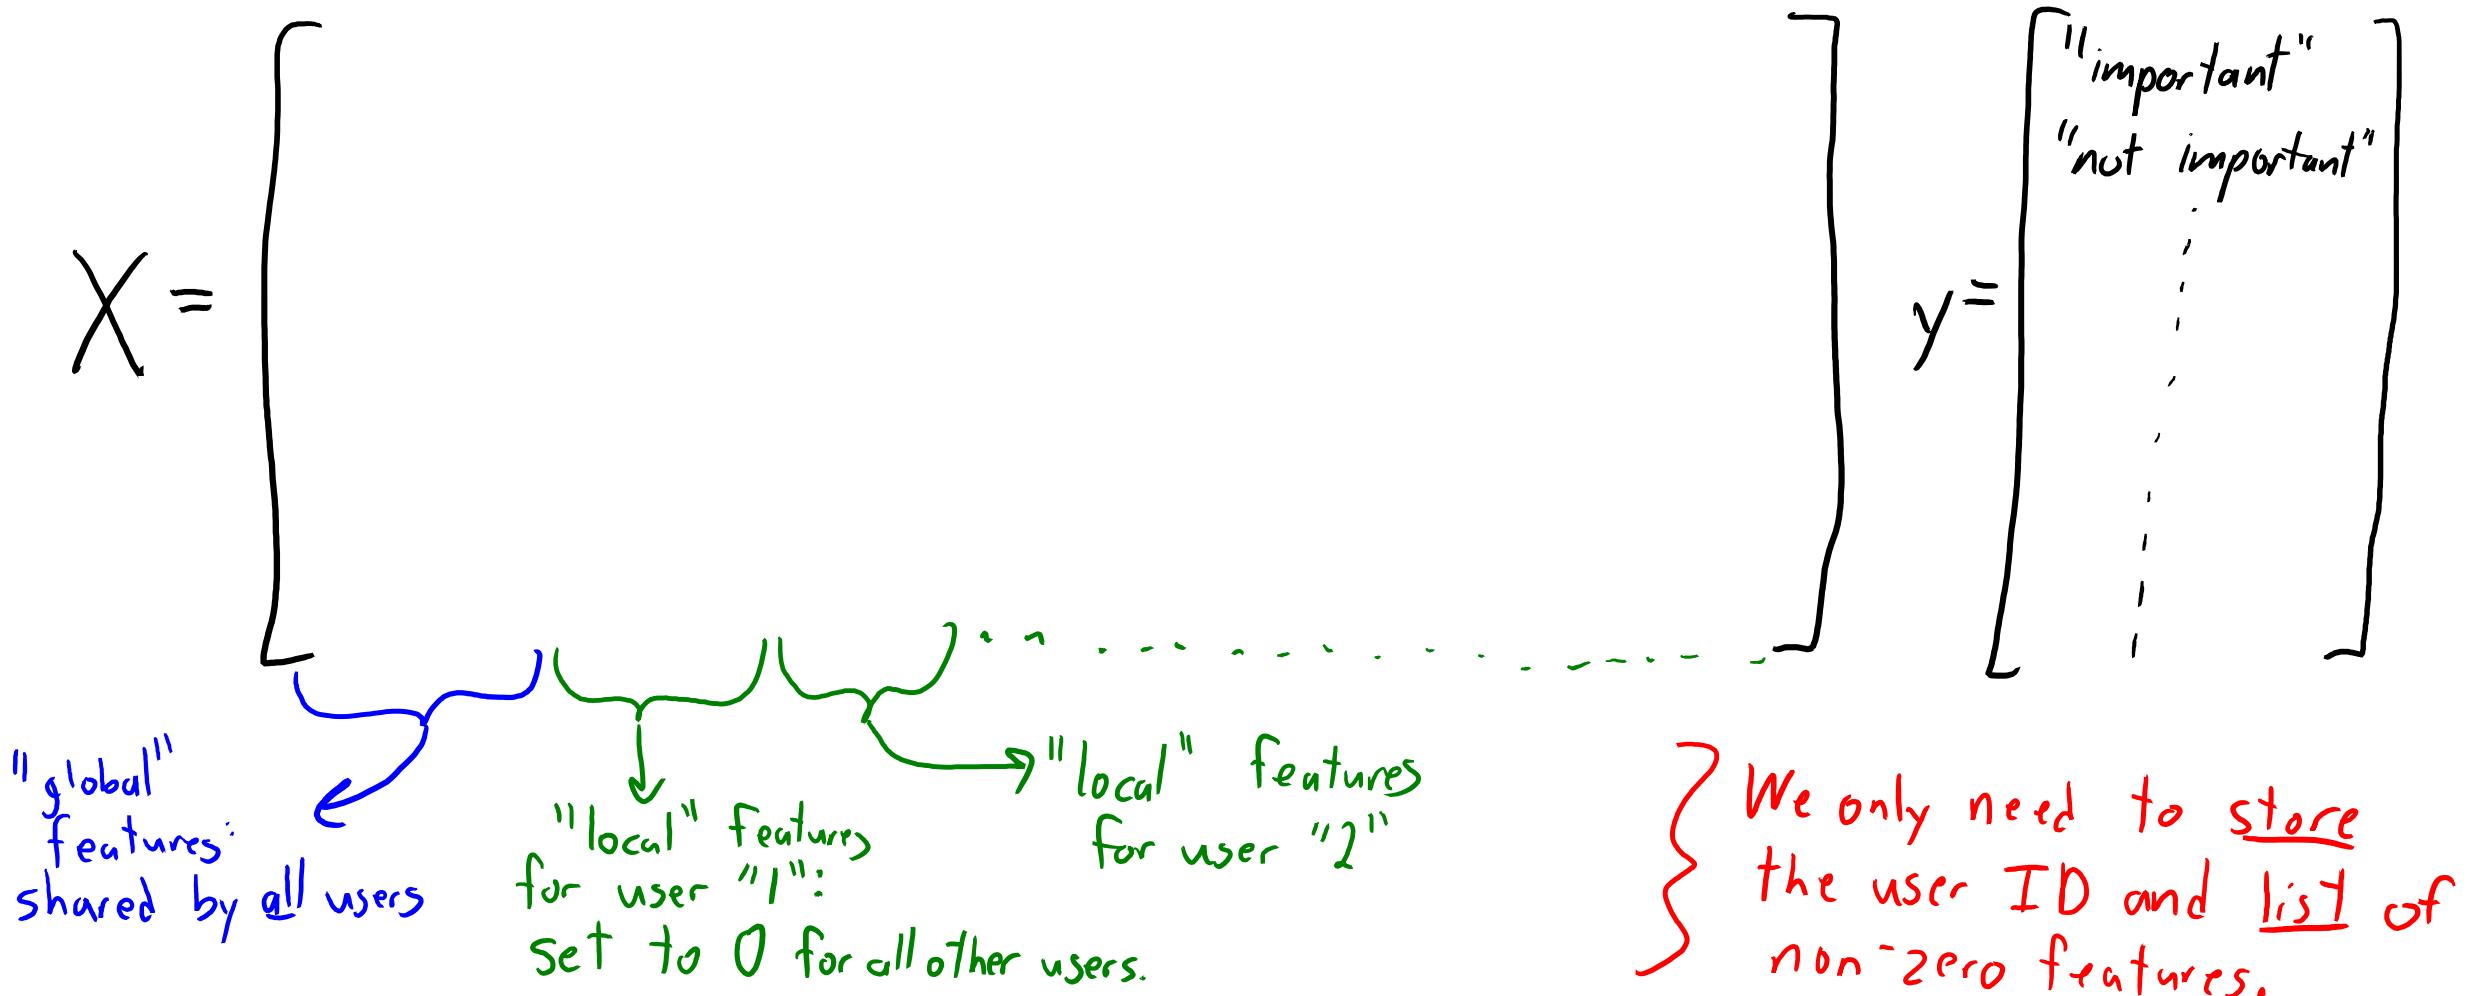
- Consider the following weird feature transformation for identifying important e-mails:

“CPSS”	“340”	“CPSC” (any user)	“340” (any user)	“CPSC” (user?)	“340” (user?)
1	0	1	0	User 1	<no “340”>
1	0	1	0	User 1	<no “340”>
1	1	1	1	User 2	User 2
0	0	0	0	<no “CPSC”>	<no “340”>
1	1	1	1	User 3	User 3

- The categorical (user?) features get expanded out into ‘k’ binary features.
 - Where ‘k’ is the number of users.
 - All those features are set to 0 if the word was not used.
- “Any user” (“global”) features increase/decrease importance of word for **every user**.
- “User” (“local”) features increase/decrease importance of word for **specific users**.
 - Lets us learn more about users where we have a lot of data

The Big Global/Local Feature Table for E-mails

- Each row is one e-mail (there are lots of rows):



Predicting Importance of E-mail For New User

- Consider a new user:
 - We start out with no information about them.
 - So we use **global** features to predict what is important to a generic user.

$$\hat{y}_i = \text{sign}(\underbrace{w_g^\top x_{ig}}_{\text{features/weights shared across users.}})$$

- Local features are initialized to zero.
- With more data, update **global** features and **user's local** features:
 - Local features make prediction *personalized*.

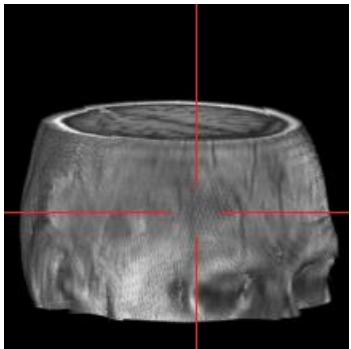
$$\hat{y}_i = \text{sign}(w_g^\top x_{ig} + \underbrace{w_u^\top x_{iu}}_{\text{features/weights specific to user.}})$$

- What is important to *this* user?
- G-mail system: classification with **logistic regression**.
 - Trained with a variant of **stochastic gradient** (later).

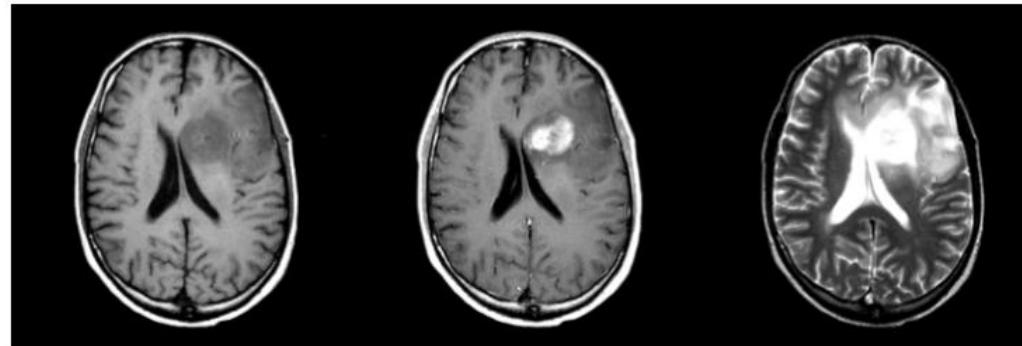
(pause)

Motivation: Automatic Brain Tumor Segmentation

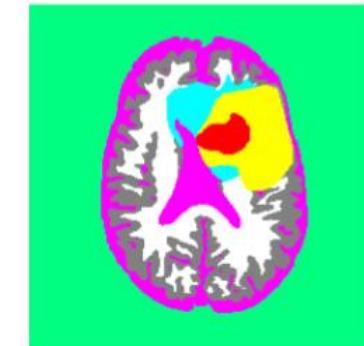
- Task: segmentation tumors and normal tissue in multi-modal MRI data.



Input:



Output:



- Applications:

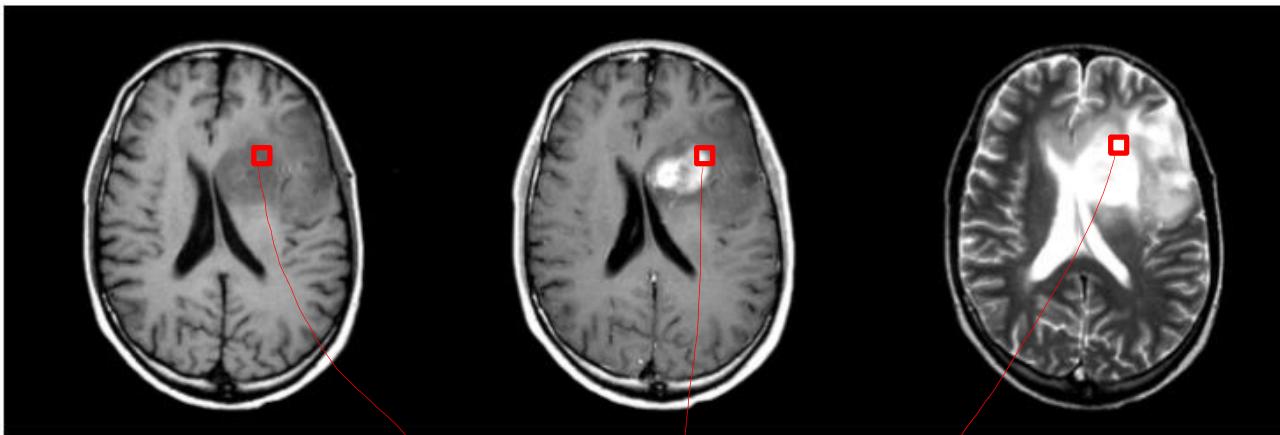
- Radiation therapy target planning, quantifying treatment responses.
 - Mining growth patterns, image-guided surgery.

- Challenges:

- Variety of tumor appearances, similarity to normal tissue.
 - “You are never going to solve this problem.”

Naïve Voxel-Level Classifier

- We could treat classifying a voxel as **supervised learning**:



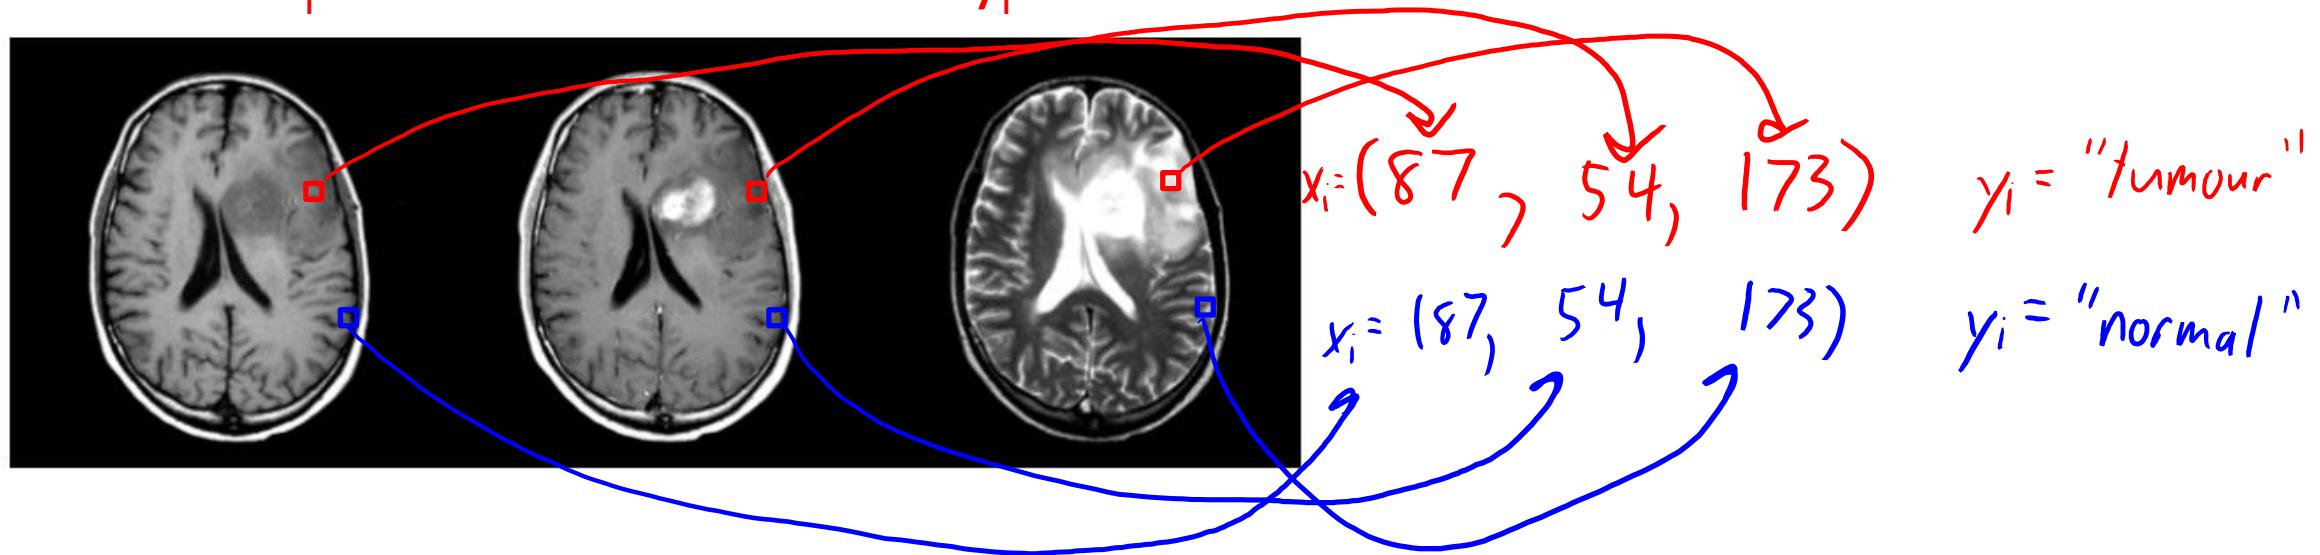
$$x_i = (98, 187, 246)$$

$y_i = \text{"tumour"}$

- We can formulate predicting y_i given x_i as supervised learning.
- But it **doesn't work** at all with these features.

Need to Summarize Local Context

- The individual voxel values are almost meaningless:
 - This x_i could lead to different y_i .



- Intensities not standardized.
- Non-trivial overlap in signal for different tissue types.
- “Partial volume” effects at boundaries of tissue types.

Need to Summarize Local Context

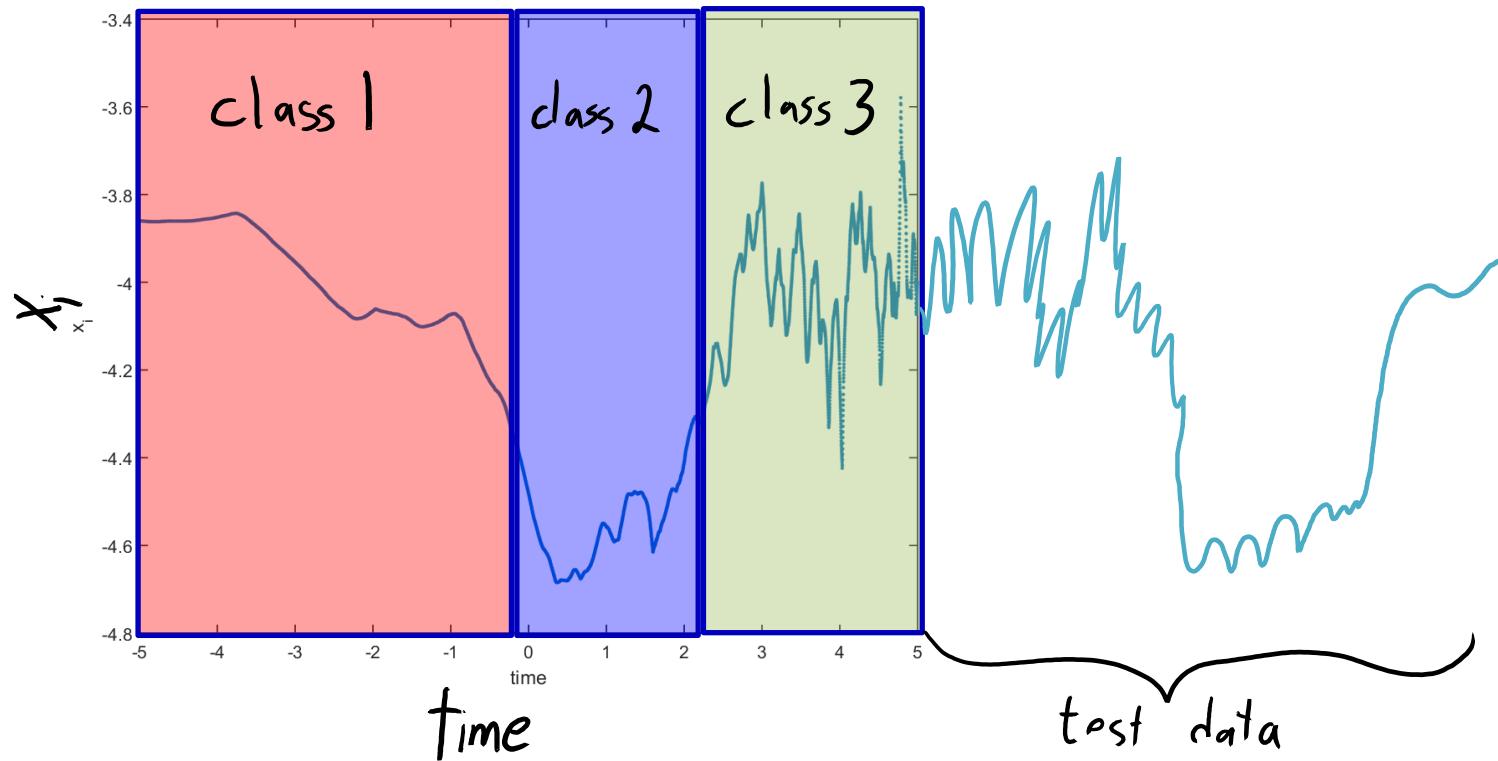
- We need to represent the spatial “context” of the voxel.



- Include all the values of **neighbouring voxels** as extra features?
 - Variation on coupon collection problem: **requires lots of data** to find patterns.
- Measure neighbourhood **summary statistics** (mean, variance, histogram)?
 - Variation on bag of words problem: loses **spatial information** present in voxels.
- Standard approach uses **convolutions** to represent neighbourhood.

Representing Neighbourhoods with Convolutions

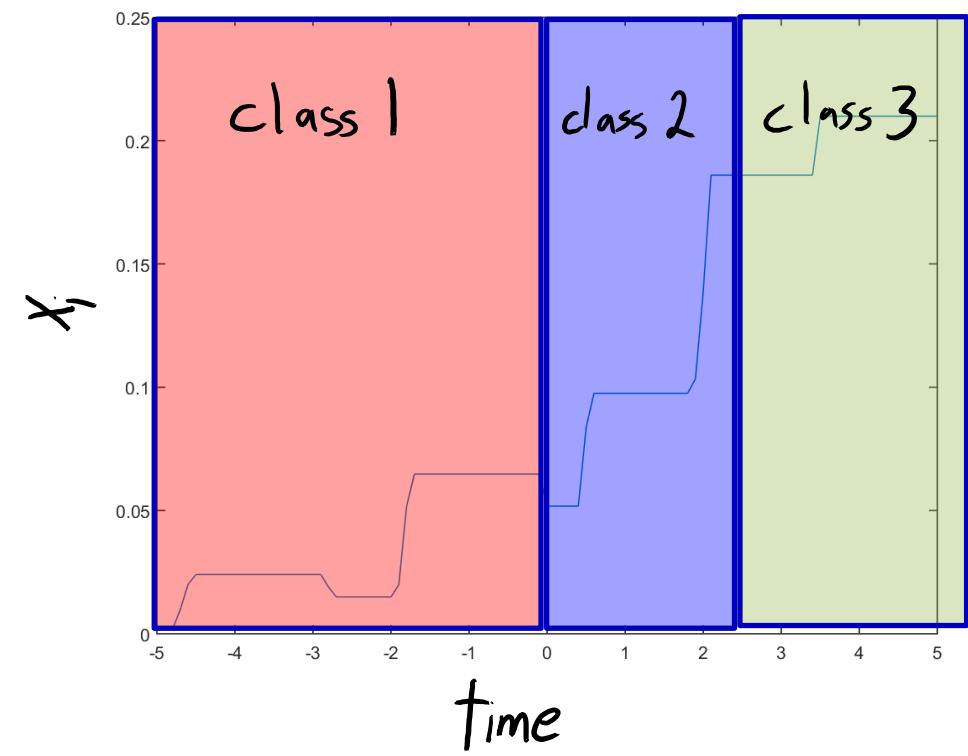
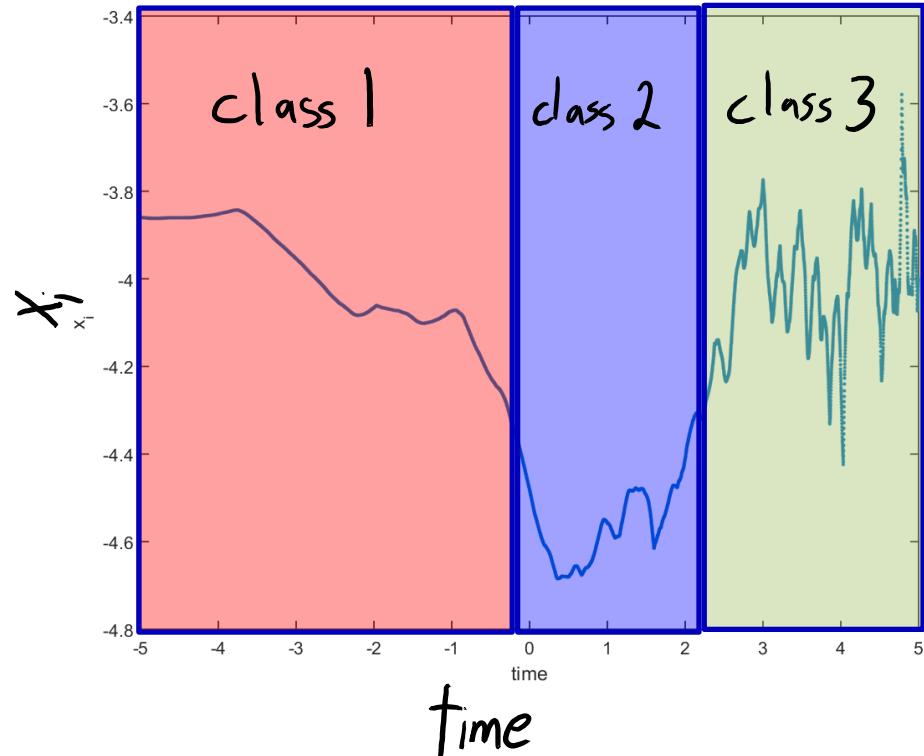
- Consider a 1D dataset:
 - Want to classify each time into y_i in $\{1,2,3\}$.
 - Example: speech data.



- Easy to distinguish class 2 from the other classes (x_i are smaller).
- Harder to distinguish between class 1 and class 3 (similar x_i range).
 - But convolutions can represent that class 3 is in “spiky” region.

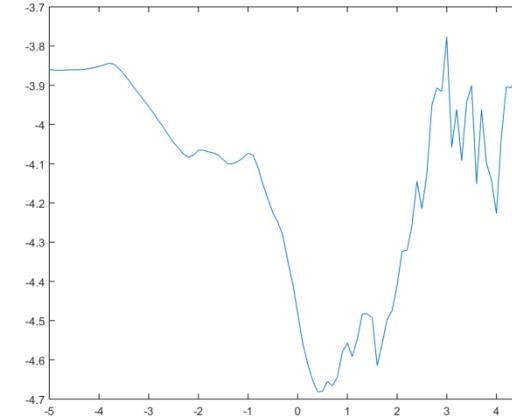
Representing Neighbourhoods with Convolutions

- Original features (left) and features from **convolutions** (right):



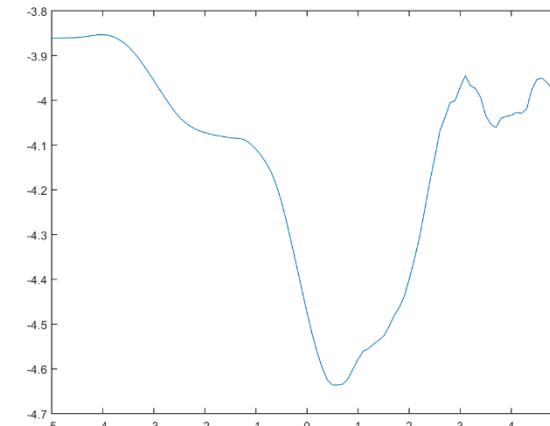
- Easy to distinguish the 3 classes with these 2 features.

1D Convolution Example

- Consider our original “signal”: 
- For each “time”:
 - Compute dot-product of signal at surrounding times with a “filter”.

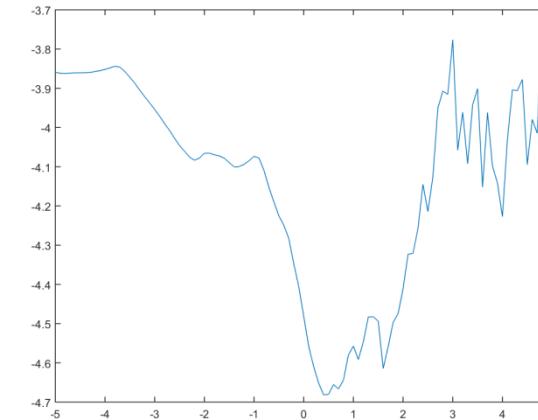
$$w = \left[\frac{1}{9} \quad \frac{1}{9} \right]$$

- This gives a new “signal”:
 - Measures a property of “neighbourhood”.
 - This particular filter shows a local “average” value.



1D Convolution Example

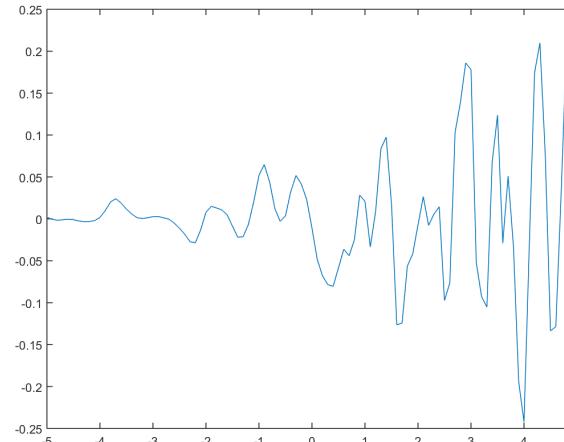
- Consider our original “signal”:



- For each “time”:
 - Compute dot-product of signal at surrounding times with a “filter”.

$$w = [-0.1416 \ -0.1781 \ -0.2746 \ 0.1640 \ 0.8607 \ 0.1640 \ -0.2746 \ -0.1781 \ -0.1416]$$

- This gives a new “signal”:
 - Measures a property of “neighbourhood”.
 - This particular filter shows a local “how spiky” value.



1D Convolution (notation is specific to this lecture)

- **1D convolution** input:

- Signal ‘x’ which is a vector length ‘n’.

- Indexed by $i=1,2,\dots,n$.

- Filter ‘w’ which is a vector of length ‘ $2m+1$ ’:

- Indexed by $i=-m,-m+1,\dots,-2,0,1,2,\dots,m-1,m$

$$x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$$

$$w = [0 \ -1 \ 2 \ -1 \ 0]$$

$w_{-2} \quad w_{-1} \quad w_0 \quad w_1 \quad w_2$

- Output is a vector of length ‘n’ with elements:

$$z_i = \sum_{j=-m}^m w_j x_{i+j}$$

- You can think of this as **centering w at position ‘i’**,
and **taking a dot product of ‘w’ with that “part” x_i** .

1D Convolution

- 1D convolution example:

- Signal:

$$x = [0 \boxed{1} \ 1 \ \textcircled{2} \ 3 \ 5 \ 8 \ 13]$$

i s n

Let's compute z_4 :

- Filter:

$$w = [0 \ -1 \ 2 \ -1 \ 0]$$

w_{-2} w_{-1} w_0 w_1 w_2

$$x_{4-2:4+2} = [1 \ 1 \ 2 \ 3 \ 5]$$

- Convolution:

$$z = [1 \ 2 \ 3 \ \textcircled{4} \ 5 \ 6 \ 7 \ 8]$$

n

Multiply element-wise

$$[0 \ -1 \ 4 \ -3 \ 0]$$

Add

$$z_4 = 0 - 1 + 4 - 3 + 0 = \underline{0}$$

1D Convolution

- 1D convolution example:

- Signal:

$$x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$$

Diagram showing the signal x as a sequence of values. A red box highlights the value 3 . Below the sequence, indices n are labeled from 1 to 8 , corresponding to each value.

- Filter:

$$w = [0 \ -1 \ 2 \ -1 \ 0]$$

Diagram showing the filter w as a sequence of weights. The weights are labeled $w_{-2}, w_{-1}, w_0, w_1, w_2$ below the sequence.

- Convolution:

$$z = [\underbrace{1 \ 2 \ 3 \ 4}_n \ 0 \ \boxed{-1} \ 5 \ 6 \ 7 \ 8]$$

Diagram showing the result of the convolution step. The output z is shown as a sequence of values. A red box highlights the value -1 .

Let's compute z_5 :

$$[1 \ 2 \ 3 \ 5 \ 8]$$

Multiply ↓

$$[0 \ -2 \ 6 \ -5 \ 6]$$

Add ↓

$$z_5 = -2 + 6 - 5 = -1$$

1D Convolution Examples

- Examples:

- “Identity”

$$\hookrightarrow w = [0 \ 1 \ 0]$$

Let $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$

$$z = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$$

$0 \cdot x_0 + 1 \cdot x_1 + 0 \cdot x_2$ $0 \cdot x_1 + 1 \cdot x_2 + 0 \cdot x_3$

- “Translation”

$$\hookrightarrow w = [0 \ 0 \ 1]$$

$$z = [1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13 \ ?]$$

$0 \cdot x_0 + 0 \cdot x_1 + 1 \cdot x_2$

1D Convolution Examples

- Examples:

- “Identity”

$$\hookrightarrow w = [0 \ 1 \ 0]$$

Let $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$

The diagram shows the input vector x as a horizontal sequence of numbers: 0, 1, 1, 2, 3, 5, 8, 13. Two green curly braces are placed under the first three elements (0, 1, 1) and the last three elements (5, 8, 13), each labeled "average". This indicates that the kernel w of size 3 is applied across overlapping windows of size 3 in the input x .

- “Local Average”

$$\hookrightarrow w = [\frac{1}{3} \ \frac{1}{3} \ \frac{1}{3}]$$

$$z = [? \ \frac{2}{3} \ \frac{1}{3} \ 2 \ \frac{3\frac{1}{3}}{3} \ \frac{5\frac{1}{3}}{3} \ \frac{8\frac{2}{3}}{3} ?]$$

Boundary Issue

- What can we do about the “?” at the edges?

If $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$ and $w = [\frac{1}{3} \ \frac{1}{3} \ \frac{1}{3}]$ then $z = [? \ \frac{2}{3} \ 1\frac{1}{3} \ 2 \ 3\frac{1}{3} \ 5\frac{1}{3} \ 8\frac{2}{3} \ ?]$

- Can assign values past the boundaries:

- “Zero”: $x = [0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 3 \ 5 \ 5 \ 8 \ 8 \ 13] \ 0 \ 0 \ 0$

- “Replicate”: $x = [0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 3 \ 5 \ 5 \ 8 \ 8 \ 13] \ 13 \ 13 \ 13$

- “Mirror”: $x = [2 \ 1 \ 1 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 3 \ 5 \ 5 \ 8 \ 8 \ 13] \ 8 \ 5 \ 3$

- Or just ignore the “?” values and return a shorter vector:

$$z = [\frac{2}{3} \ 1\frac{1}{3} \ 2 \ 3\frac{1}{3} \ 5\frac{1}{3} \ 8\frac{2}{3}]$$

Formal Convolution Definition

- We've defined the convolution as:

$$z_i = \sum_{j=-m}^m w_j x_{i+j}$$

- In other classes you may see it defined as:

$$z_i = \sum_{j=-m}^m w_j x_{i-j}$$

(reverses 'w')

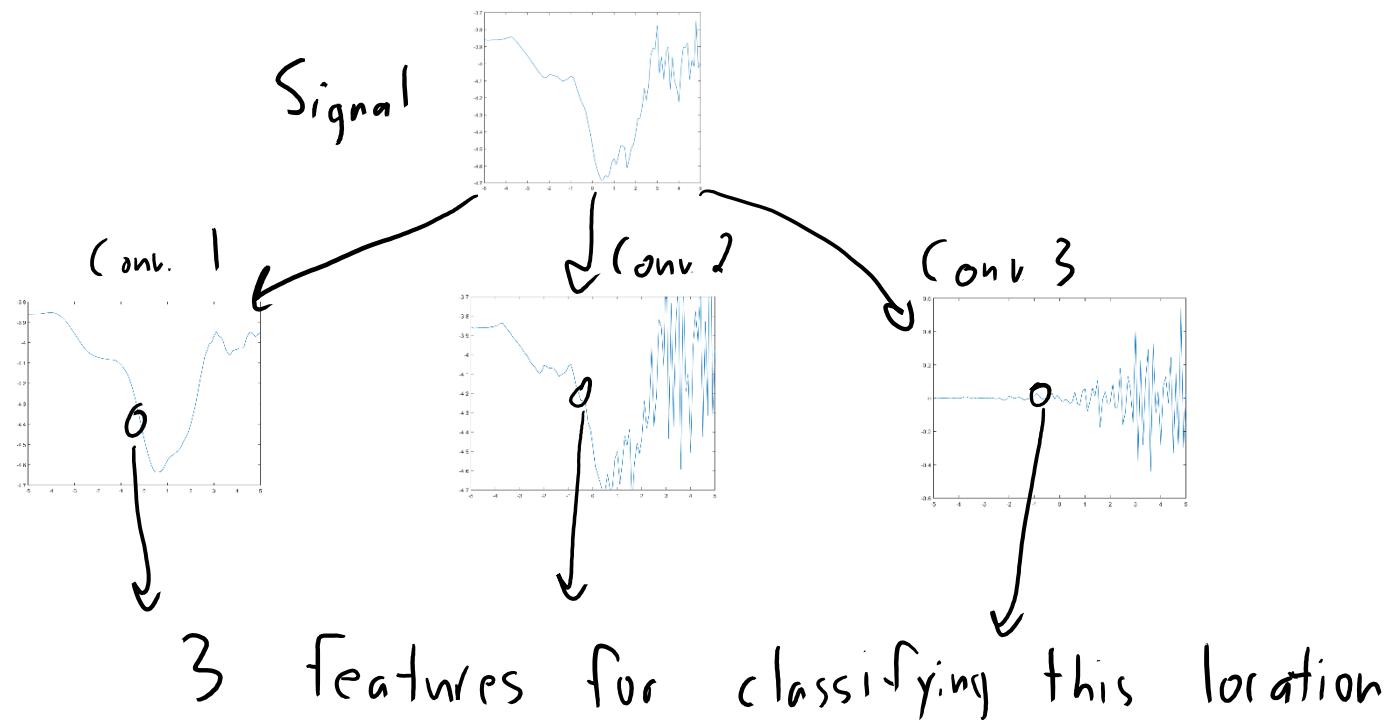
$$z_i = \int_{-\infty}^{\infty} w_j x_{i-j} dj$$

(assumes signal + filter are continuous)

- For simplicity we're skipping the "reverse" step, and assuming 'w' and 'x' are sampled at discrete points (not functions).
- But keep this mind if you read about convolutions elsewhere.

Convolutions: Big Picture

- How do you use convolutions to get features?
 - Apply **several different convolutions** to your signal/image.
 - Each convolution gives a different “signal/image” value at each location.
 - Use **theses different signal/image values** to give **features** at each location.



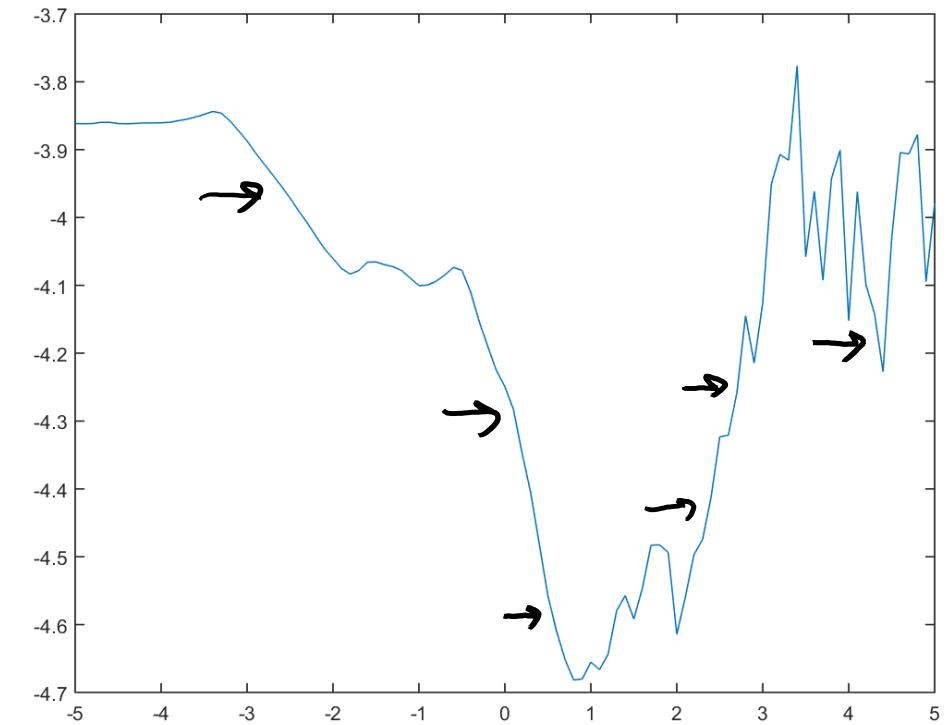
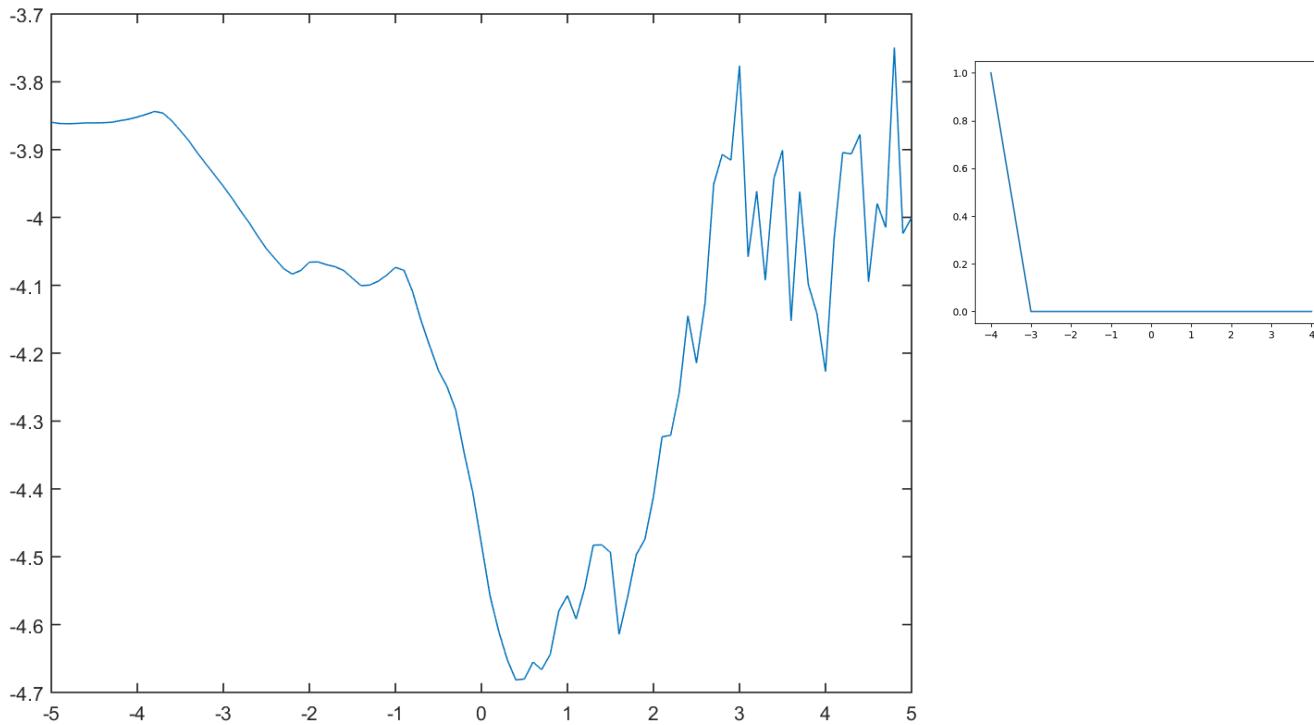
Convolutions: Big Picture

- What can features coming from convolutions represent?
 - Some filters give you an **average value of the neighbourhood**.
 - Some filters **approximate the “first derivative”** in the neighbourhood.
 - “Is there a change from low to high (or dark to bright)?”
 - Some filters **approximate the “second derivative”** in the neighbourhood.
 - “Is there a spike or is the signal speeding up?”
- Hope: we can characterize **“what happens in a neighbourhood”**, with just a few numbers.

1D Convolution Examples

- Translation convolution shift signal:
 - “What is my neighbour’s value?”

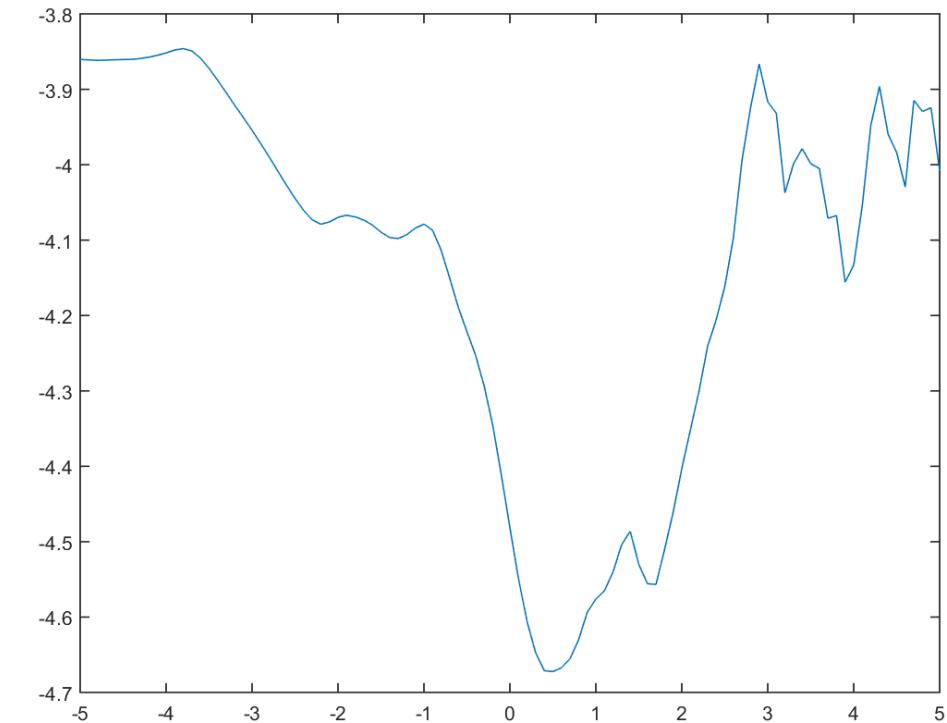
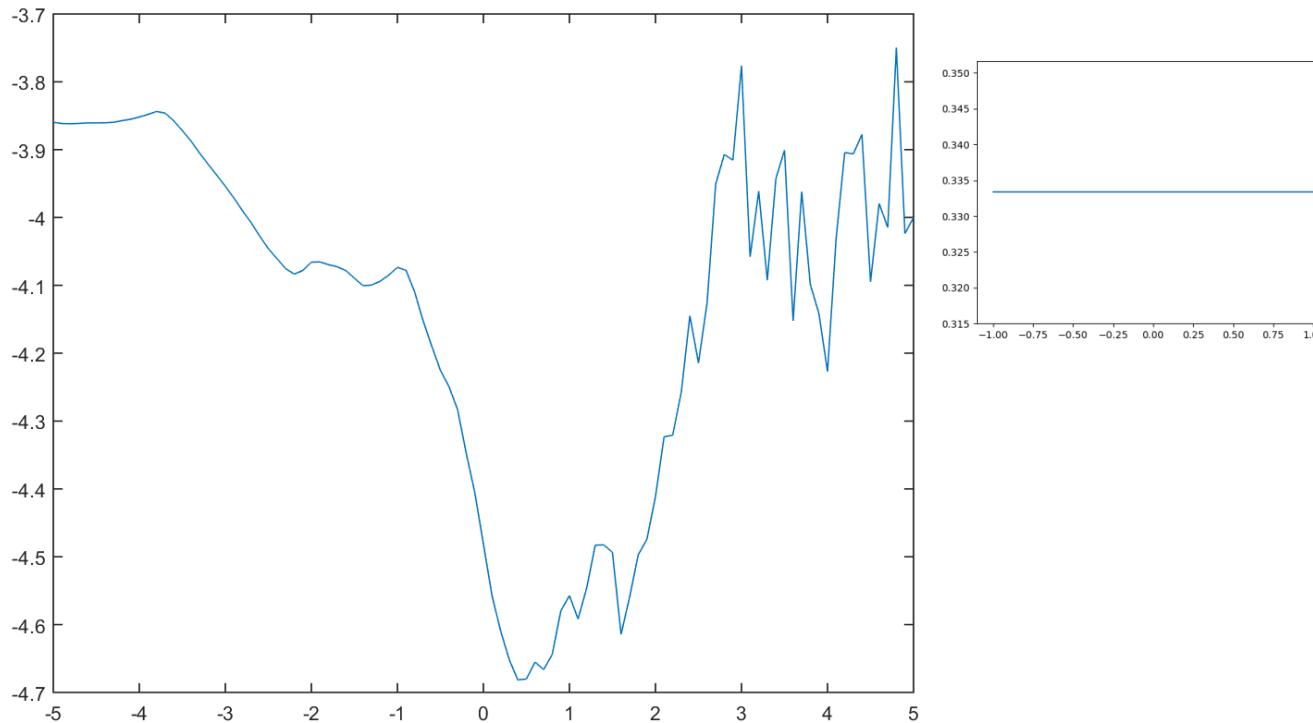
$$W = [1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$$



1D Convolution Examples

- **Averaging** convolution computes local mean:
 - This is **less sensitive to noise** (or spikes) than taking the raw signal value.

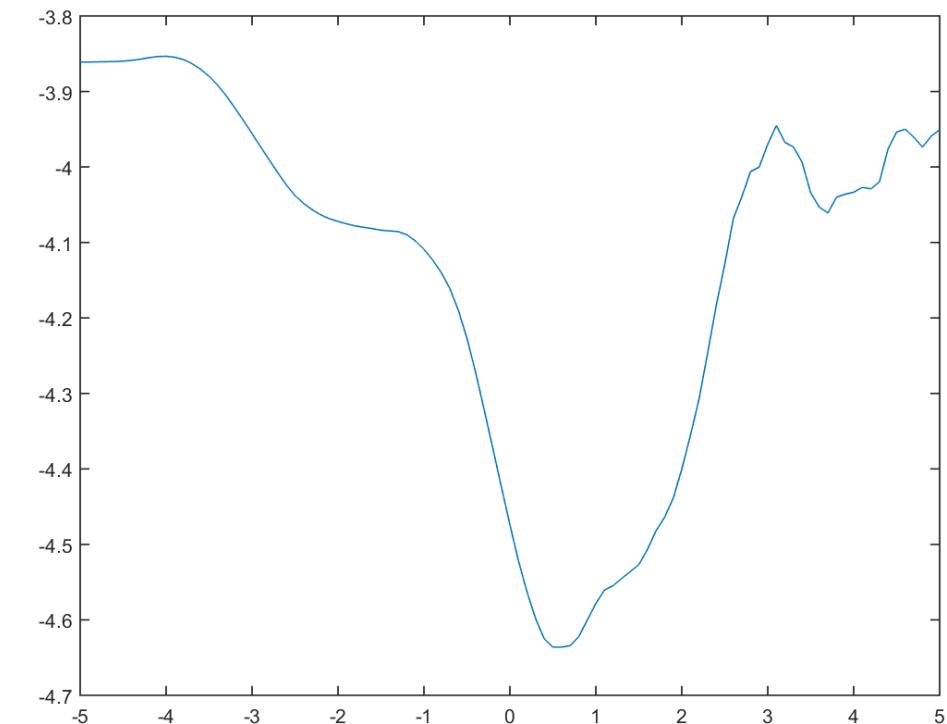
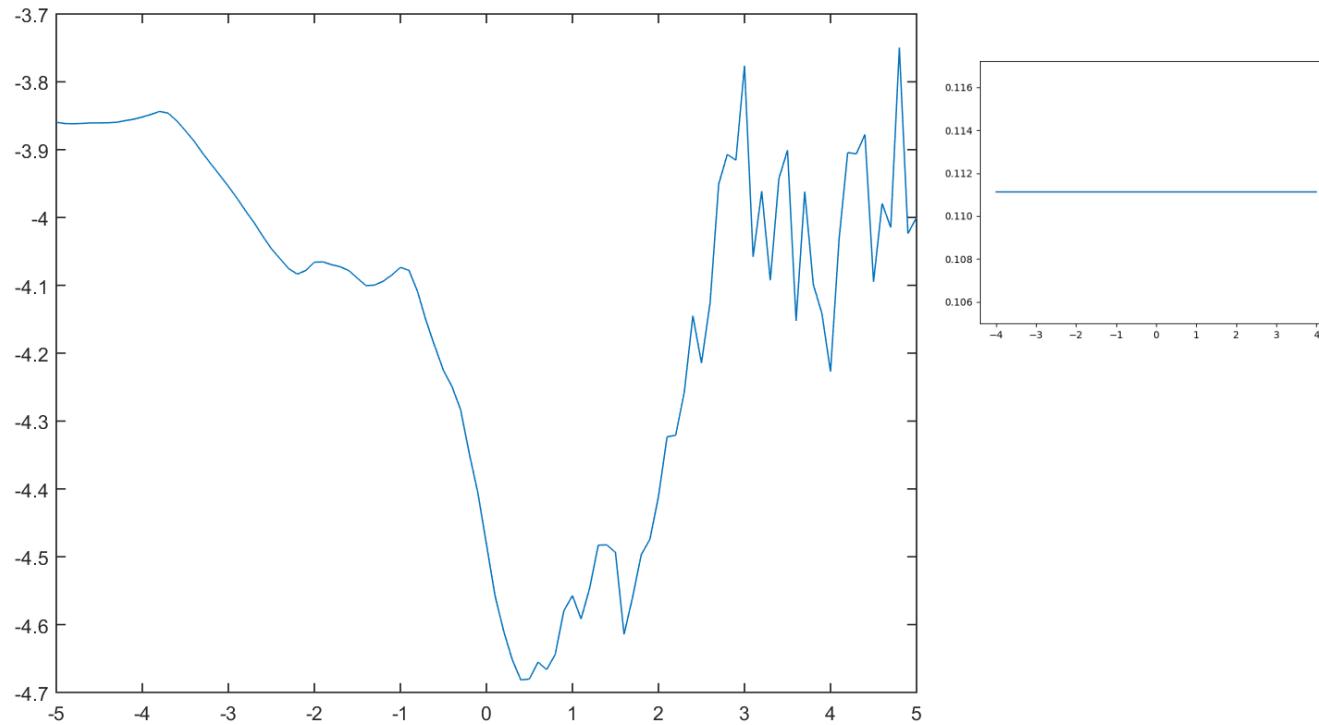
$$w = \left[\frac{1}{3} \quad \frac{1}{3} \quad \frac{1}{3} \right]$$



1D Convolution Examples

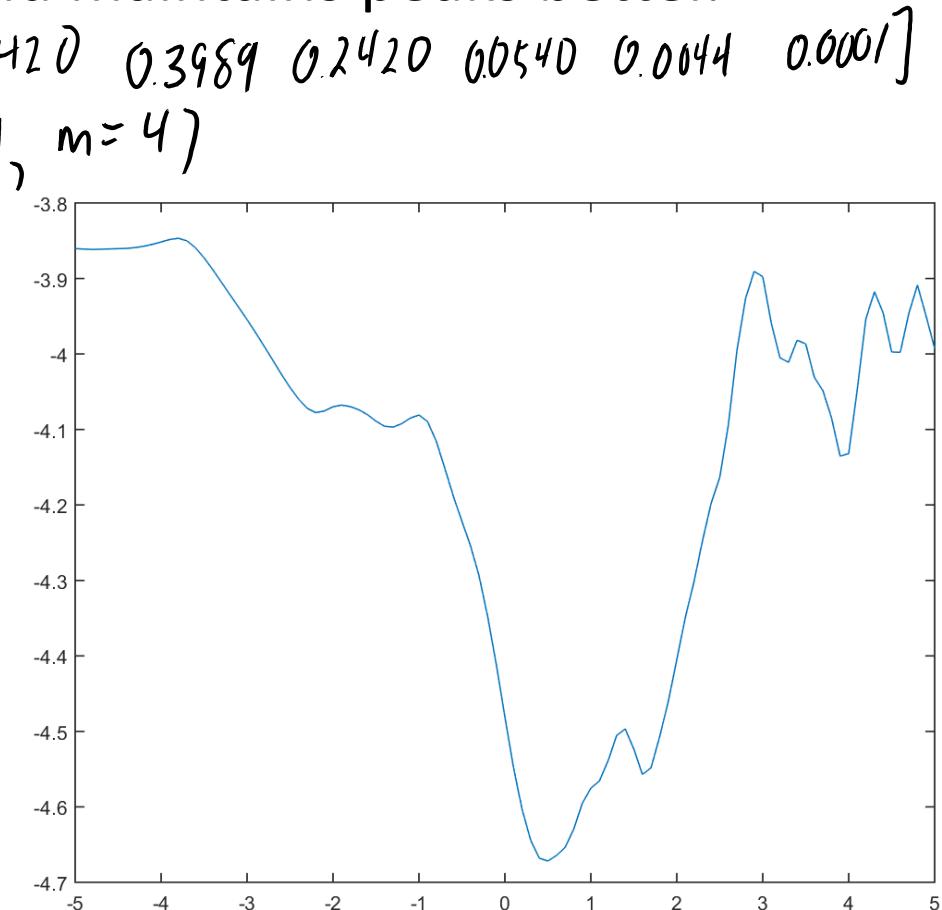
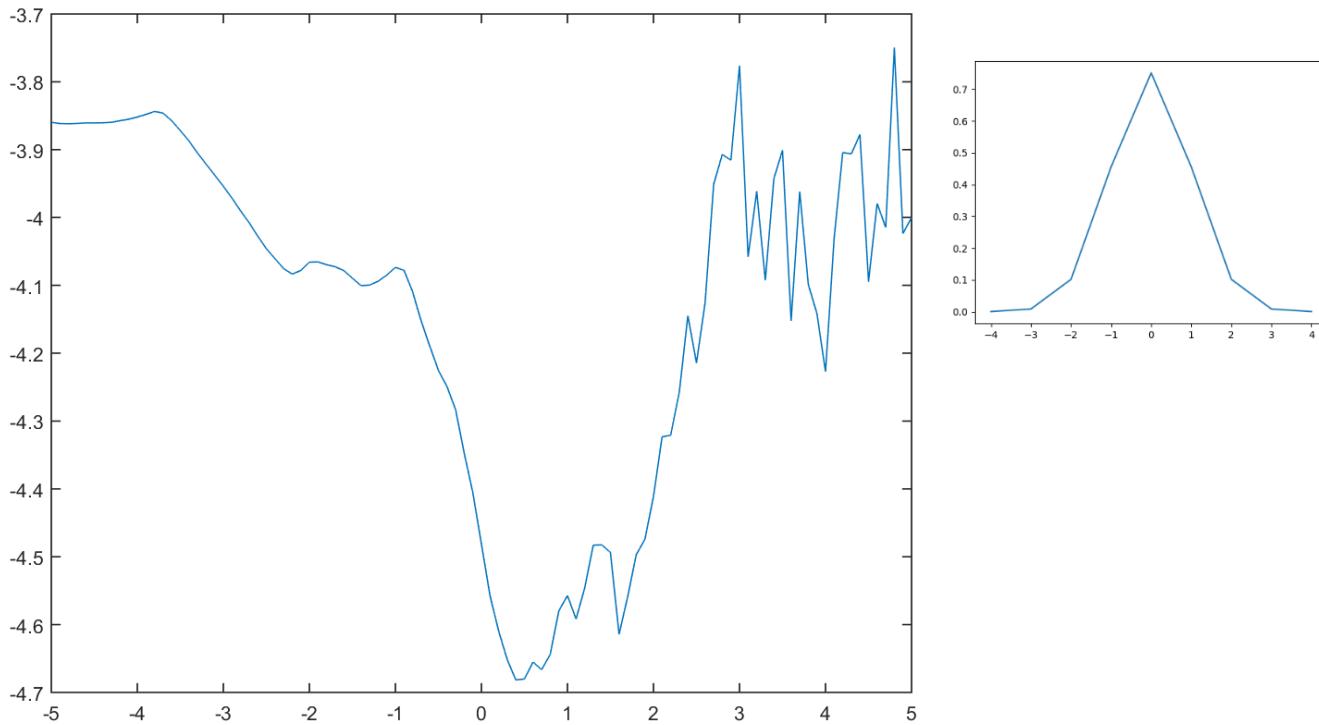
- **Averaging** over bigger window gives coarser view of signal:
 - “Is this the signal generally high in this region?”

$$w = \left[\frac{1}{9} \quad \frac{1}{9} \right]$$



1D Convolution Examples

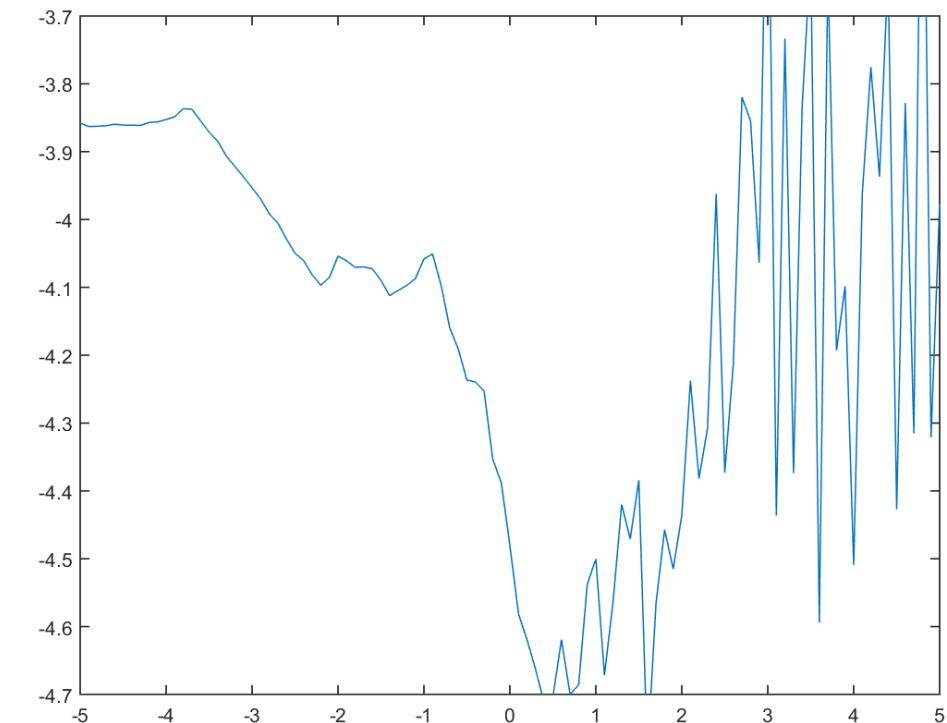
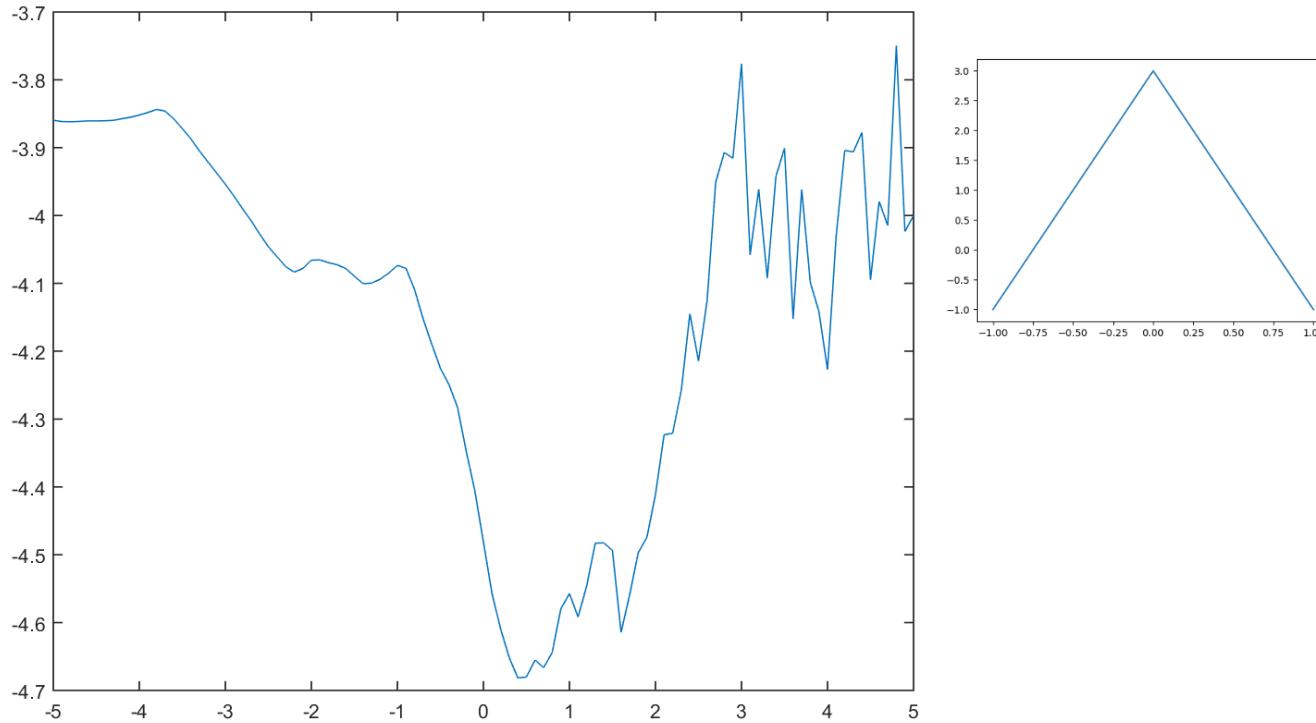
- **Gaussian convolution** “blurs” signal: $w_i \propto \exp\left(-\frac{i^2}{2\sigma^2}\right)$
 - Compared to averaging it's more smooth and maintains peaks better.
- $w = [0.0001 \ 0.0644 \ 0.0540 \ 0.1410 \ 0.3989 \ 0.2420 \ 0.0540 \ 0.0044 \ 0.0001]$
 $(\sigma = 1, m = 4)$



1D Convolution Examples

- **Sharpen** convolution enhances peaks.
 - An “average” that places **negative weights** on the surrounding pixels.

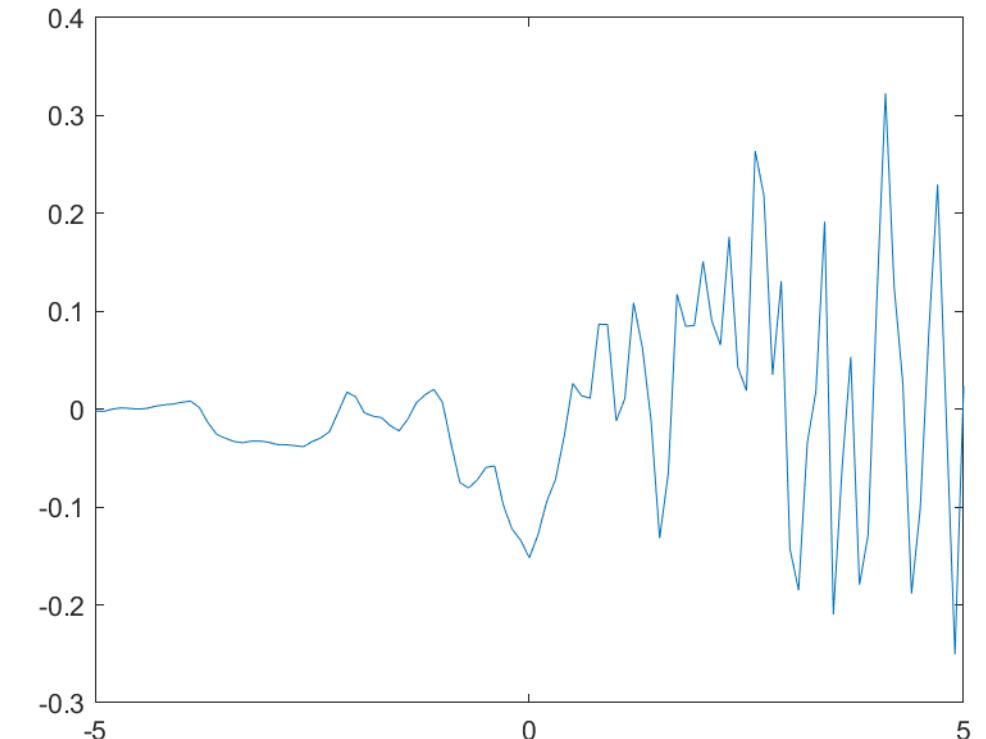
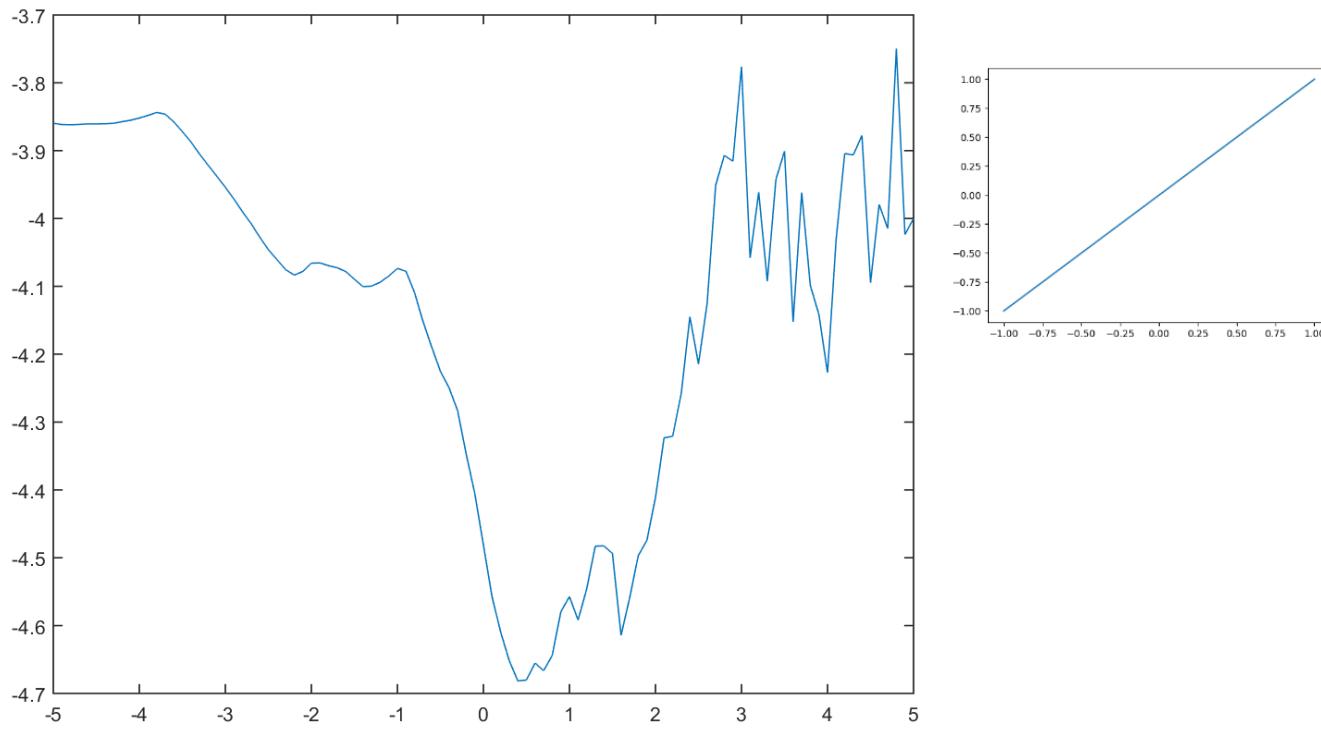
$$w = [-1 \quad 3 \quad -1]$$



1D Convolution Examples

- Centered difference convolution approximates first derivative:
 - Positive means change from low to high (negative means high to low).

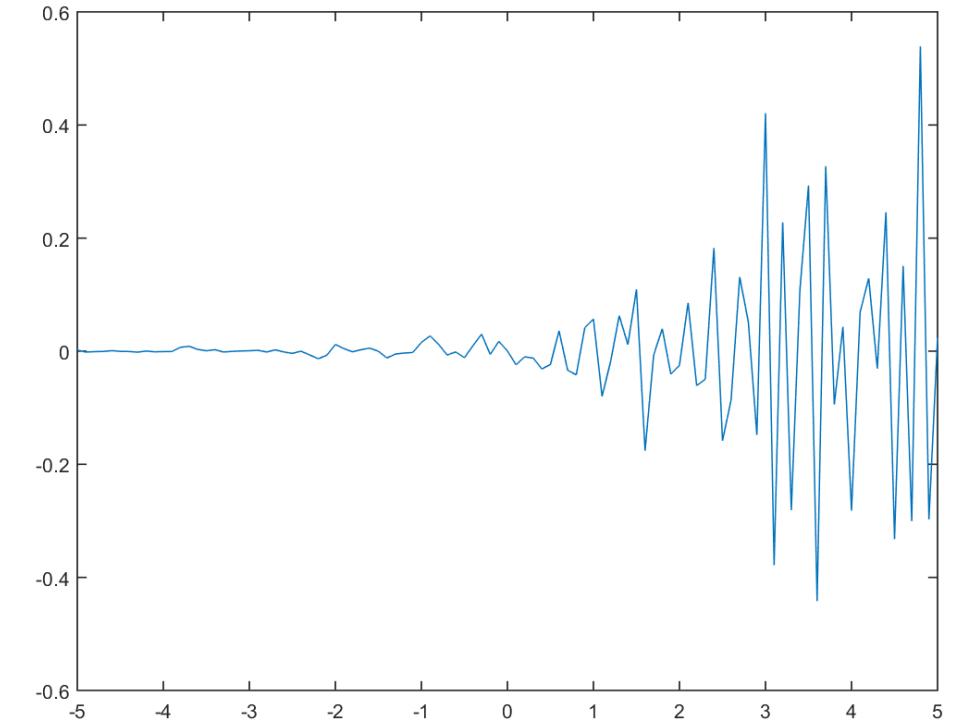
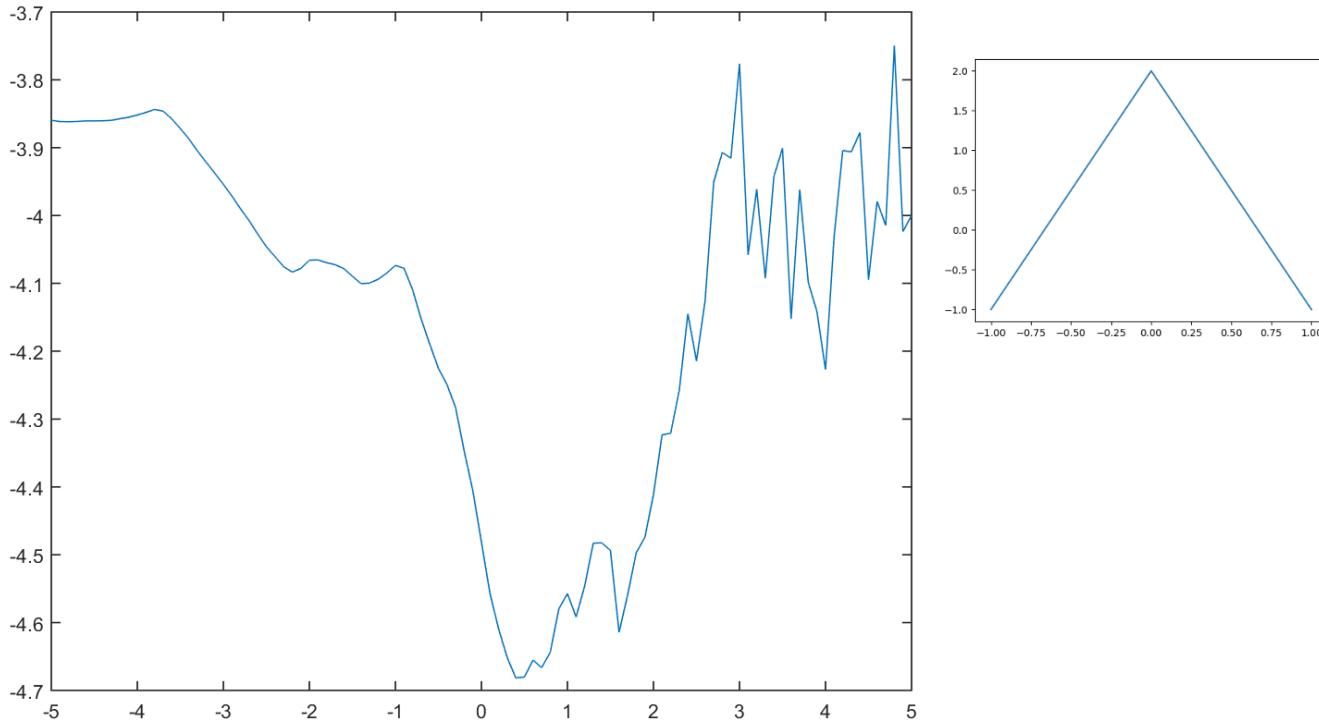
$$w = [-1 \quad 0 \quad 1]$$



1D Convolution Examples

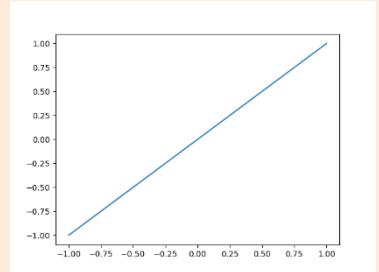
- Laplacian convolution approximates second derivative:
 - “Sum to zero” filters “respond” if input vector looks like the filter

$$w = [-1 \quad 2 \quad -1]$$

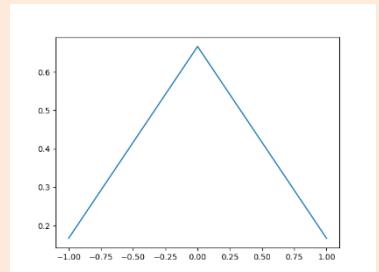


Digression: Derivatives and Integrals

- Numerical derivative approximations can be viewed as filters:
 - Centered difference: $[-1, 0, 1]$ (derivativeCheck in findMin).



- Numerical integration approximations can be viewed as filters:
 - “Simpson’s” rule: $[1/6, 4/6, 1/6]$ (a bit like Gaussian filter).



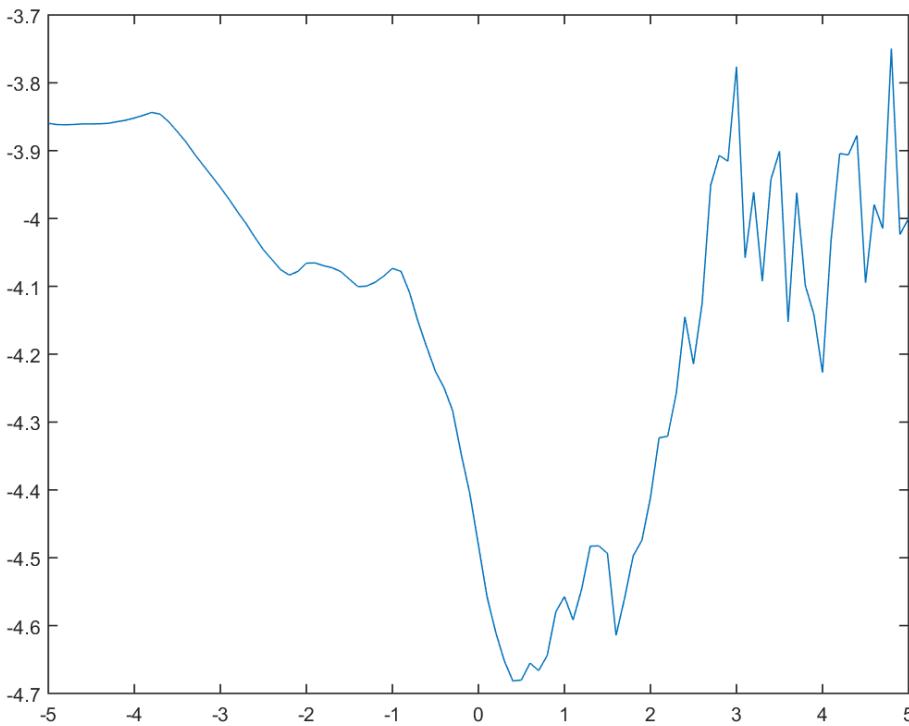
- Derivative filters add to 0, integration filters add to 1,
 - For constant function, derivative should be 0 and average = constant.

Laplacian of Gaussian Filter

- Laplacian of Gaussian is a smoothed 2nd-derivative approximation:

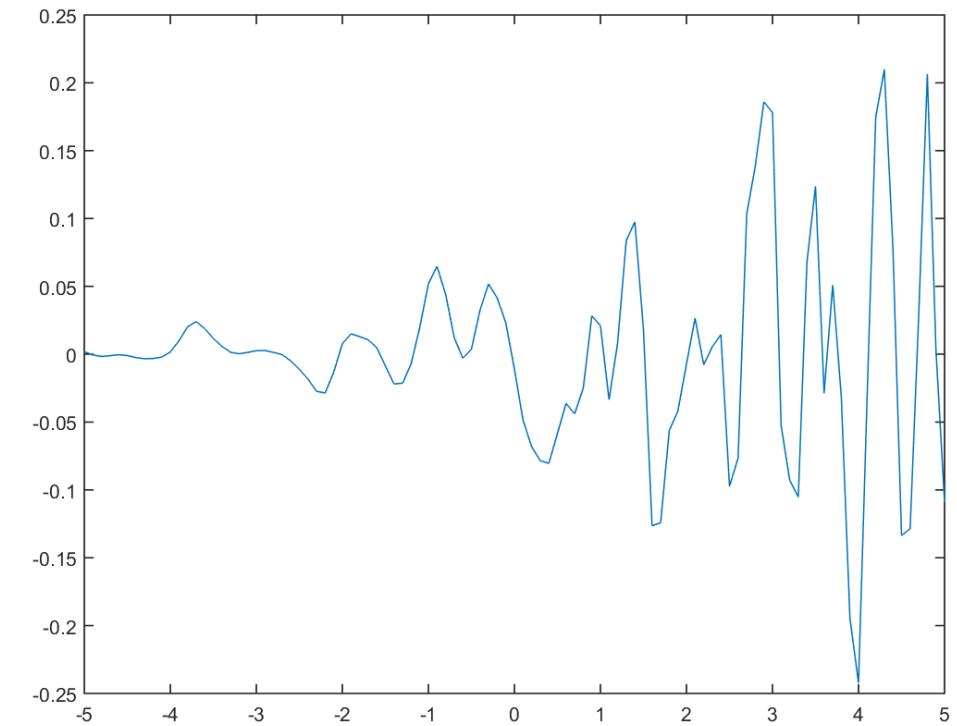
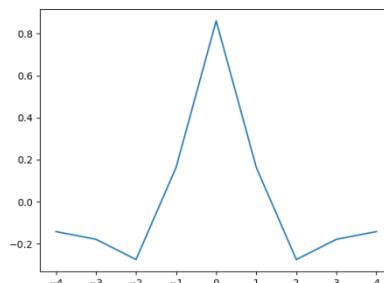
$$w_i = \left(1 - \frac{i^2}{2\sigma^2}\right) \exp\left(-\frac{i^2}{2\sigma^2}\right)$$

(then subtract mean)



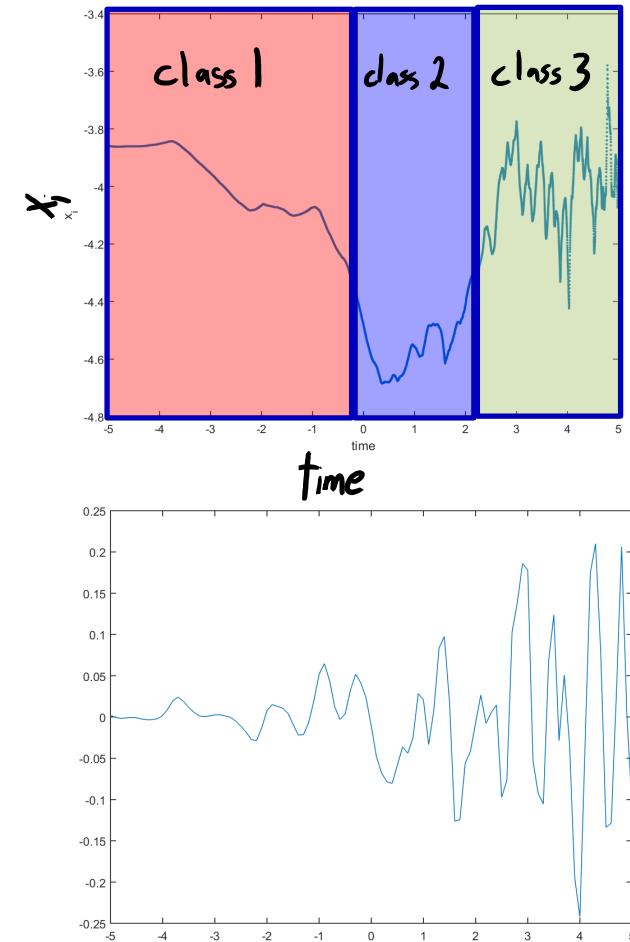
$$w = [-0.1416 \ -0.1781 \ -0.2746 \ 0.1640 \ 0.8607 \ 0.1640 \ -0.2746 \ -0.1781 \ -0.1416]$$

$(\sigma^2 = 1, m = 4)$



Taking Maximums of Convolutions

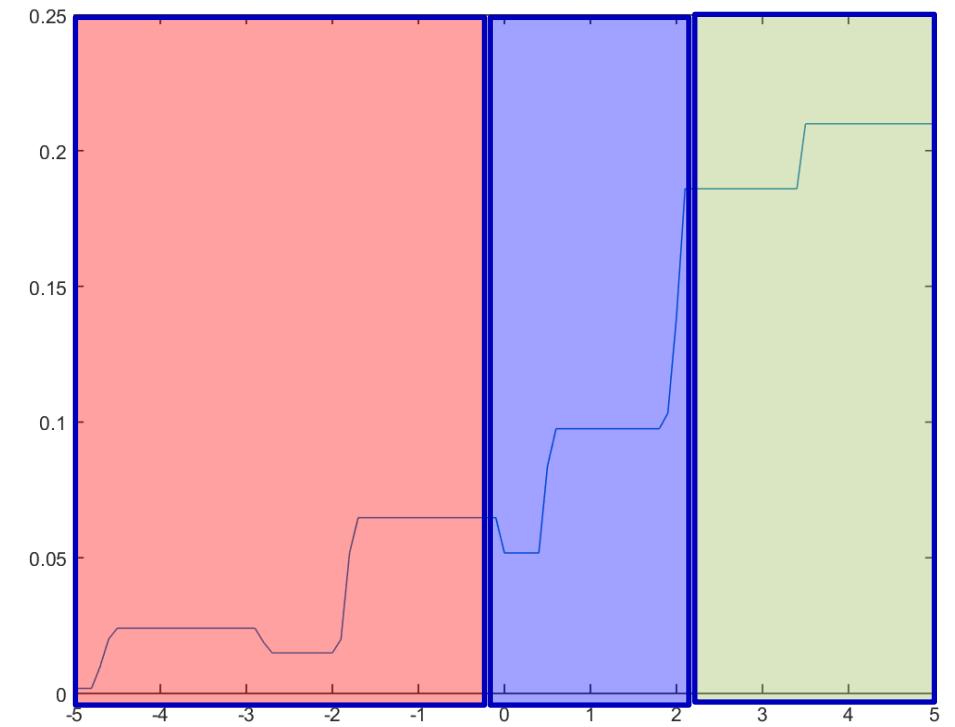
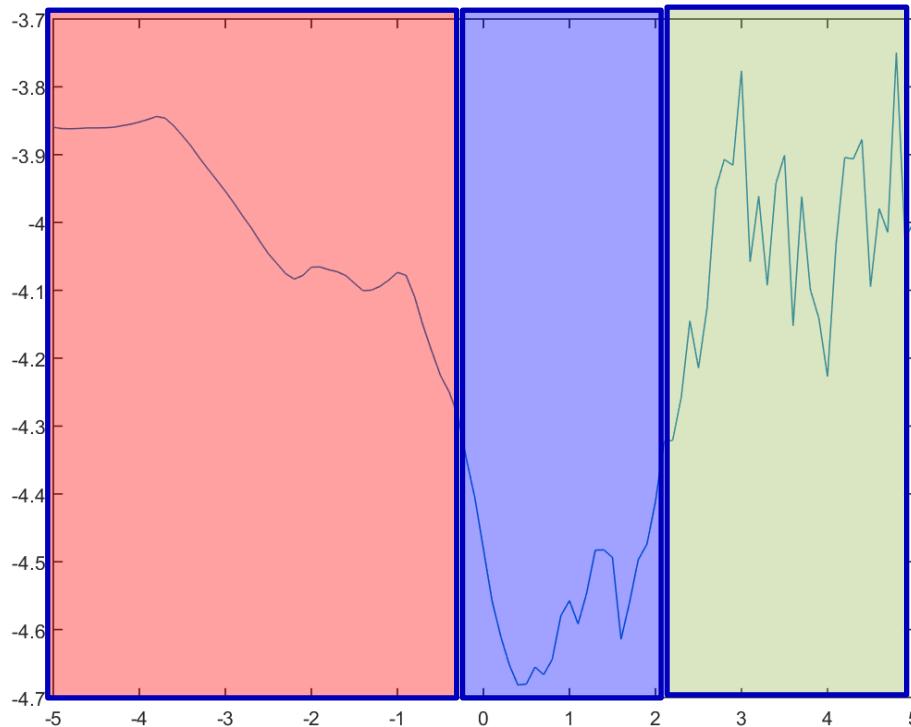
- Remember our motivation example:



- Laplacian of Gaussian filters looks useful:
 - Class 1 and 3 usually often have different values.
 - Close to zero for class 1, often far from zero for class 3.
 - But class 3 values are still sometimes close to 0.
- What if I take maximum absolute value over 16 KNNs in this signal?

Taking Maximums of Convolutions

- We often use **maximum over several convolutions** as features:
 - On right is the **maximum of Laplacian of Gaussian at ‘i’ and its 16 KNNs**.
 - We can **solve** the problem with just the 2 features below at each location.



Images and Higher-Order Convolution

- **2D convolution:**
 - Signal ‘x’ is the pixel intensities in an ‘n’ by ‘n’ image.
 - Filter ‘w’ is the pixel intensities in a ‘2m+1’ by ‘2m+1’ image.
- The **2D convolution** is given by:

$$z[i_1, i_2] = \sum_{j_1=-m}^m \sum_{j_2=-m}^m w[j_1, j_2] x[i_1 + j_1, i_2 + j_2]$$

- **3D and higher-order convolutions** are defined similarly.

$$z[i_1, i_2, i_3] = \sum_{j_1=-m}^m \sum_{j_2=-m}^m \sum_{j_3=-m}^m w[j_1, j_2, j_3] x[i_1 + j_1, i_2 + j_2, i_3 + j_3]$$

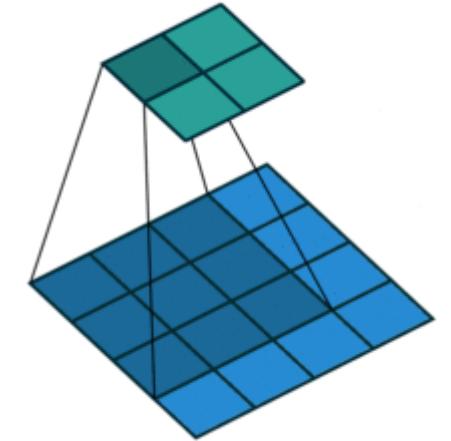
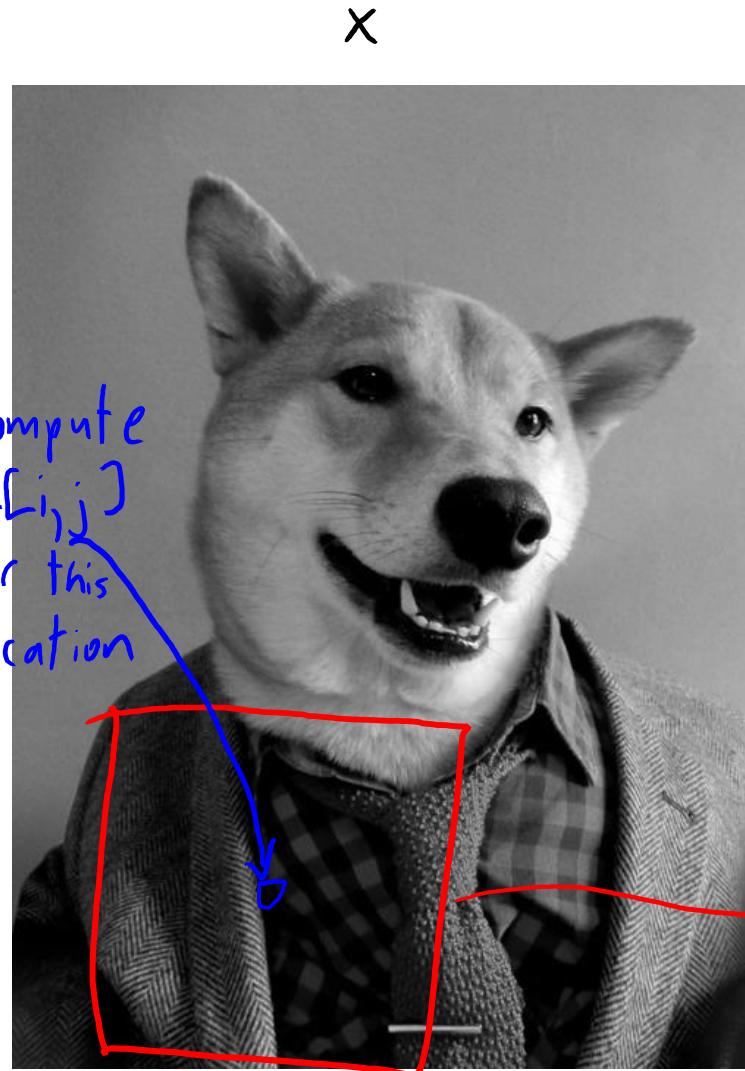


Image Convolution Examples



Identity convolution:
(zeroes with a '1' at $w_{0,0}$)

$$w \quad *$$

multiply element-wise
and add up result to get

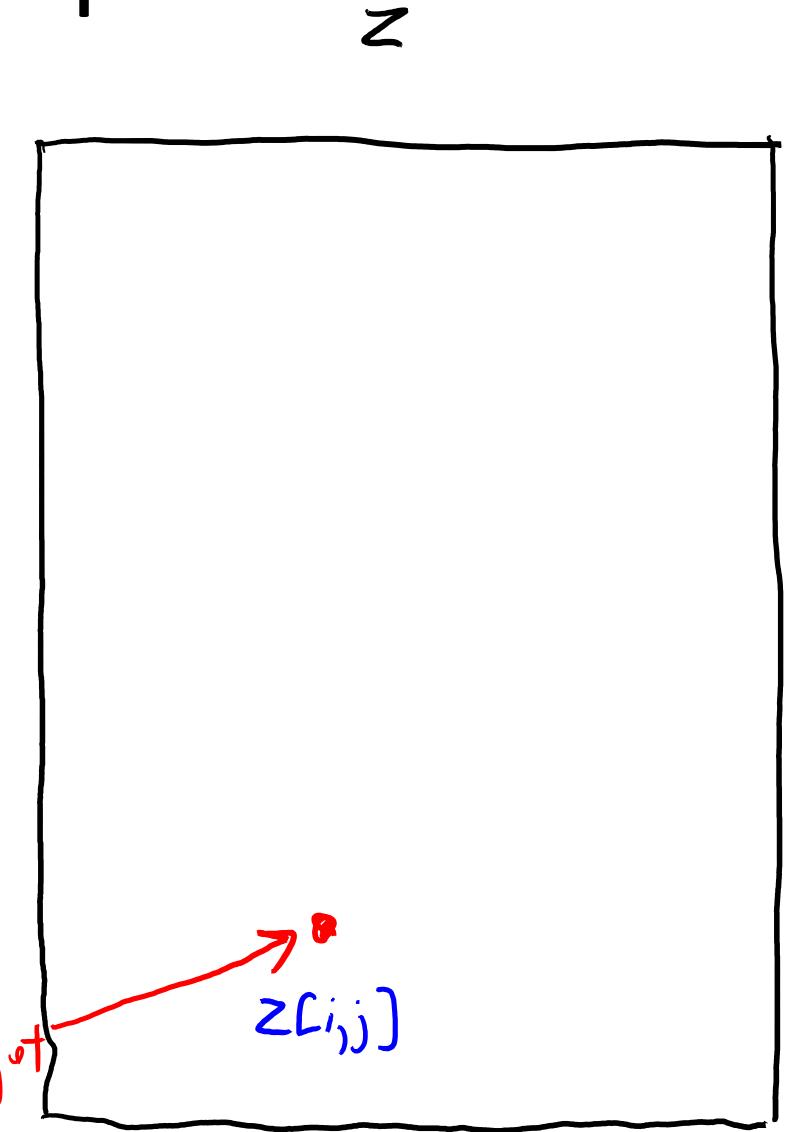
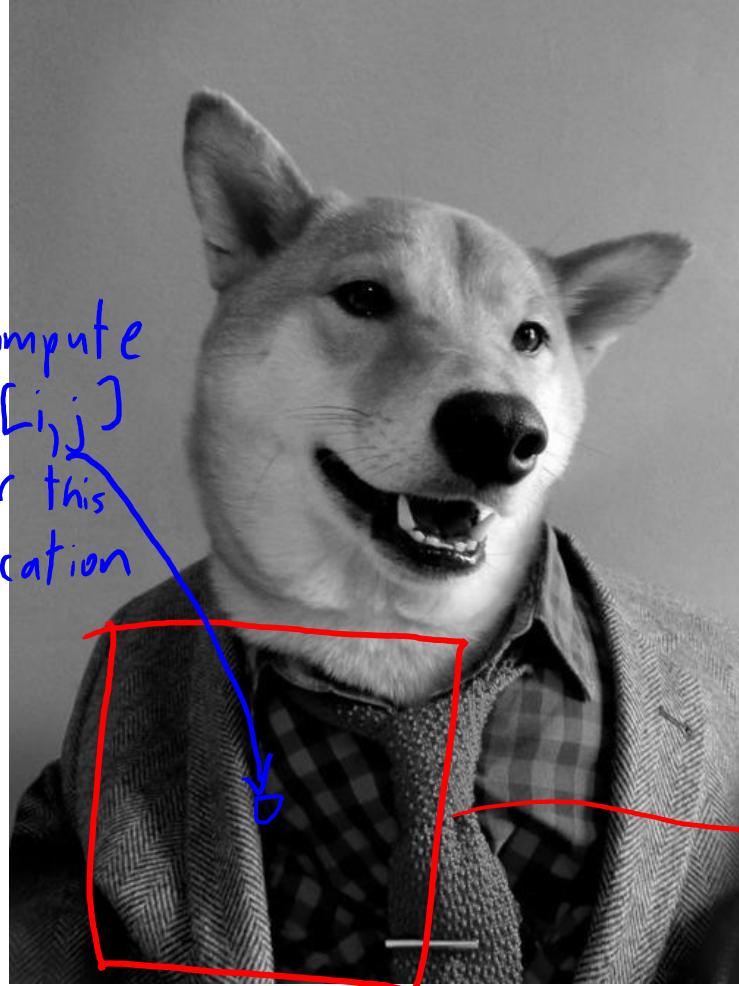


Image Convolution Examples

x



Identity convolution:
(zeroes with a '1' at $w_{0,0}$)

$$w \quad *$$



multiply element-wise
and add up result to get



Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{Input Image} \end{matrix} = \text{Output Image}$$

Boundary: "zero"



Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{Input Image} \end{matrix} \quad = \quad \text{Output Image}$$

Boundary: "replicate"



repents

Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{Boundary: "mirror"} \end{matrix} \quad \begin{matrix} \text{Translation Convolution:} \\ \text{= } \end{matrix} \quad \begin{matrix} \text{Boundary: "mirror"} \\ \text{= } \end{matrix}$$

A diagram illustrating a convolution operation. On the left, a small black square represents a kernel. A red circle highlights the top-left corner of this kernel. To its left is a symbol resembling a star with a diagonal line through it. To its right is an equals sign. To the right of the equals sign is another symbol resembling a star with a diagonal line through it. Below these symbols is the text "Boundary: 'mirror'".

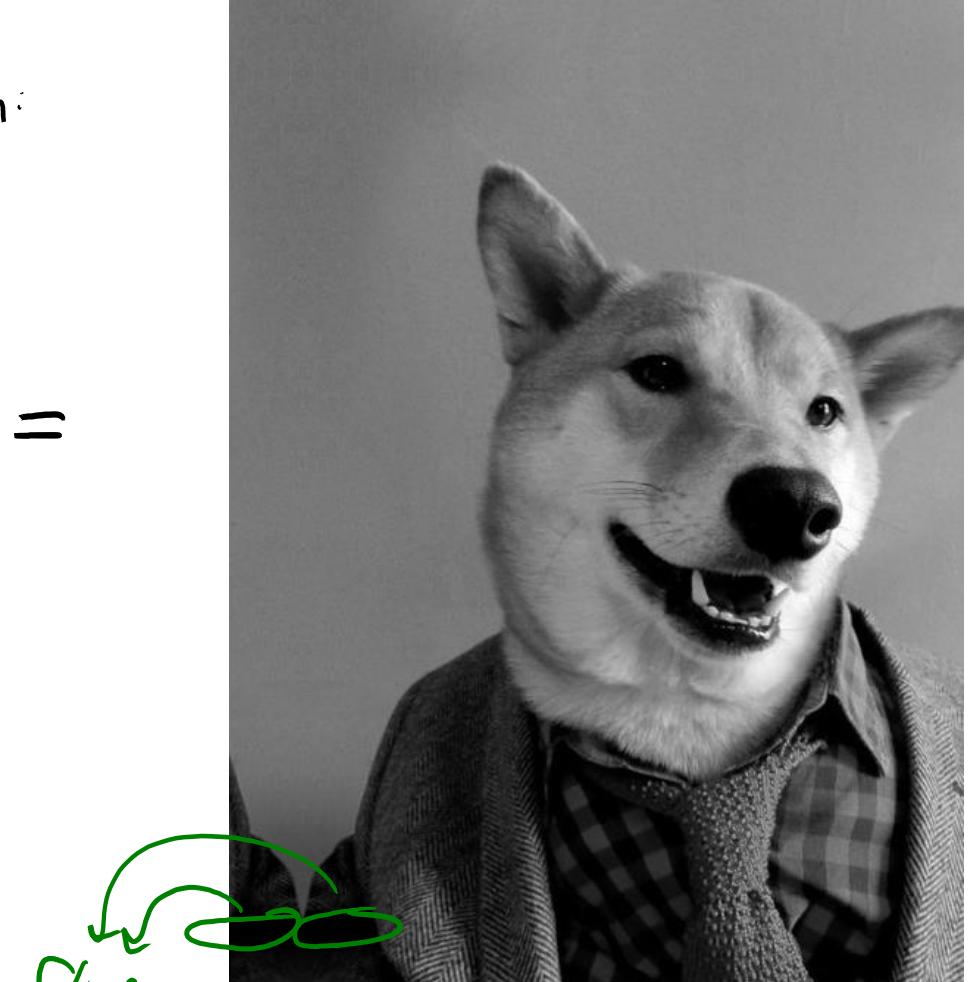


Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * & \begin{matrix} \textcircled{1} \\ \text{black square} \end{matrix} & = & \begin{matrix} \text{original image} \end{matrix} \end{matrix}$$

Boundary: "ignore"



Image Convolution Examples



Average convolution:

$$\ast \frac{1}{5} \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ 1 & 1 & 1 & \ddots & | \\ 1 & 1 & 1 & \ddots & | \\ \vdots & \vdots & \vdots & \ddots & | \\ 1 & 1 & 1 & \cdots & 1 \end{bmatrix} =$$



Convolutions as Features

- Classic vision methods uses **convolutions as features**:
 - Usually have different types/variances/orientations.
 - Can take maxes across locations/orientations/scales.

- Notable convolutions:
 - Gaussian (blurring/averaging).
 - Laplace of Gaussian (second-derivative).
 - Gabor filters (directional first- or higher-derivative).

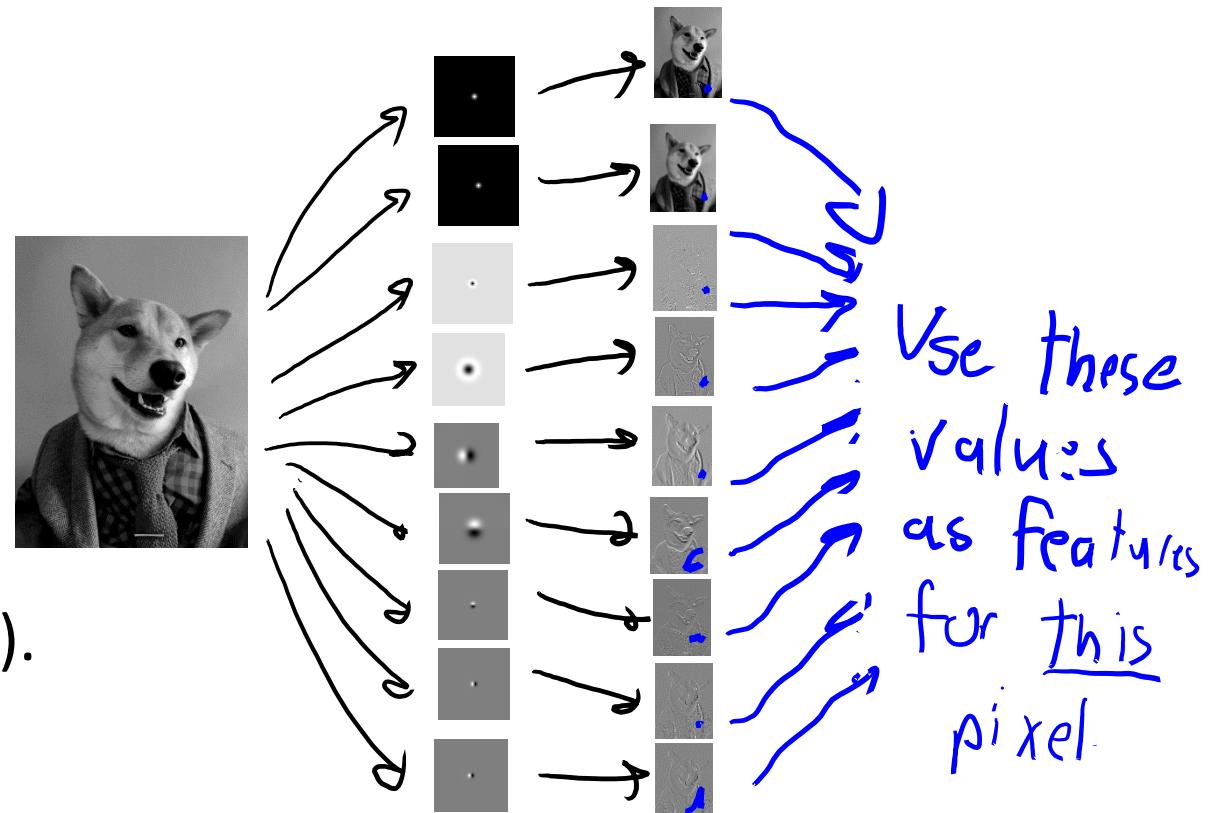
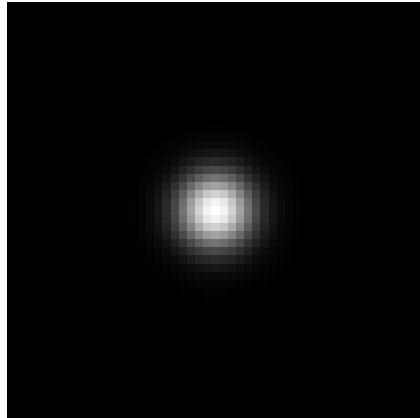


Image Convolution Examples



Gaussian Convolution:

*



=

blurs image to represent
average
(smoothing)

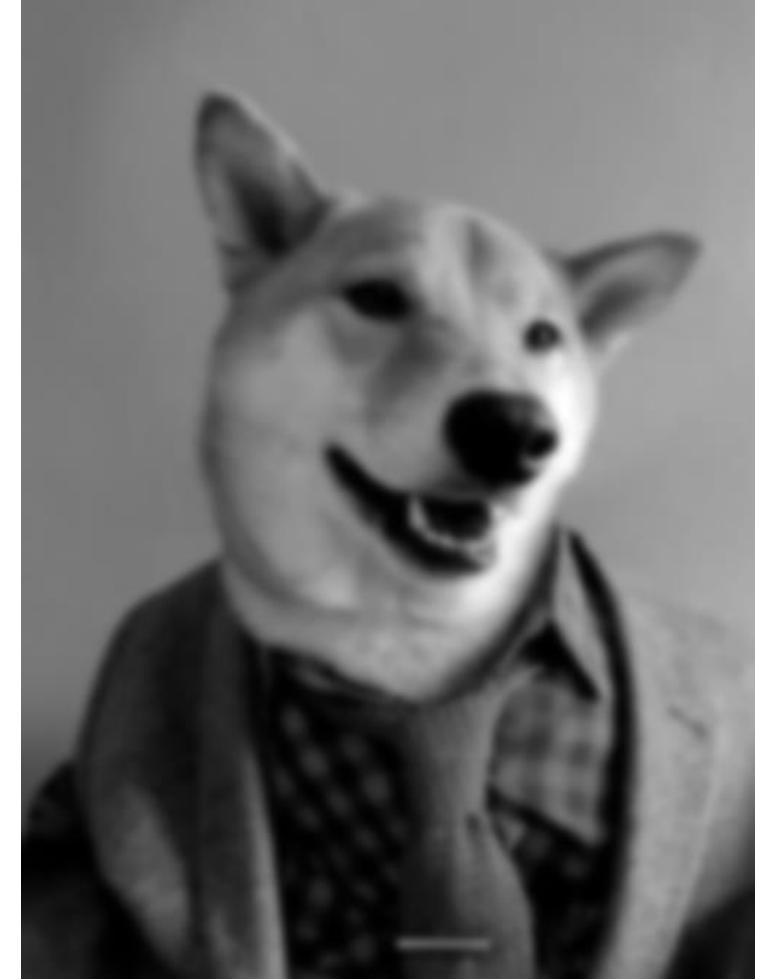
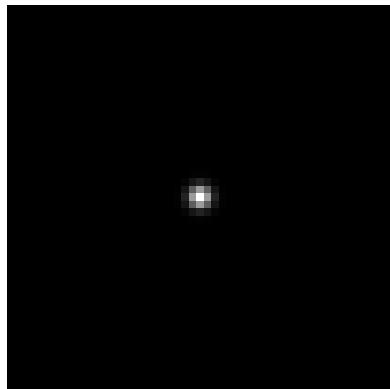


Image Convolution Examples



Gaussian Convolution:

*



=

(smaller variance)

blurs image to represent

average
(smoothing)



Image Convolution Examples



Laplacian of Gaussian

$$\begin{matrix} * & & \\ & \text{---} & \\ & \quad \square \quad & \\ & \text{---} & \end{matrix} =$$

"How much does it look
like a black dot
surrounded by white?"

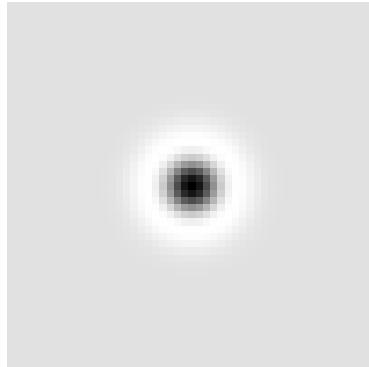


Image Convolution Examples



Laplacian of Gaussian

*



=

(larger variance)

Similar preprocessing may be
done in basal ganglia and LGN.



Image Convolution Examples



"Emboss" filter:

$$\ast \begin{bmatrix} -2 & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & 2 \end{bmatrix} =$$

Many Photoshop effects
are just convolutions.

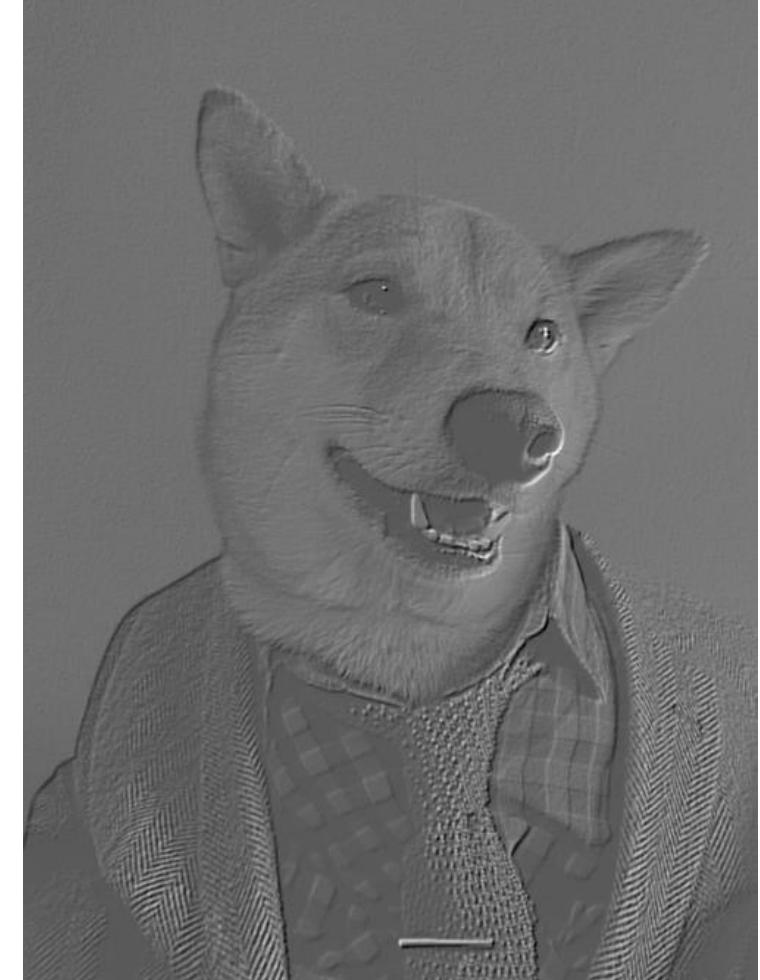


Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

$$\ast \quad \begin{matrix} \text{ } \\ \text{ } \end{matrix} = \quad \begin{matrix} \text{ } \\ \text{ } \end{matrix}$$

$$\text{II} \quad \begin{matrix} \text{ } \\ \text{ } \end{matrix} \quad \ast \quad \begin{matrix} \text{ } \\ \text{ } \end{matrix}$$

Gaussian Parallel Sine functions

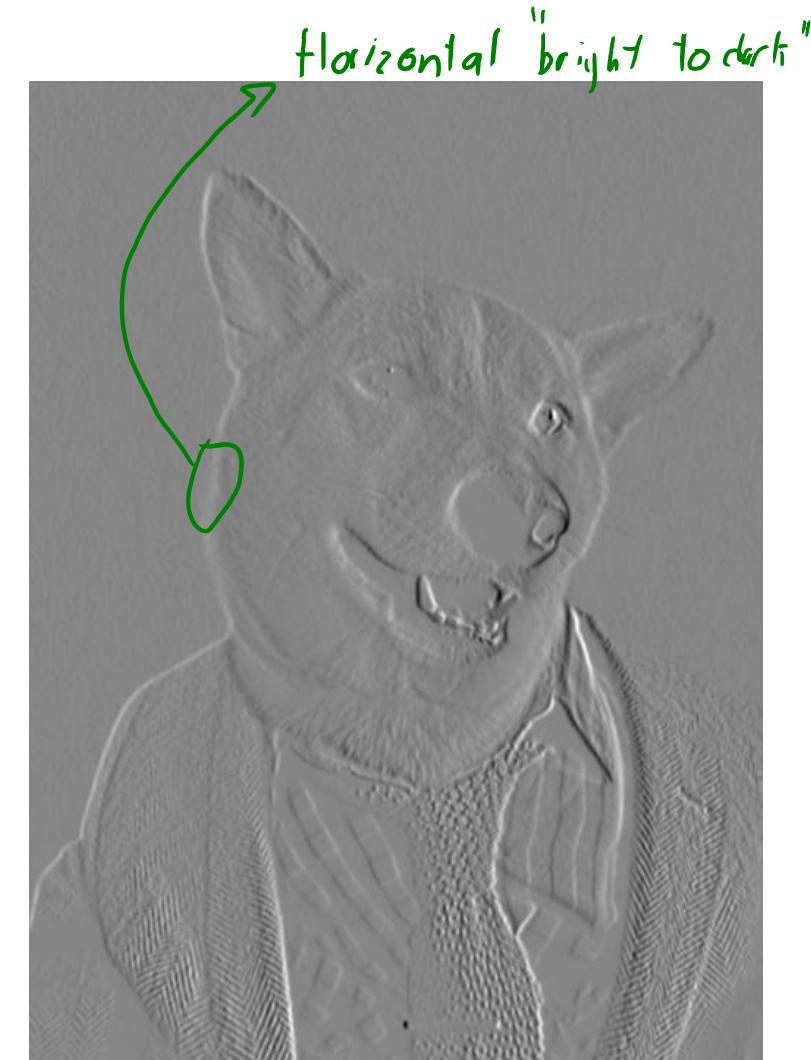
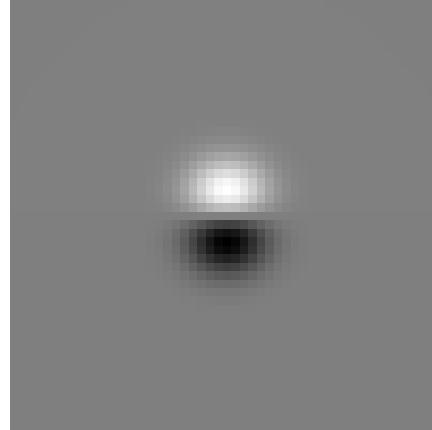


Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=



Different orientations of
the sine/cosine let us
detect changes with different

orientations.

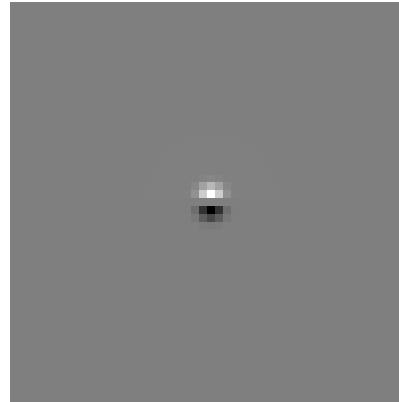
→ 2d derivatives have a direction.

Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=

(smaller variance)

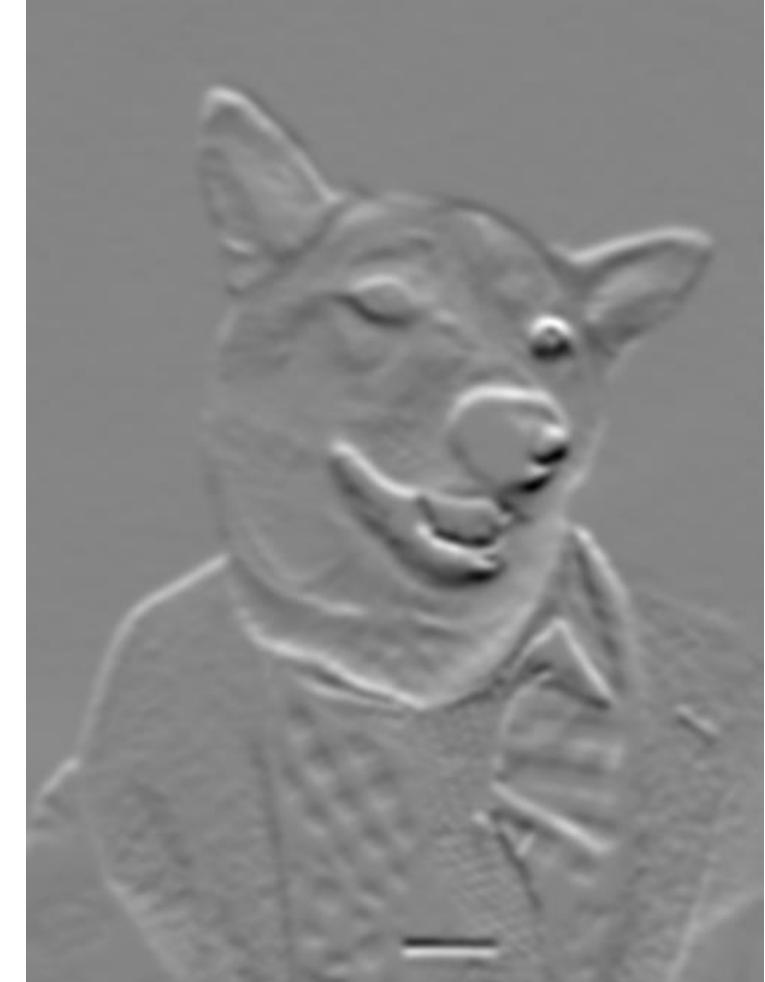
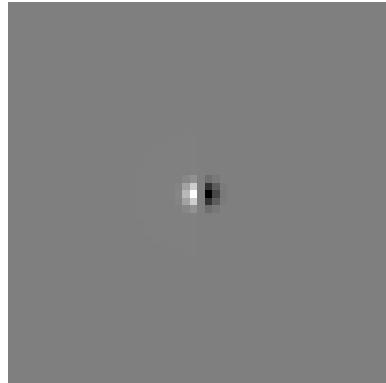


Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=

(smaller variance)

Vertical orientation

- Can obtain other orientations by rotating.

- May be similar to effect of V1 "simple cells."

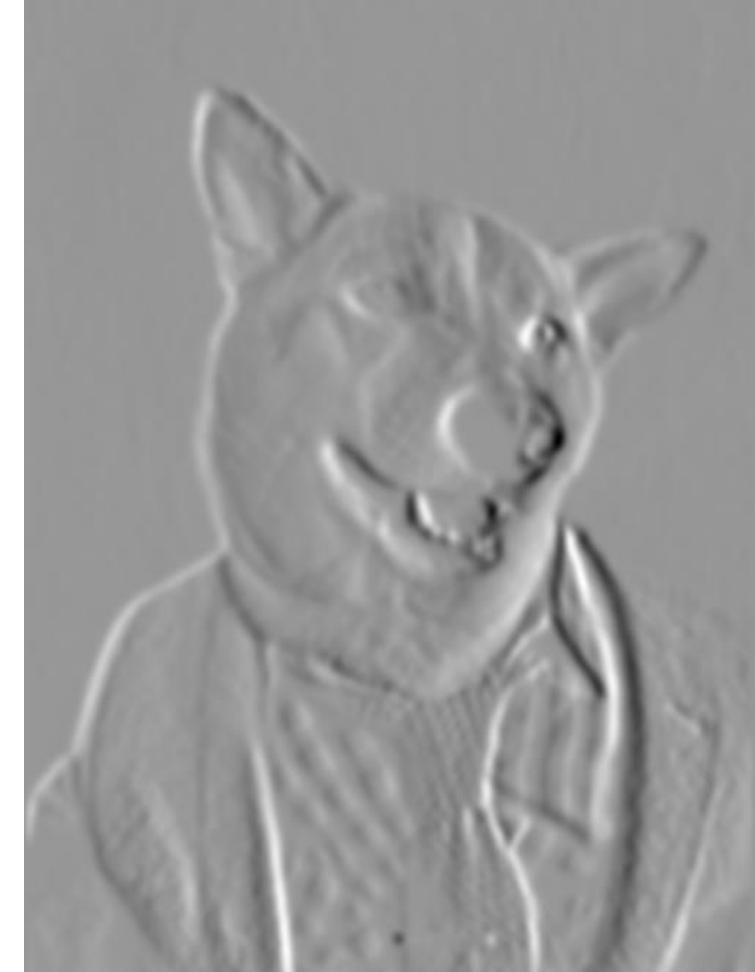
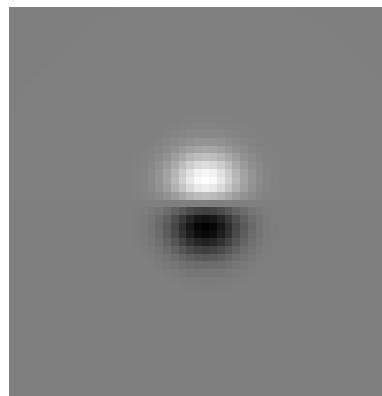
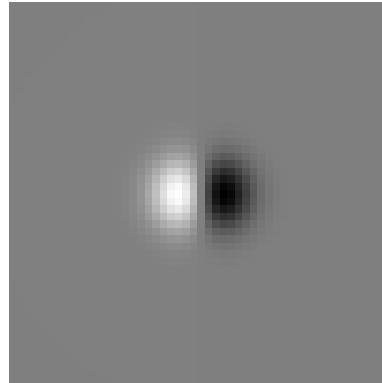


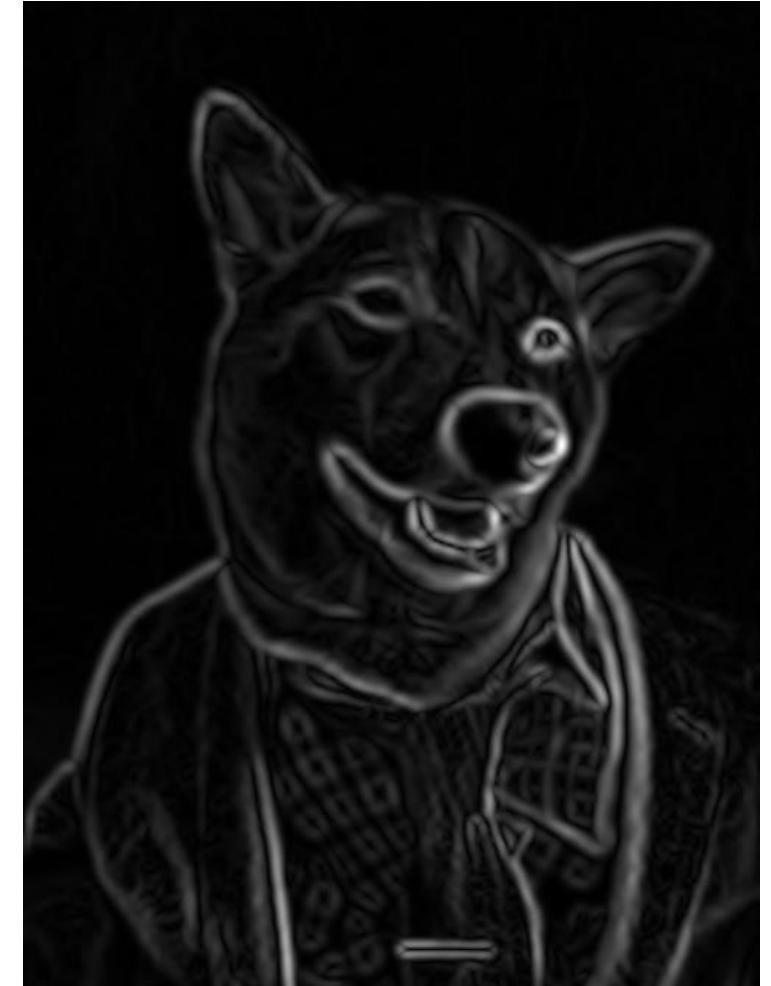
Image Convolution Examples



Max absolute value
between horizontal and
vertical Gabor:



maximum
absolute
value

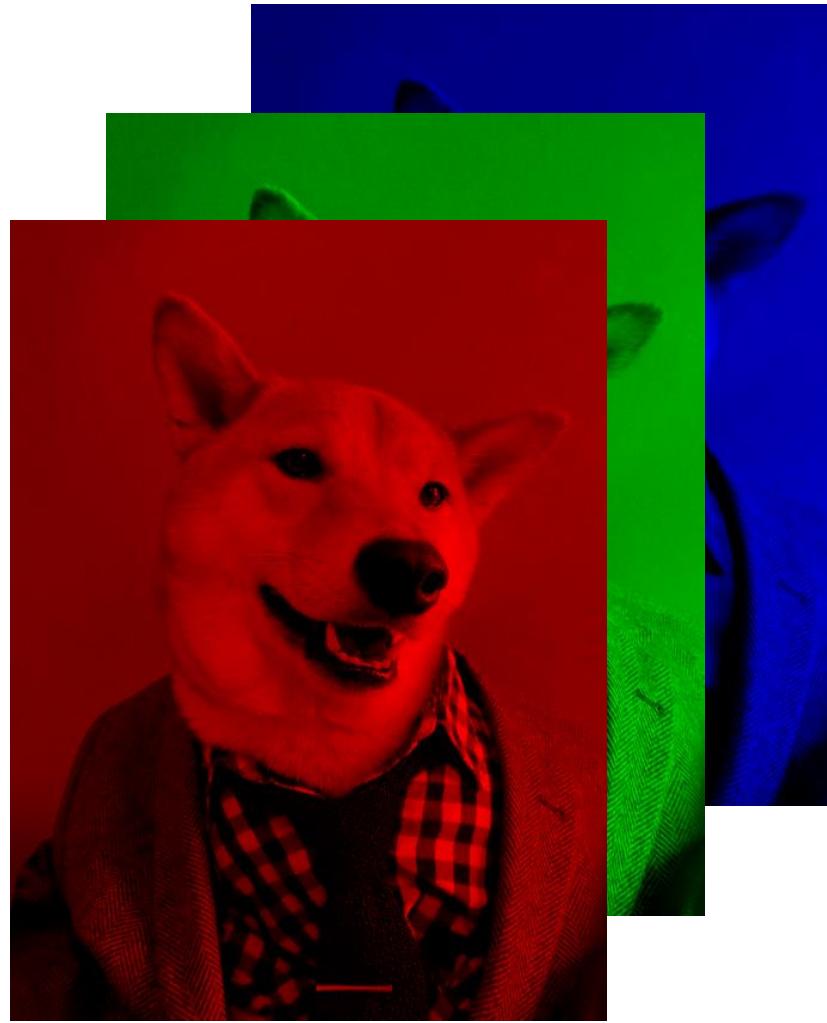


"Horizontal/vertical edge detector"

3D Convolution



Represent
as RGB

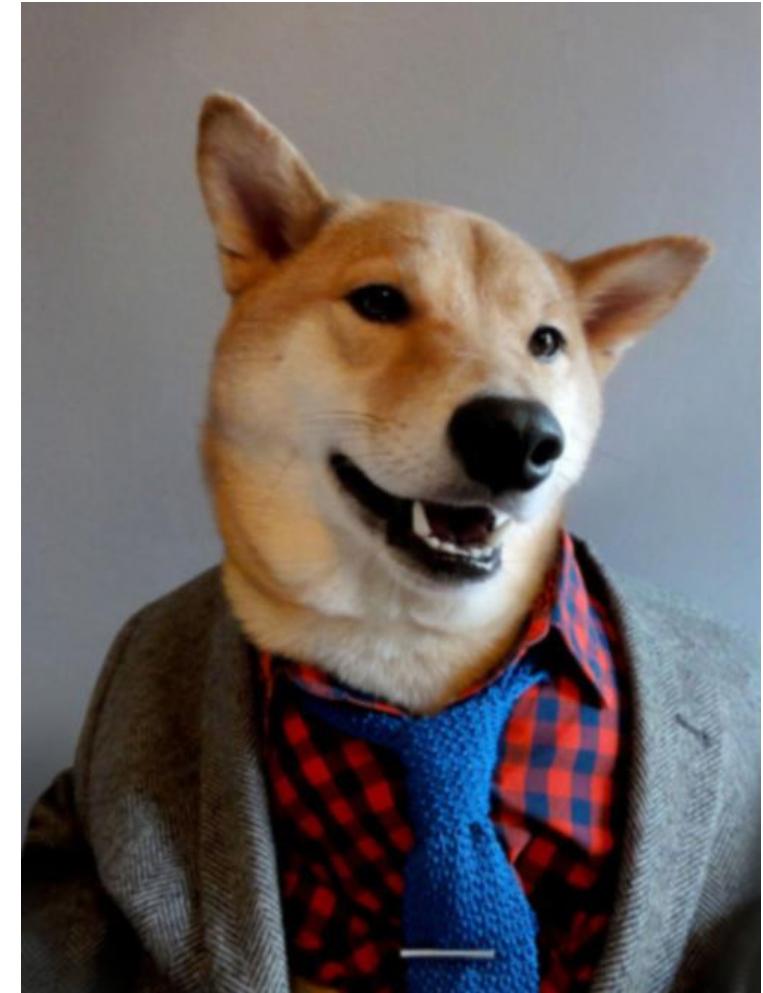


Can apply 3D
convolutions

3D Convolution



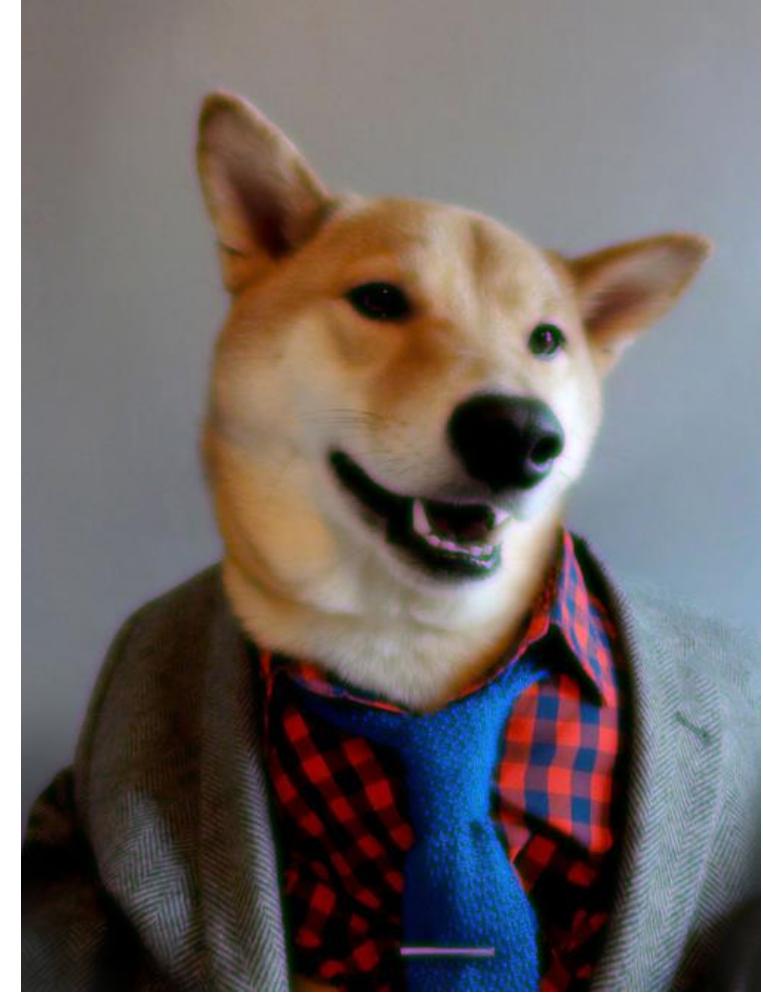
Gaussian filter



3D Convolution



Gaussian filter
(higher variance on
green channel)



3D Convolution



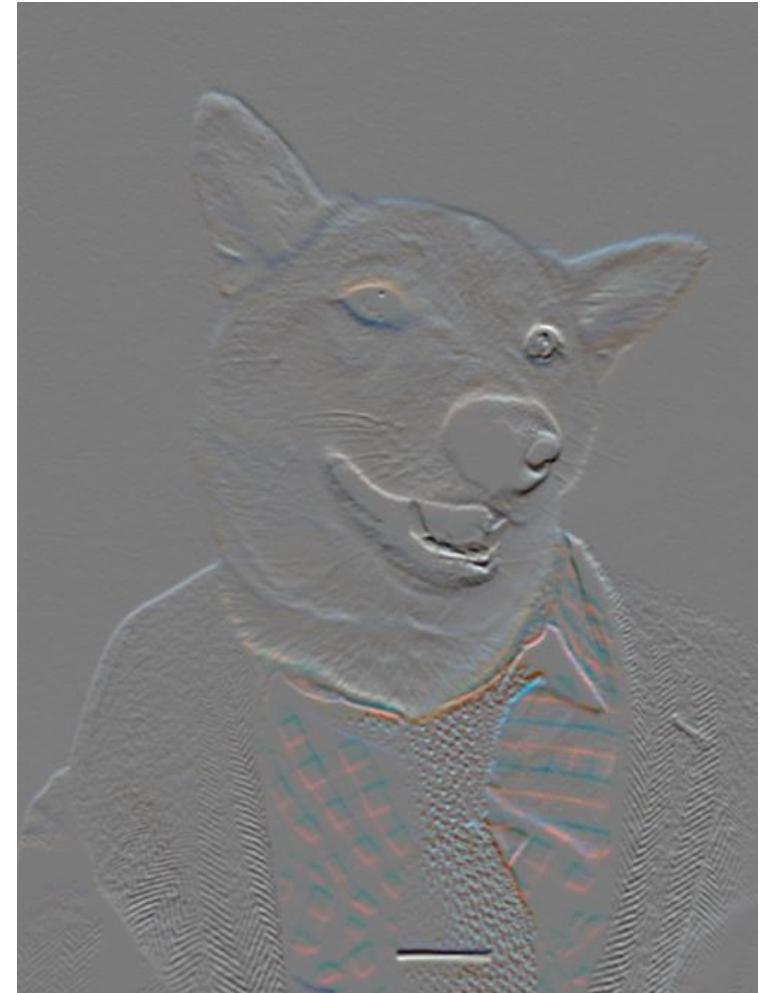
Sharpen the blue channel.



3D Convolution

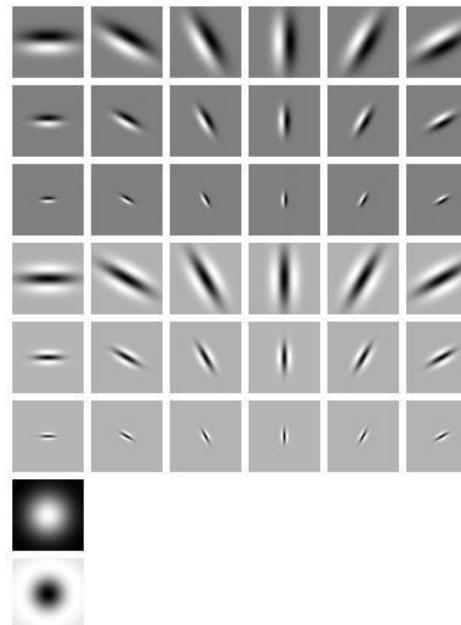


Gabor filter on
each channel.



Filter Banks

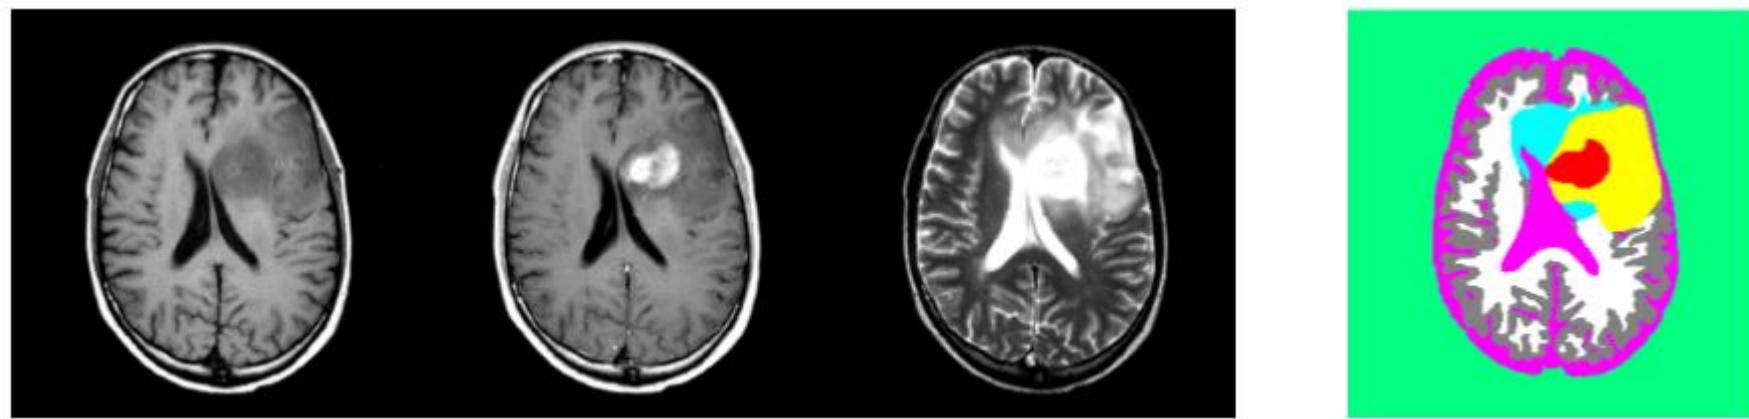
- To characterize context, we used to use **filter banks** like “MR8”:
 - 1 Gaussian filter, 1 Laplacian of Gaussian filter.
 - 6 max(Gabor) filters: 3 scales of sine/cosine (maxed over orientations).



- **Convolutional neural networks** (Part 5) are replacing filter banks.

Image Coordinates

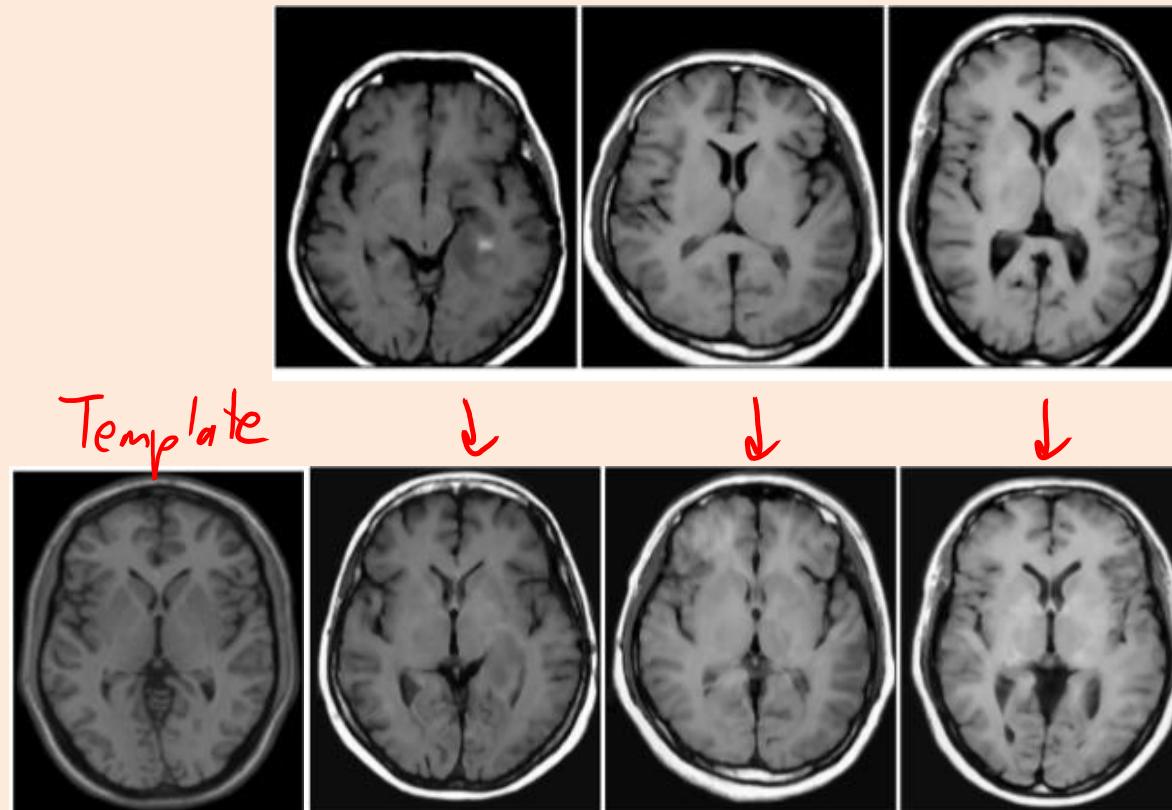
- Should we use the image coordinates?
 - E.g., the pixel is at location (124, 78) in the image.



- Considerations:
 - Is the interpretation different in different areas of the image?
 - Are you using a linear model?
 - Would “distance to center” be more logical?
 - Do you have enough data to learn about all areas of the image?

Alignment-Based Features

- The position in the image is important in brain tumour application.
 - But we didn't have much data, so **coordinates didn't make sense**.
- We aligned the images with a “template image”.



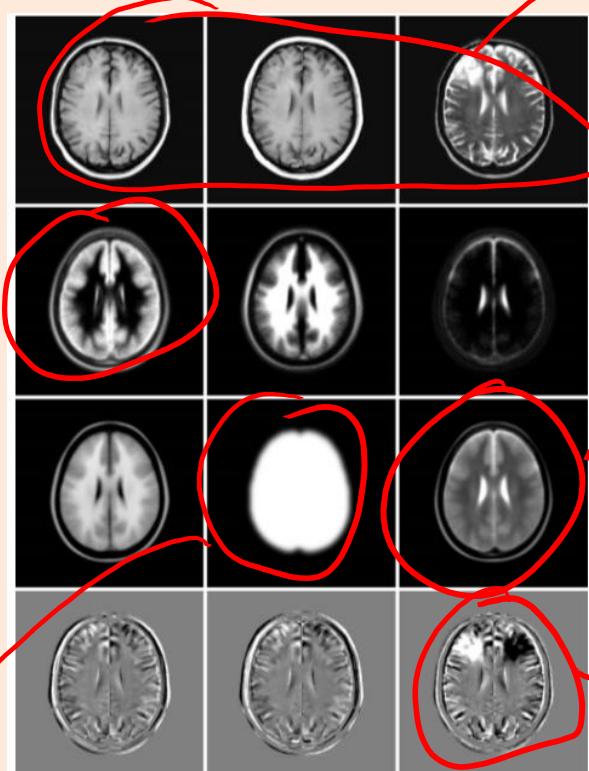
(Look different because
we're showing middle slice
and alignment is in 3D.)

Alignment-Based Features

- The position in the image is important in brain tumour application.
 - But we didn't have much data, so **coordinates didn't make sense**.
- We aligned the images with a “template image”.
 - Allowed “alignment-based” features:

Probability of
gray matter at
this pixel among
tons of people aligned
with template.

Probability of
being brain pixel.



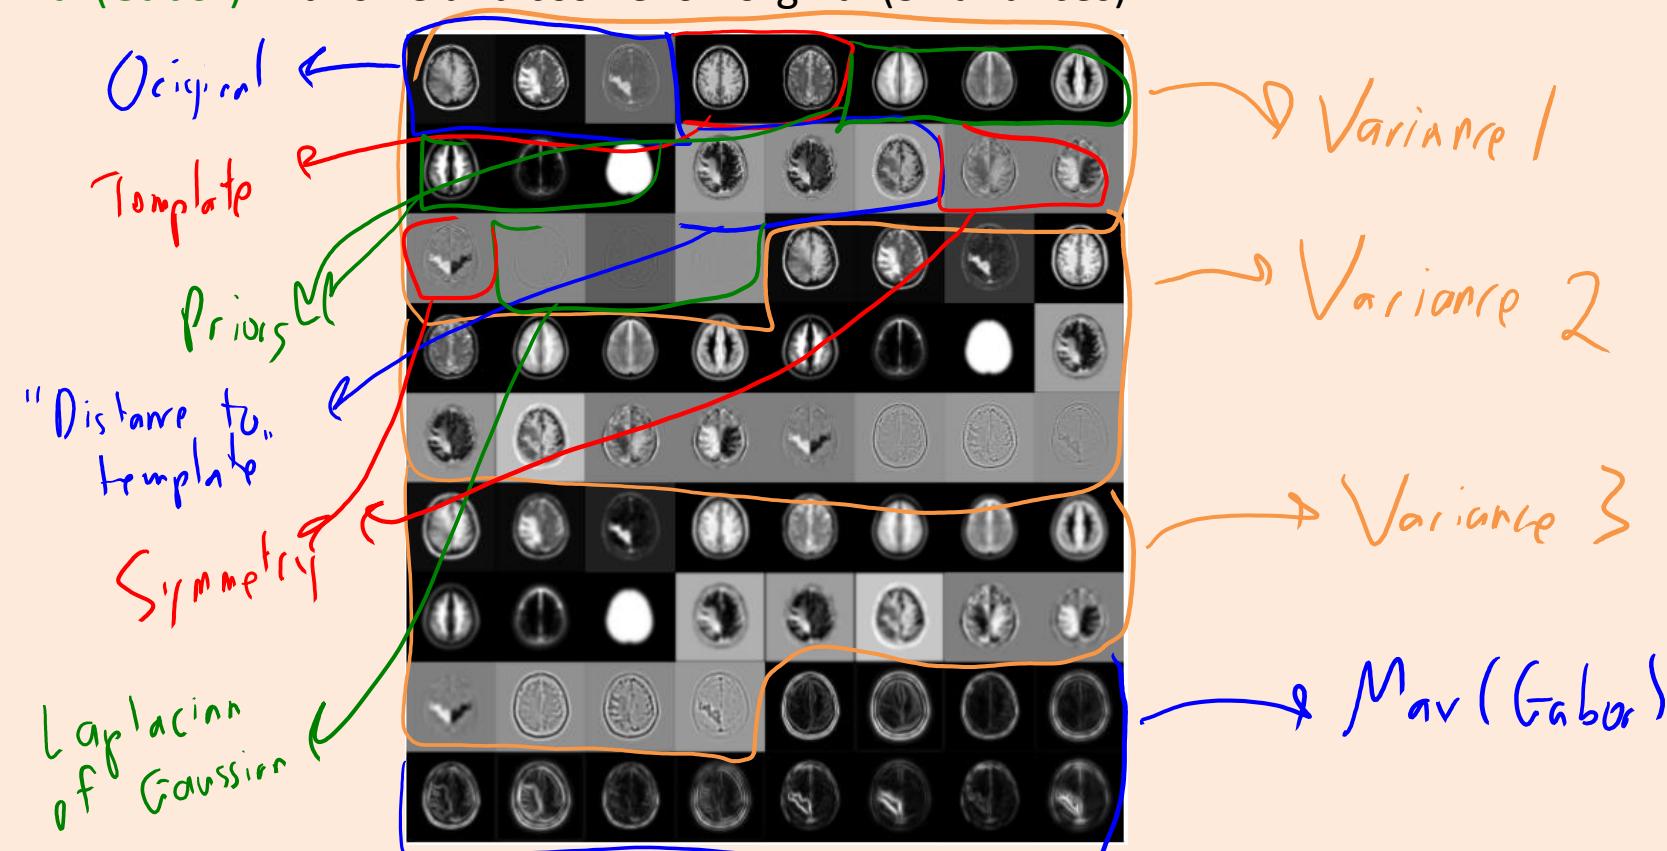
Original pixel
values

Actual pixel
value of template
image at this
location.

Left-right
symmetry difference.

Motivation: Automatic Brain Tumor Segmentation

- Final features for brain tumour segmentation:
 - Gaussian convolution of original/template/priors/symmetry, Laplacian of Gaussian on original.
 - All with 3 variances.
 - Max(Gabor) with sine and cosine on orginal (3 variances).



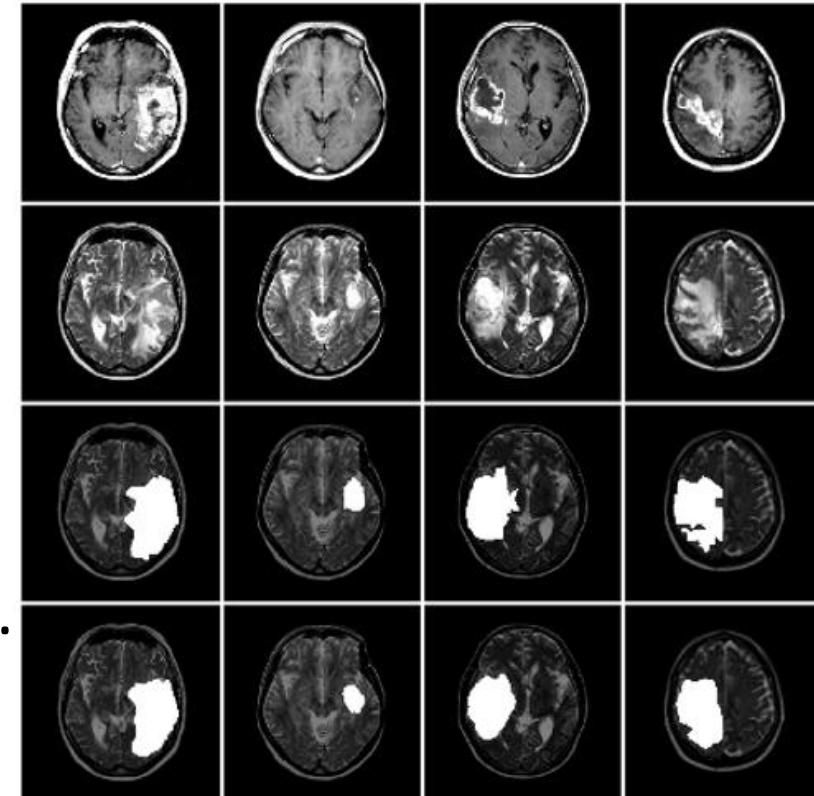
Motivation: Automatic Brain Tumour Segmentation

- Logistic regression and SVMs among best methods.

- When using these 72 features from last slide.
 - If you used all features I came up with, it overfit.

- Possible solutions to overfitting:

- Forward selection was too slow.
 - Just one image gives 8 million training examples.
 - I did manual feature selection (“guess and check”).
 - L2-regularization with all features also worked.
 - But this is slow at test time.
 - L1-regularization gives best of regularization and feature selection.



Summary

- **Convolutions** are flexible class of signal/image transformations.
 - Can approximate directional derivatives and integrals at different scales.
- **Max(convolutions)** can yield features that make classification easy.
- **Filter banks:**
 - Make features for a vision problem by takin a bunch of convolutions.
- Next time: feature engineering for image and sound data.

Global and Local Features for Domain Adaptation

- Suppose you want to solve a classification task, where you have very little labeled data from your domain.
- But you have access to a huge dataset with the same labels, from a different domain.
- Example:
 - You want to label POS tags in medical articles, and pay a few \$\$\$ to label some.
 - You have access the thousands of examples of Wall Street Journal POS labels.
- **Domain adaptation:** using data from different domain to help.

Global and Local Features for Domain Adaptation

- “Frustratingly easy domain adaptation”:
 - Use “global” features across the domains, and “local” features for each domain.
 - “Global” features let you learn patterns that occur across domains.
 - Leads to sensible predictions for new domains without any data.
 - “Local” features let you learn patterns specific to each domain.
 - For linear classifiers this would look like:

$$\hat{y}_i = \text{sign}(w_g^\top x_{ig} + w_d^\top x_{id})$$

features used across domains features/weights specific to domain

FFT implementation of convolution

- Convolutions can be implemented using fast Fourier transform:
 - Take FFT of image and filter, multiply elementwise, and take inverse FFT.
- It has faster asymptotic running time but there are some catches:
 - You need to be using periodic boundary conditions for the convolution.
 - Constants matter: it may not be faster in practice.
 - Especially compared to using GPUs to do the convolution in hardware.
 - The gains are largest for larger filters (compared to the image size).

SIFT Features

- Scale-invariant feature transform (SIFT):
 - Features used for object detection (“is particular object in the image”?)
 - Designed to detect unique visual features of objects at multiple scales.
 - Proven useful for a variety of object detection tasks.

