

CPSC 340: Machine Learning and Data Mining

Stochastic Gradient

Fall 2019

Last Time: Stochastic Gradient

- Stochastic gradient minimizes average of smooth functions:

$$f(w) = \frac{1}{n} \sum_{i=1}^n f_i(w)$$

- Function $f_i(w)$ is error on example 'i'.

- For example, with squared error we would have: $f_i(w) = \frac{1}{2}(w^T x_i - y_i)^2$.

- Iterations perform gradient descent on one random example 'i':

$$w^{t+1} = w^t - \alpha^t \nabla f_i(w^t)$$

- Cheap iterations even when 'n' is large.

- With 1 billion training examples, this iteration is 1 billion times faster.

Stochastic Gradient (SG)

- Stochastic gradient is an iterative optimization algorithm:
 - We start with some initial guess, w^0 .
 - Generate new guess by moving in the negative gradient direction:

$$w^1 = w^0 - \alpha^0 \nabla f_i(w^0)$$

- For a random training example 'i'.
- Repeat to successively refine the guess:

$$w^{t+1} = w^t - \alpha^t \nabla f_i(w^t) \quad \text{for } t = 1, 2, 3, \dots$$

- For a random training example 'i'.

Problem where we can use Stochastic Gradient

- Stochastic gradient applies when minimizing averages:

$$f(w) = \frac{1}{n} \sum_{i=1}^n (w^T x_i - y_i)^2 \quad (\text{squared error})$$

$$f(w) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) \quad (\text{logistic regression})$$

$$f(w) = \frac{1}{n} \sum_{i=1}^n \left[\log(1 + \exp(-y_i w^T x_i)) + \frac{\lambda}{2} \|w\|^2 \right] \quad (l_2\text{-regularized logistic})$$

$$f(w) = \frac{1}{n} \sum_{i=1}^n f_i(w) \quad (\text{our notation for the general case})$$

- Basically, all our regression losses except “brittle” regression.
 - Recall: multiplying by positive constant doesn’t change location of optimal ‘w’.

Why Does Stochastic Gradient Work / Not Work?


- Main problem with stochastic gradient:
 - Gradient of random example might **point in the wrong direction**.
- Does this have any hope of working?
 - The average of the random gradients is the full gradient.

Mean over $\nabla f_i(w^t)$ is $\frac{1}{n} \sum_{i=1}^n \nabla f_i(w^t)$ which is $\nabla f(w^t)$


- The algorithm is going in the **right direction on average**.

Gradient Descent vs. Stochastic Gradient (SG)

- Gradient descent:




- Stochastic gradient:



Gradient Descent in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$



Stochastic Gradient in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$


$$f_1(w) = (w^T x_1 - y_1)^2$$

$$f_2(w) = (w^T x_2 - y_2)^2$$

$$f_3(w) = (w^T x_3 - y_3)^2$$

$$f_4(w) = (w^T x_4 - y_4)^2$$

$$f_5(w) = (w^T x_5 - y_5)^2$$



Stochastic Gradient in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$



$$f_1(w) = (w^T x_1 - y_1)^2$$

$$f_2(w) = (w^T x_2 - y_2)^2$$

$$f_3(w) = (w^T x_3 - y_3)^2$$

$$f_4(w) = (w^T x_4 - y_4)^2$$

$$f_5(w) = (w^T x_5 - y_5)^2$$



Stochastic
gradient
minimizes
average
value.

Stochastic Gradient in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$



$$f_1(w) = (w^T x_1 - y_1)^2$$

$$f_2(w) = (w^T x_2 - y_2)^2$$

$$f_3(w) = (w^T x_3 - y_3)^2$$

$$f_4(w) = (w^T x_4 - y_4)^2$$

$$f_5(w) = (w^T x_5 - y_5)^2$$



Stochastic
gradient
minimizes
average
value.

Stochastic Gradient in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$



$$f_1(w) = (w^T x_1 - y_1)^2$$

$$f_2(w) = (w^T x_2 - y_2)^2$$

$$f_3(w) = (w^T x_3 - y_3)^2$$

$$f_4(w) = (w^T x_4 - y_4)^2$$

$$f_5(w) = (w^T x_5 - y_5)^2$$



Stochastic
gradient
minimizes
average
value.

Stochastic Gradient in Action

$$f(w) = \frac{1}{5} \sum_{i=1}^5 (w^T x_i - y_i)^2$$



$$f_1(w) = (w^T x_1 - y_1)^2$$

$$f_2(w) = (w^T x_2 - y_2)^2$$


$$f_3(w) = (w^T x_3 - y_3)^2$$

$$f_4(w) = (w^T x_4 - y_4)^2$$

$$f_5(w) = (w^T x_5 - y_5)^2$$



Effect of 'w' Location on Progress



- We'll still make good progress if most gradients point in right direction.

Variance of the Random Gradients

- The “confusion” is captured by a kind of **variance of the gradients**:


$$\frac{1}{n} \sum_{i=1}^n \|\nabla f_i(w^t) - \nabla f(w^t)\|^2$$

- If the variance is 0, **every step goes in the right direction**.
 - We’re outside of the region of confusion.
- If the variance is small, **most steps point in the direction**.
 - We’re just inside region of confusion.
- If the variance is large, **many steps will point in the wrong direction**.
 - Middle of region of confusion, where w^* lives.


Effect of the Step-Size

- We can reduce the effect of the variance with the step size.
 - Variance slows progress by amount proportional to square of step-size.
 - So as the step size gets smaller, the variance has less of an effect.
- For a fixed step-size, SG makes progress until variance is too big.
- This leads to two “phases” when we use a constant step-size:
 1. Rapid progress when we are far from the solution.
 2. Erratic behaviour confined to a “ball” around solution.
(Radius of ball is proportional to the step-size.)


Stochastic Gradient with Constant Step Size




Stochastic Gradient with Constant Step Size



Stochastic Gradient with Constant Step Size



Stochastic Gradient with Constant Step Size



Stochastic Gradient with Decreasing Step Sizes

- To get convergence, we need a **decreasing step size**.
 - Shrinks size of ball to zero so we converge to w^* .
- But it **can't shrink too quickly**:
 - Otherwise, we don't move fast enough to reach the ball.
- Stochastic gradient **converges to a stationary point** if:
 - Ratio of sum of squared step-sizes over sum of step-sizes converges to 0.

$$\begin{array}{l} \text{"how much noise affects you"} \rightarrow \frac{\sum_{t=1}^{\infty} (\alpha^t)^2}{\sum_{t=1}^{\infty} \alpha^t} = 0 \\ \text{"how far you can get"} \rightarrow \end{array}$$

- This choice **also works for non-smooth** functions like SVMs.
 - Function must be continuous and not "too crazy" (we're still figuring it out for non-convex).

Stochastic Gradient with Decreasing Step Sizes

- For convergence step-sizes need to satisfy: $\sum_{t=1}^{\infty} (\alpha^t)^2 / \sum_{t=1}^{\infty} \alpha^t = 0$
- Classic solution is to use a step-size sequence like $\alpha^t = O(1/t)$.

$$\sum_{t=1}^{\infty} \alpha^t = \sum_{t=1}^{\infty} \frac{1}{t} = \infty$$

"we can get everywhere"

$$\sum_{t=1}^{\infty} (\alpha^t)^2 = \sum_{t=1}^{\infty} \frac{1}{t^2} < \infty$$

"effect of variance goes to zero"

- E.g., $\alpha^t = .001/t$.
- Unfortunately, this often **works badly in practice**:
 - Steps get really small really fast.
 - Some authors add extra parameters like $\alpha^t = \gamma/(t + \Delta)$, which helps a bit.
 - One of the only cases where this works well: binary SVMs with $\alpha^t = 1/\lambda t$.

Stochastic Gradient with Decreasing Step Sizes

- How do we pick step-sizes satisfying

$$\sum_{t=1}^{\infty} (\alpha^t)^2 / \sum_{t=1}^{\infty} \alpha^t = 0$$

- Better solution is to use a step-size sequence like $\alpha^t = O(1/\sqrt{t})$.

$$\sum_{t=1}^{\infty} \alpha^t = \sum_{t=1}^{\infty} \frac{1}{\sqrt{t}} = O(\sqrt{t}) \quad \sum_{t=1}^{\infty} (\alpha^t)^2 = \sum_{t=1}^{\infty} \frac{1}{t} = O(\log t)$$

- E.g., use $\alpha^t = .001/\sqrt{t}$
- Both sequences diverge, but **denominator diverges faster**.
- This approach (roughly) **optimizes rate** that it goes to zero.
 - Better worst-case theoretical properties (and more robust to step-size).
 - Often better in practice too.

Stochastic Gradient with Constant Step Sizes?

- Alternately, could we just use a **constant step-size**.
 - E.g., use $\alpha^t = .001$ for all 't'.
- This **will not converge** to a stationary point in general.
 - However, do we need it to converge?
- What if you **only care about the first 2-3 digits of the test error**?
 - Who cares if you aren't able to get 10 digits of optimization accuracy?
- **There is a step-size small enough to achieve any fixed accuracy.**
 - Just need radius of “ball” to be small enough.

Mini-batches: Using more than 1 example

- Does it make sense to use **more than 1 random example**?
 - Yes, you can use a “mini-batch” B^t of examples.

$$w^{t+1} = w^t - \alpha^t \frac{1}{|B^t|} \sum_{i \in B^t} \nabla f_i(w^t)$$

Random “batch” of examples.

- Radius of ball **is inversely proportional to the mini-batch size**.
 - If you double the batch size, you half the radius of the ball.
 - Big gains for going from 1 to 2, less big gains from going from 100 to 101.
 - You **can use a bigger step size as the batch size increases**.
 - Gets you to the ball faster (though diverges if step-size is too big).
- Useful for vectorizing/parallelizing code.
 - Evaluate **one gradient on each core**.

Polyak-Ruppert Iterate Averaging

- Another practical/theoretical trick is **averaging of the iterations**.
 1. Run the stochastic gradient algorithm with $\alpha^t = O(1/\sqrt{t})$ or α^t constant.
 2. Take some **weighted average** of the w^t values.

$$\bar{w}^t = \sum_{k=1}^t v^k w^k$$


here, v^k is a scalar
"weight" of iteration 'k'

Uniform average:
 $\bar{w}^t = \frac{1}{n} \sum_{k=1}^t w^k$

$v^k = \frac{1}{n}$

- Average does not affect the algorithm, it's just "watching".
- Surprising result shown by Polyak and by Ruppert in the 1980s:
 - Asymptotically converges as fast **as stochastic Newton's method**.


Stochastic Gradient with Averaging




A Practical Strategy for Deciding When to Stop

- In gradient descent, we can stop when gradient is close to zero.
- In stochastic gradient:
 - Individual gradients don't necessarily go to zero.
 - We **can't see full gradient**, so we **don't know when to stop**.
- Practical trick:
 - Every 'k' iterations (for some large 'k'), **measure validation set error**.
 - **Stop if the validation set error "isn't improving"**.
 - We don't check the gradient, since it takes a lot longer for the gradient to get small.
 - This "early stopping" can also **reduce overfitting**.


Gradient Descent vs. Stochastic Gradient



- 2012: methods with **cost of stochastic gradient, progress of full gradient**.
 - Key idea: if 'n' is finite, you **can use a memory** instead of having α_t go to zero.
 - First was stochastic average gradient (SAG), “low-memory” version is SVRG.




This graph shows how algorithms have become fast and more efficient over time. The horizontal axis represents time and the vertical axis represents error. Older algorithms (yellow) were very slow but had very little error. Faster algorithms were created by only analyzing some of the data (orange). The method was faster but had an accuracy limit. Schmidt's algorithm is faster and has no accuracy limit. *Aiken Lao / The Ubyyssey*





Machine Learning with “ $n = \infty$ ”

- Here are some scenarios where you effectively have “ $n = \infty$ ”:
 - A dataset that is so **large we cannot even go through it once** (Gmail).
 - A function **you want to minimize that you can't measure without noise**.
 - You want to encourage invariance with a **continuous set of transformation**:
 - You consider infinite number of translations/rotations instead of a fixed number.



- Learning from simulators with random numbers (physics/chem/bio):



Stochastic Gradient with Infinite Data

- Amazing property of stochastic gradient:
 - The classic convergence analysis does not rely on ‘ n ’ being finite.
- Previous slide gives examples with an infinite sequence of IID samples.
- Approach 1 (exact optimization on finite ‘ n ’):
 - Grab ‘ n ’ data points, for some really large ‘ n ’.
 - Fit a regularized model on this fixed dataset (“empirical risk minimization”).
- Approach 2 (stochastic gradient for ‘ n ’ iterations):
 - Run stochastic gradient iteration for ‘ n ’ iterations.
 - Each iteration considers a new example, never re-visiting any example.

Stochastic Gradient with Infinite Data

- Approach 2 **only looks at a data point once**:
 - Each example is an unbiased approximation of test data.
- So Approach 2 is doing **stochastic gradient on test error**:
 - It cannot overfit.
- Up to a constant, **Approach 1 and 2 have same test error bound**.
 - This is sometimes used to justify SG as the “ultimate” learning algorithm.
 - “Optimal test error by computing gradient of each example once!”
 - In practice, Approach 1 usually gives lower test error.
 - The constant factor matters!

Summary

- **Step-size in stochastic gradient** is a huge pain:
 - Needs to go to zero to get convergence, but classic $O(1/t)$ steps are bad.
 - $O(1/\sqrt{t})$ works better, but still pretty slow.
 - Constant step-size is fast, but only up to a certain point.
- **SGD practical issues**: mini-batching, averaging, termination.
- **SAG** and other methods fix SG convergence for finite datasets.
- **Infinite datasets** can be used with SG and do not overfit.
- Next time:
 - An algorithm that has been dominating Kaggle ML competitions.

A Practical Strategy For Choosing the Step-Size

- All these step-sizes have a constant factor in the “O” notation.
 - E.g., $\alpha^t = \frac{\gamma}{\sqrt{t}}$ ← How do we choose this constant?
- We **don't know how to set step size as we go** in the stochastic case.
 - And choosing wrong γ can destroy performance.
- Common practical trick:
 - Take a **small amount of data** (maybe 5% of the original data).
 - Do a **binary search for γ** that most improves objective on this subset.