

CPSC 340: Machine Learning and Data Mining

Multi-Dimensional Scaling

Fall 2019

Last Time: Multi-Dimensional Scaling

- **PCA** for visualization:
 - We're **using PCA** to get the location of the z_i values.
 - We then plot the z_i values as locations in a scatterplot.
- **Multi-dimensional scaling (MDS)** is a crazy idea:
 - Let's directly optimize the pixel locations of the z_i values.
 - "Gradient descent on the points in a scatterplot".
 - Needs a "cost" function saying how "good" the z_i locations are.
 - Traditional **MDS cost function**:

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (\|z_i - z_j\| - \|x_i - x_j\|)^2$$

Try to make scatterplot distances match high-dimensional distance

Distance between points in original 'd' dimensions

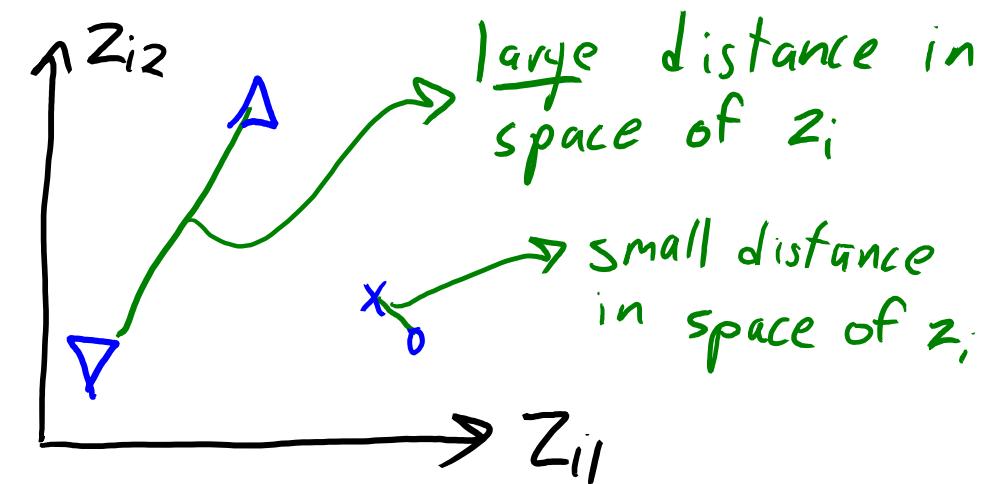
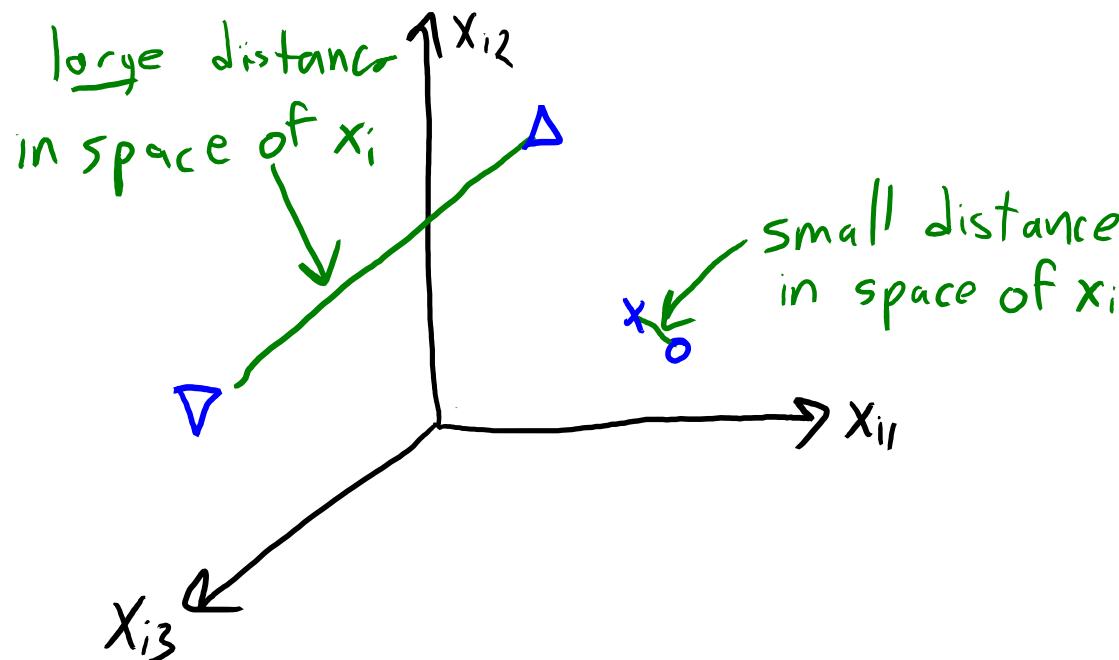
sum over pairs of examples

distance in scatterplot

Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (||z_i - z_j|| - ||x_i - x_j||)^2$$

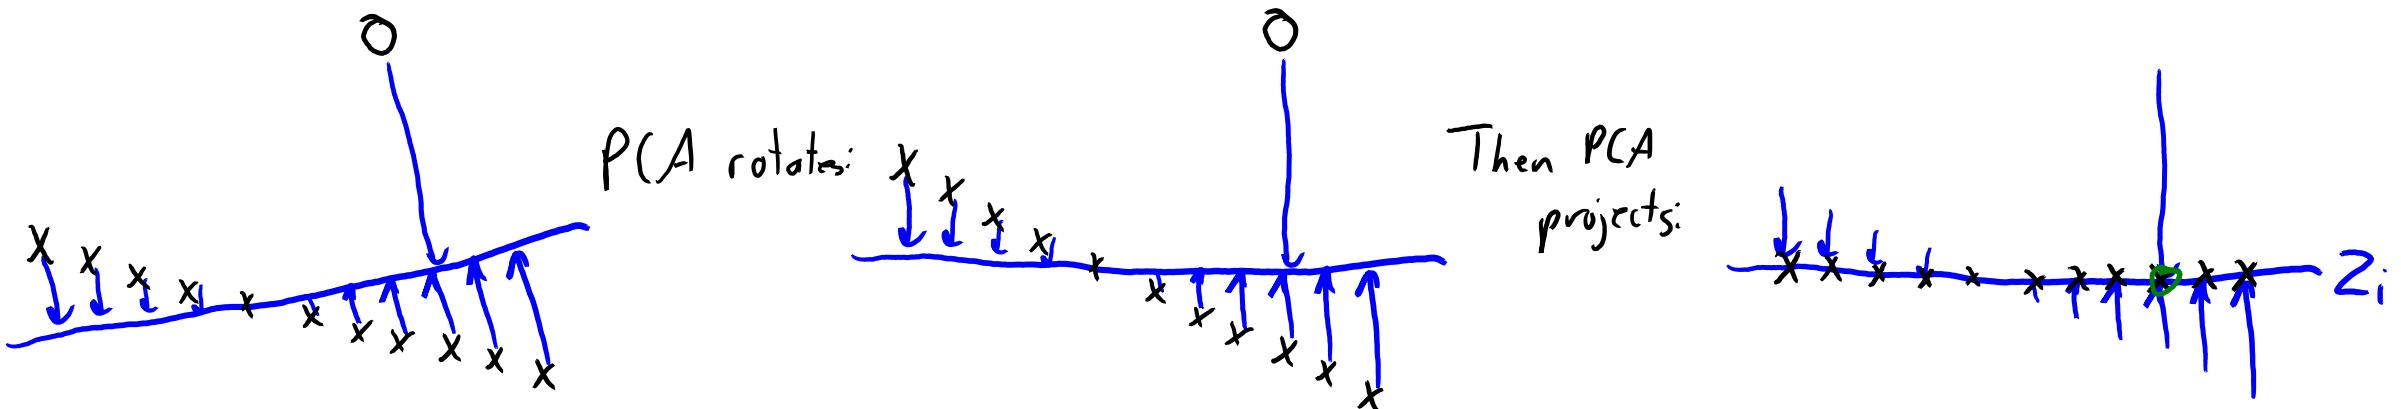


Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (||z_i - z_j|| - ||x_i - x_j||)^2$$

- Non-parametric dimensionality reduction and visualization:
 - No ‘W’: just trying to make z_i preserve high-dimensional distances between x_i .

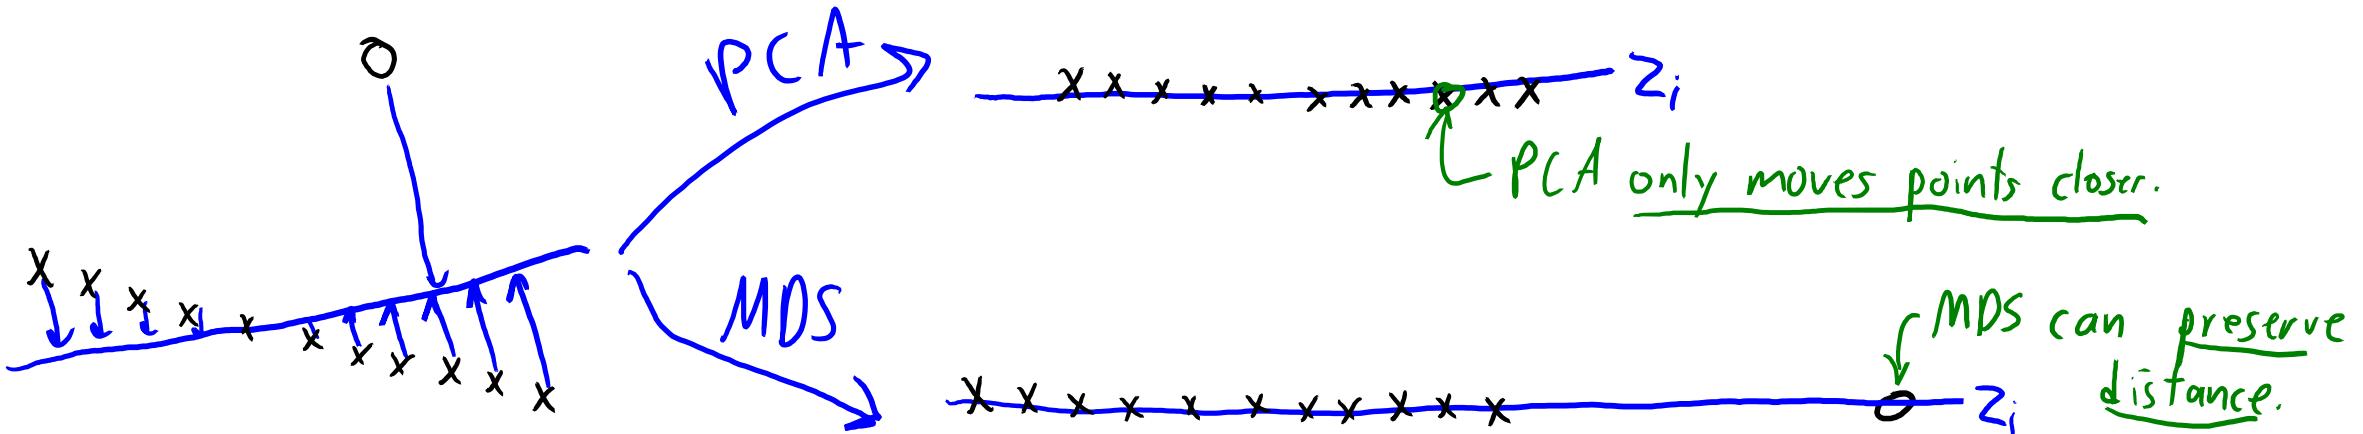


Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (\|z_i - z_j\| - \|x_i - x_j\|)^2$$

- Non-parametric dimensionality reduction and visualization:
 - No 'W': just trying to make z_i preserve high-dimensional distances between x_i .

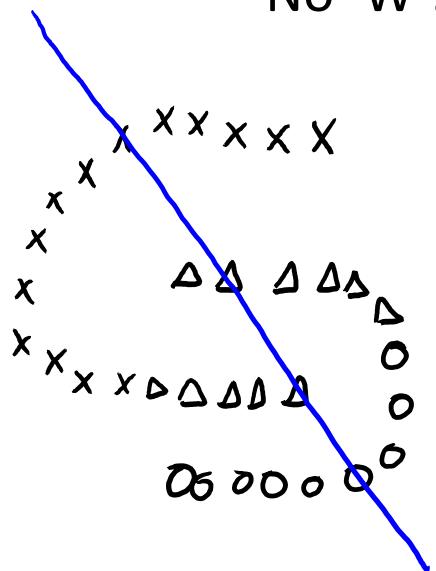


Multi-Dimensional Scaling

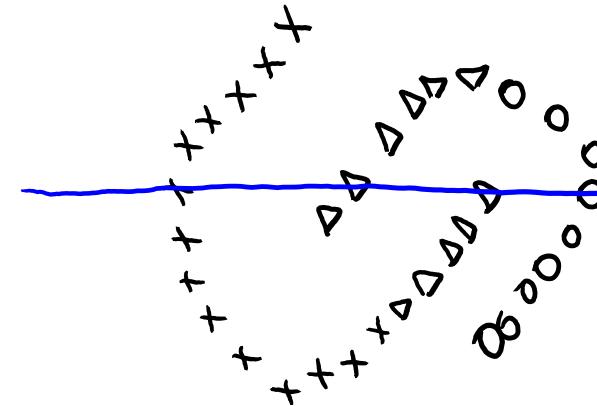
- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (||z_i - z_j|| - ||x_i - x_j||)^2$$

- Non-parametric dimensionality reduction and visualization:
 - No ‘W’: just trying to make z_i preserve high-dimensional distances between x_i .



PCA rotation:



PCA projection:

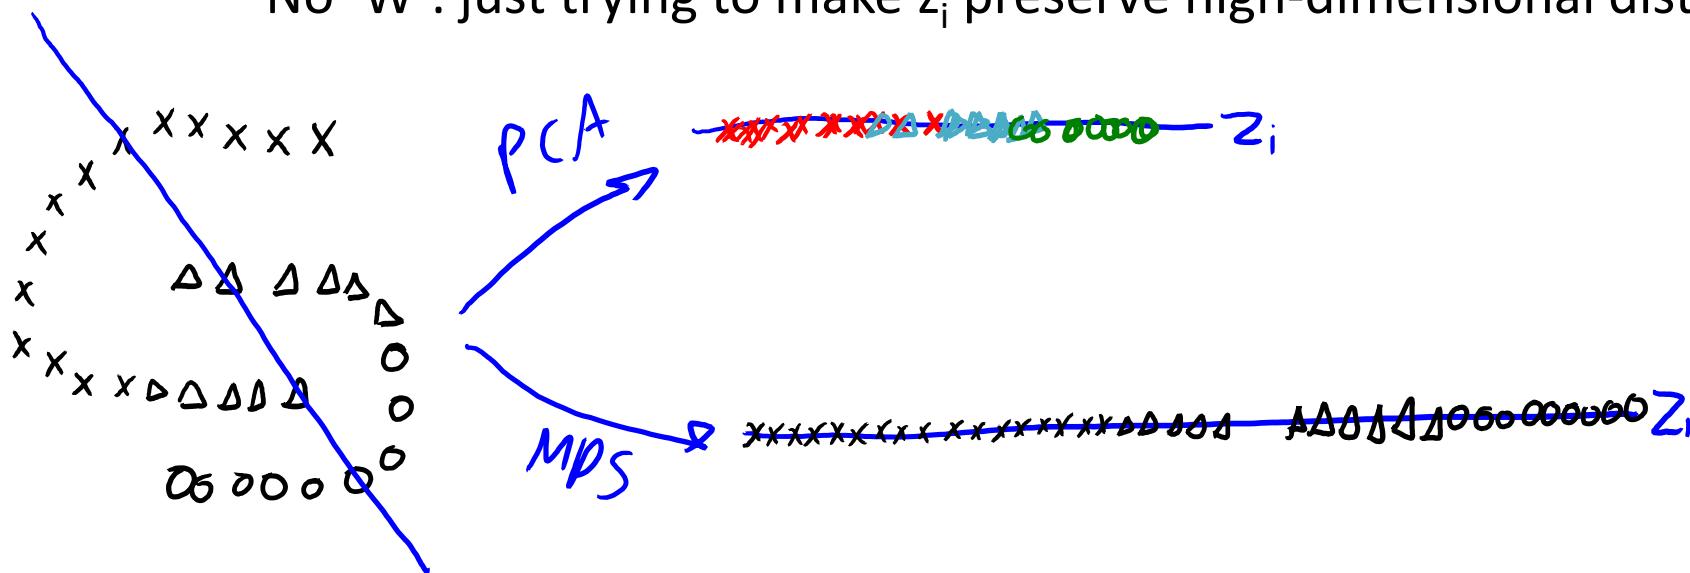


Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (||z_i - z_j|| - ||x_i - x_j||)^2$$

- Non-parametric dimensionality reduction and visualization:
 - No ‘W’: just trying to make z_i preserve high-dimensional distances between x_i .



Multi-Dimensional Scaling

- Multi-dimensional scaling (MDS):
 - Directly optimize the final locations of the z_i values.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (\|z_i - z_j\| - \|x_i - x_j\|)^2$$

- Cannot use SVD to compute solution:
 - Instead, do gradient descent on the z_i values.
 - You “learn” a scatterplot that tries to visualize high-dimensional data.
 - Not convex and sensitive to initialization.
 - And solution is not unique due to various factors like translation and rotation.

Different MDS Cost Functions

- **MDS** default objective: squared difference of Euclidean norms:

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n (\|z_i - z_j\| - \|x_i - x_j\|)^2$$

- But we can make z_i match **different distances/similarities**:

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n d_3(d_2(z_i, z_j) - d_1(x_i, x_j))$$

- Where the functions are **not necessarily the same**:

- d_1 is the high-dimensional distance we want to match.
- d_2 is the low-dimensional distance we can control.
- d_3 controls how we compare high-/low-dimensional distances.

Different MDS Cost Functions

- MDS default objective function with general distances/similarities:

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n d_3(d_2(z_i, z_j) - d_1(x_i, x_j))$$

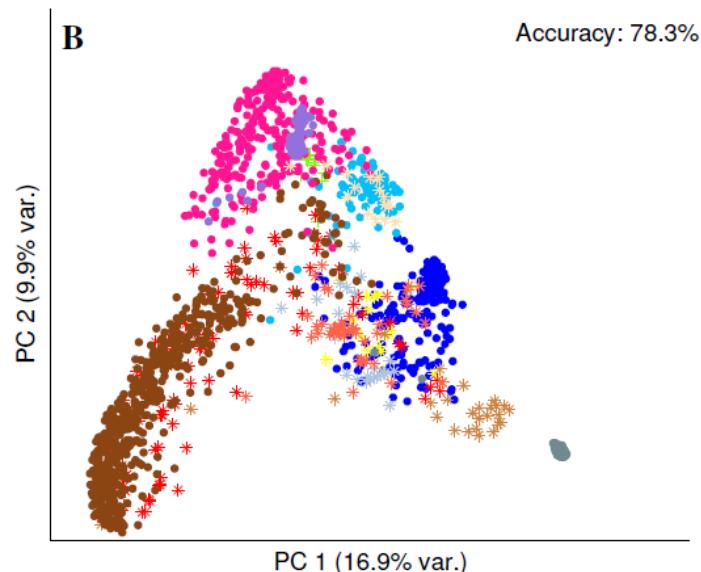
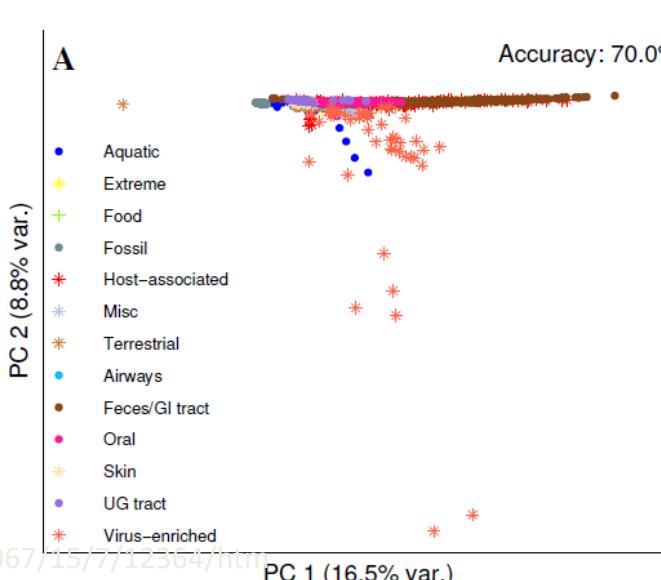
- A possibility is “classic” MDS with $d_1(x_i, x_j) = x_i^T x_j$ and $d_2(z_i, z_j) = z_i^T z_j$.
 - We obtain PCA in this special case (centered x_i , d_3 as the squared L2-norm).
 - Not a great choice because it’s a linear model.

Different MDS Cost Functions

- MDS default objective function with general distances/similarities:

$$f(z) = \hat{\sum}_{i=1}^n \sum_{j=i+1}^n d_3(d_2(z_i, z_j) - d_1(x_i, x_j))$$

- Another possibility: $d_1(x_i, x_j) = ||x_i - x_j||_1$ and $d_2(z_i, z_j) = ||z_i - z_j||$.
 - The z_i approximate the high-dimensional L_1 -norm distances.



Sammon's Mapping

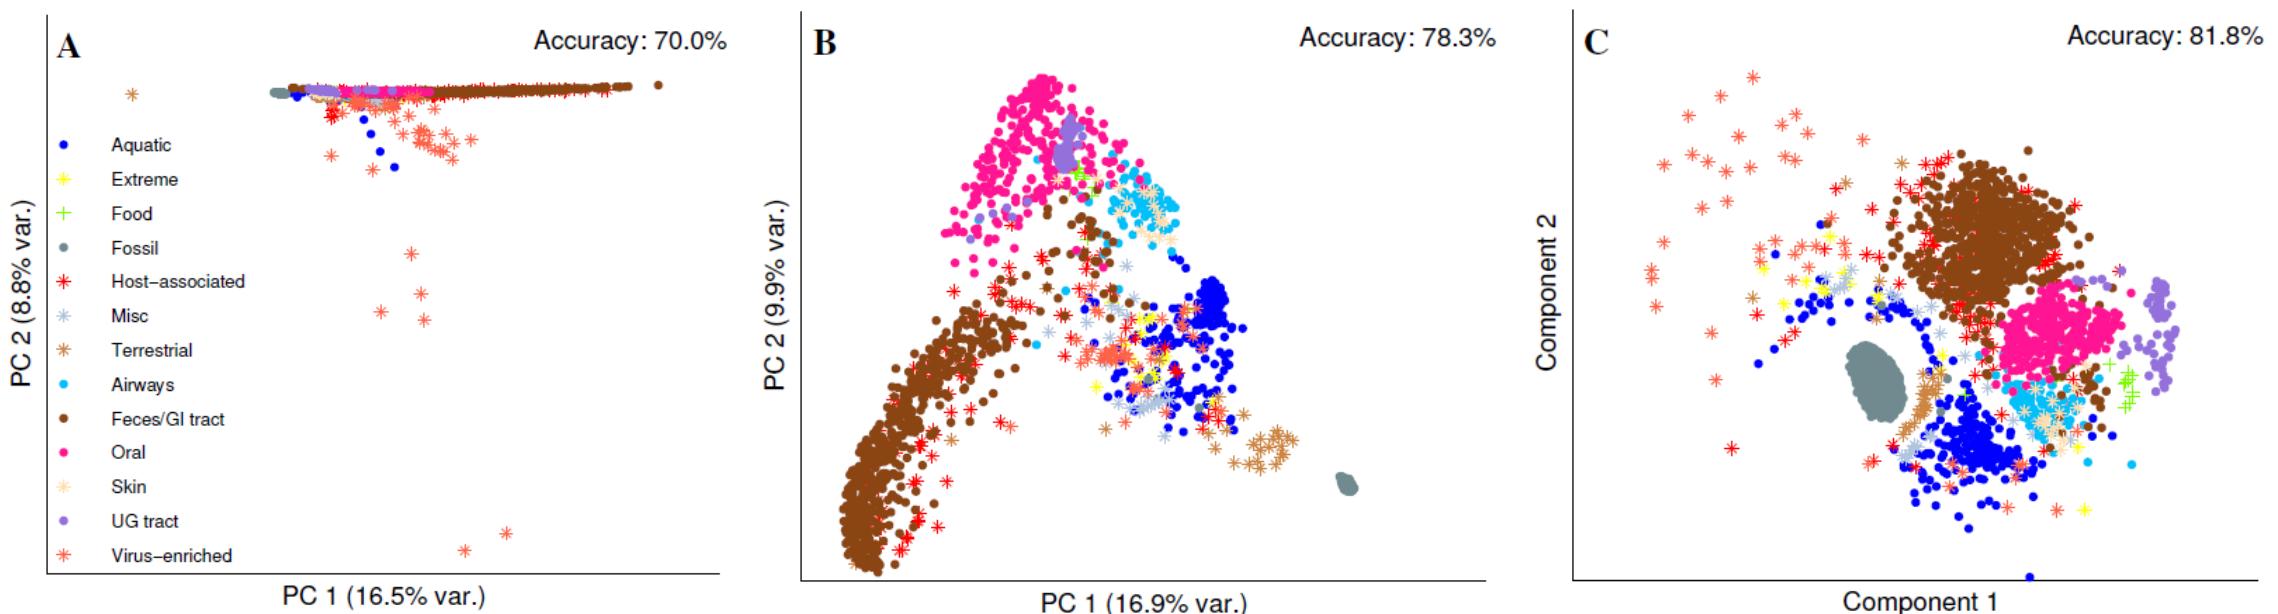
- Challenge for most MDS models: they **focus on large distances**.
 - Leads to “crowding” effect like with PCA.
- Early attempt to address this is **Sammon’s mapping**:
 - **Weighted MDS** so large/small distances are more comparable.

$$f(z) = \sum_{i=1}^n \sum_{j=i+1}^n \left(\frac{d_2(z_i, z_j) - d_1(x_i, x_j)}{d_1(x_i, x_j)} \right)^2$$

- Denominator reduces focus on large distances.

Sammon's Mapping

- Challenge for most MDS models: they **focus on large distances**.
 - Leads to “crowding” effect like with PCA.
- Early attempt to address this is **Sammon’s mapping**:
 - Weighted MDS** so large/small distances are more comparable.

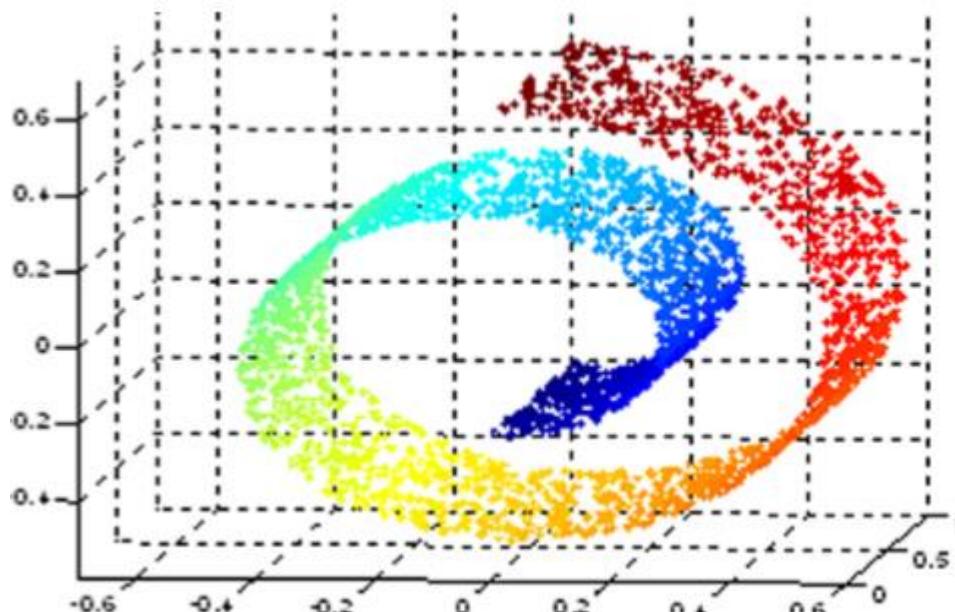


(pause)

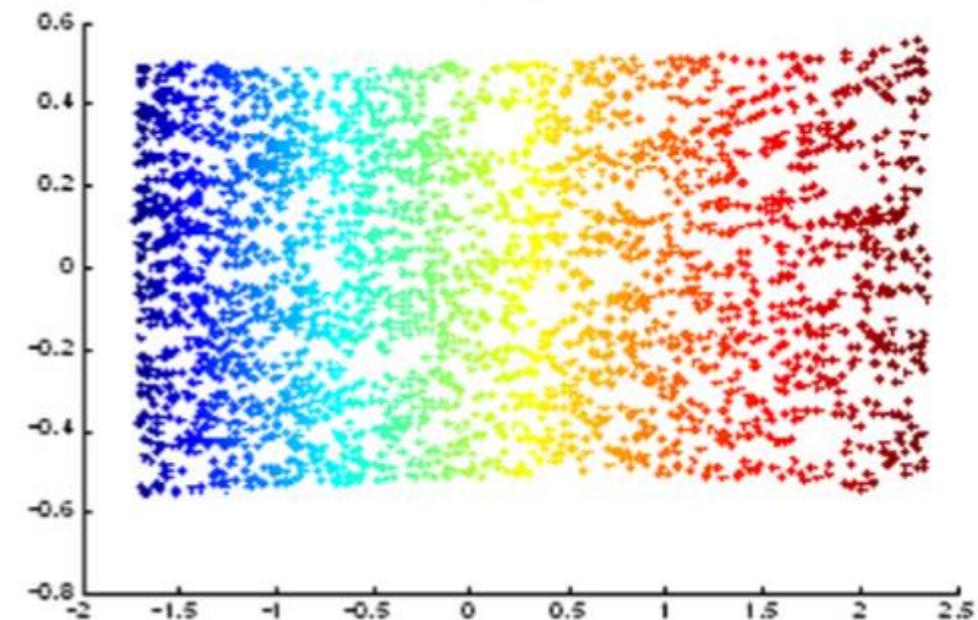
Learning Manifolds

- Consider data that lives on a **low-dimensional “manifold”**.
- Example is the ‘Swiss roll’:

Original data

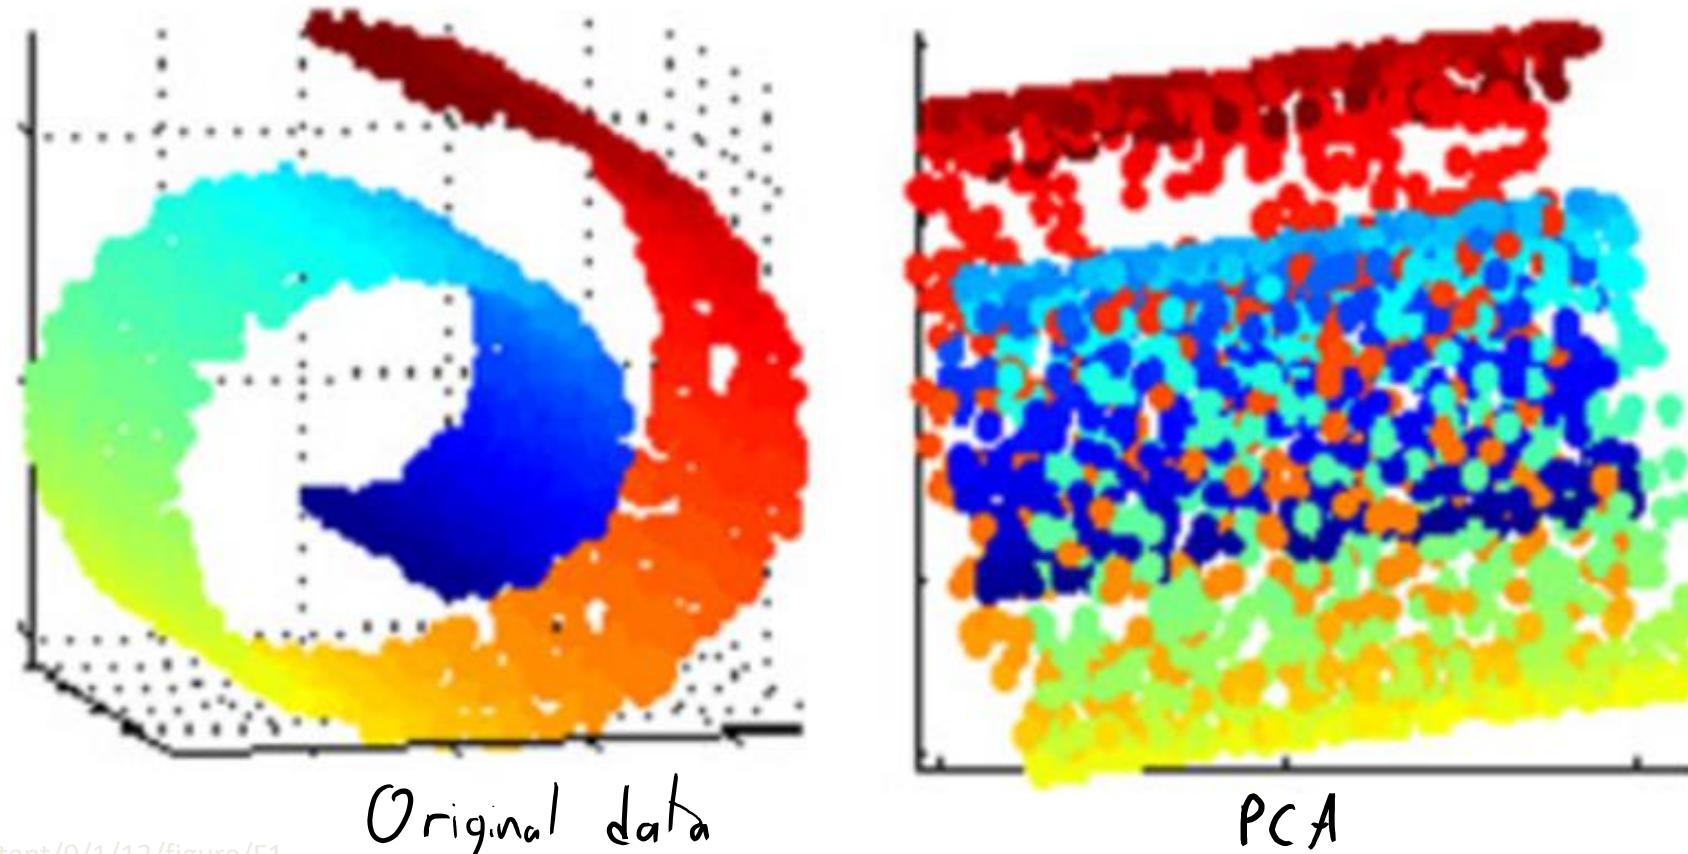


Two-dimensional manifold



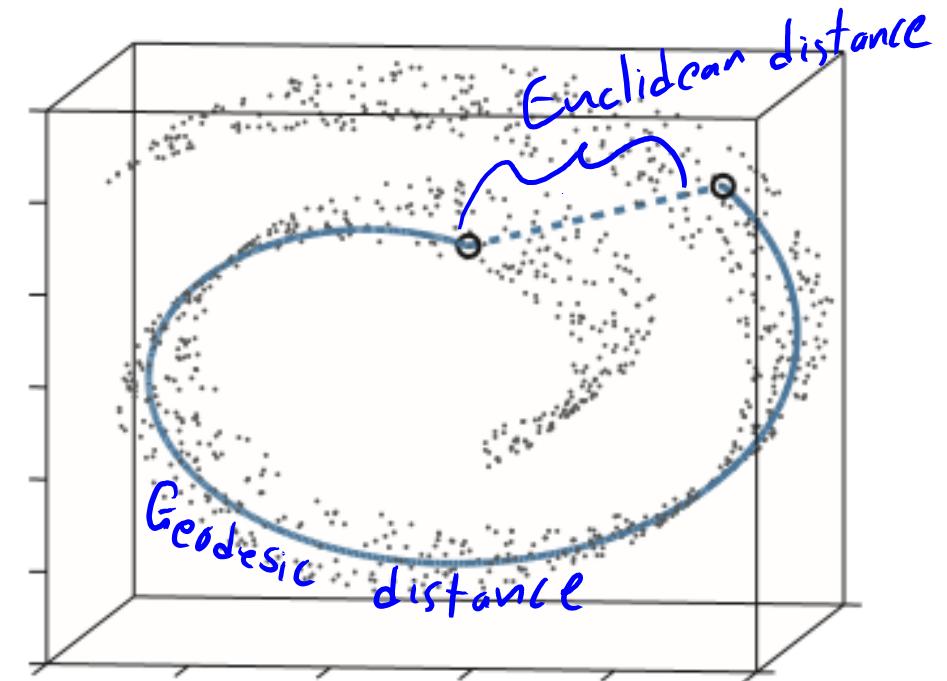
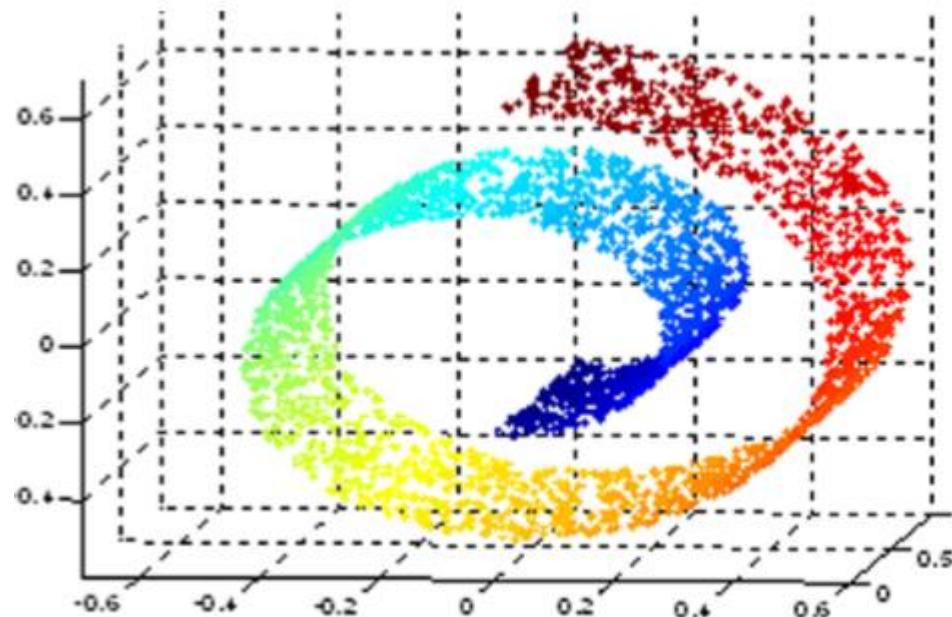
Learning Manifolds

- Consider data that lives on a **low-dimensional “manifold”**.
 - With usual distances, **PCA/MDS will not discover non-linear manifolds.**



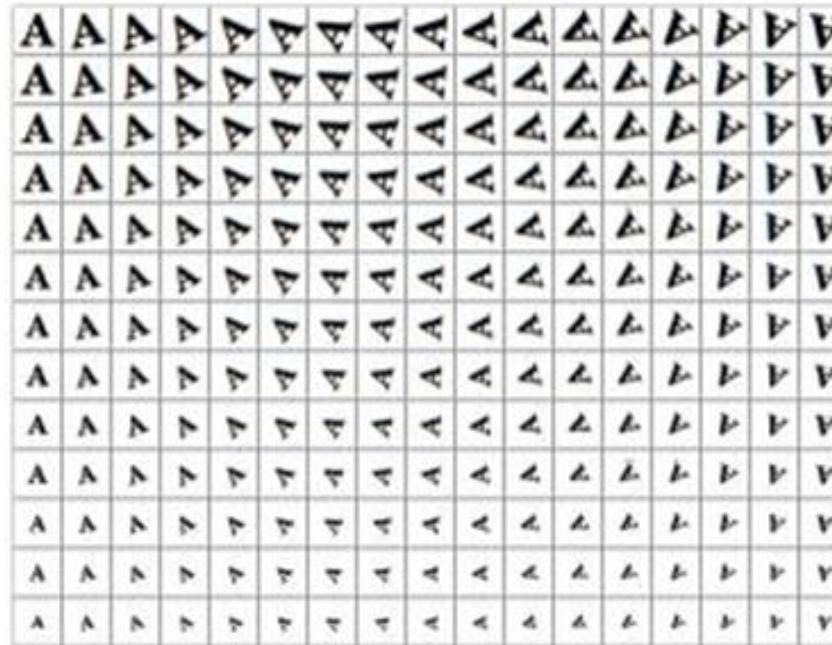
Learning Manifolds

- Consider data that lives on a **low-dimensional “manifold”**.
 - With usual distances, **PCA/MDS will not discover non-linear manifolds.**
- We need **geodesic distance**: the distance *through* the manifold.



Manifolds in Image Space

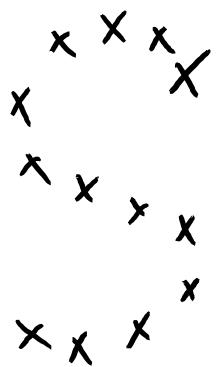
- Consider slowly-varying transformation of image:

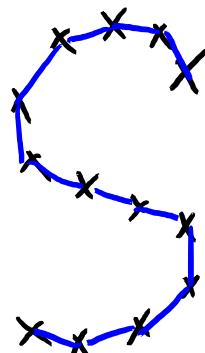


- Images are on a manifold in the high-dimensional space.
 - Euclidean distance doesn't reflect manifold structure.
 - Geodesic distance is distance through space of rotations/resizings.

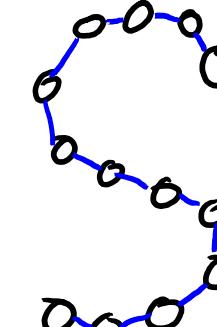
ISOMAP

- ISOMAP is latent-factor model for visualizing data on manifolds:

 find "neighbours" of each point



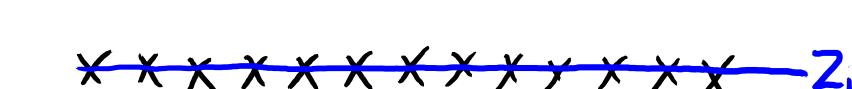
Represent points and neighbours as a weighted graph.



"weight" on each edge is distance between points

↓ Approximate geodesic distance by shortest path through graph.

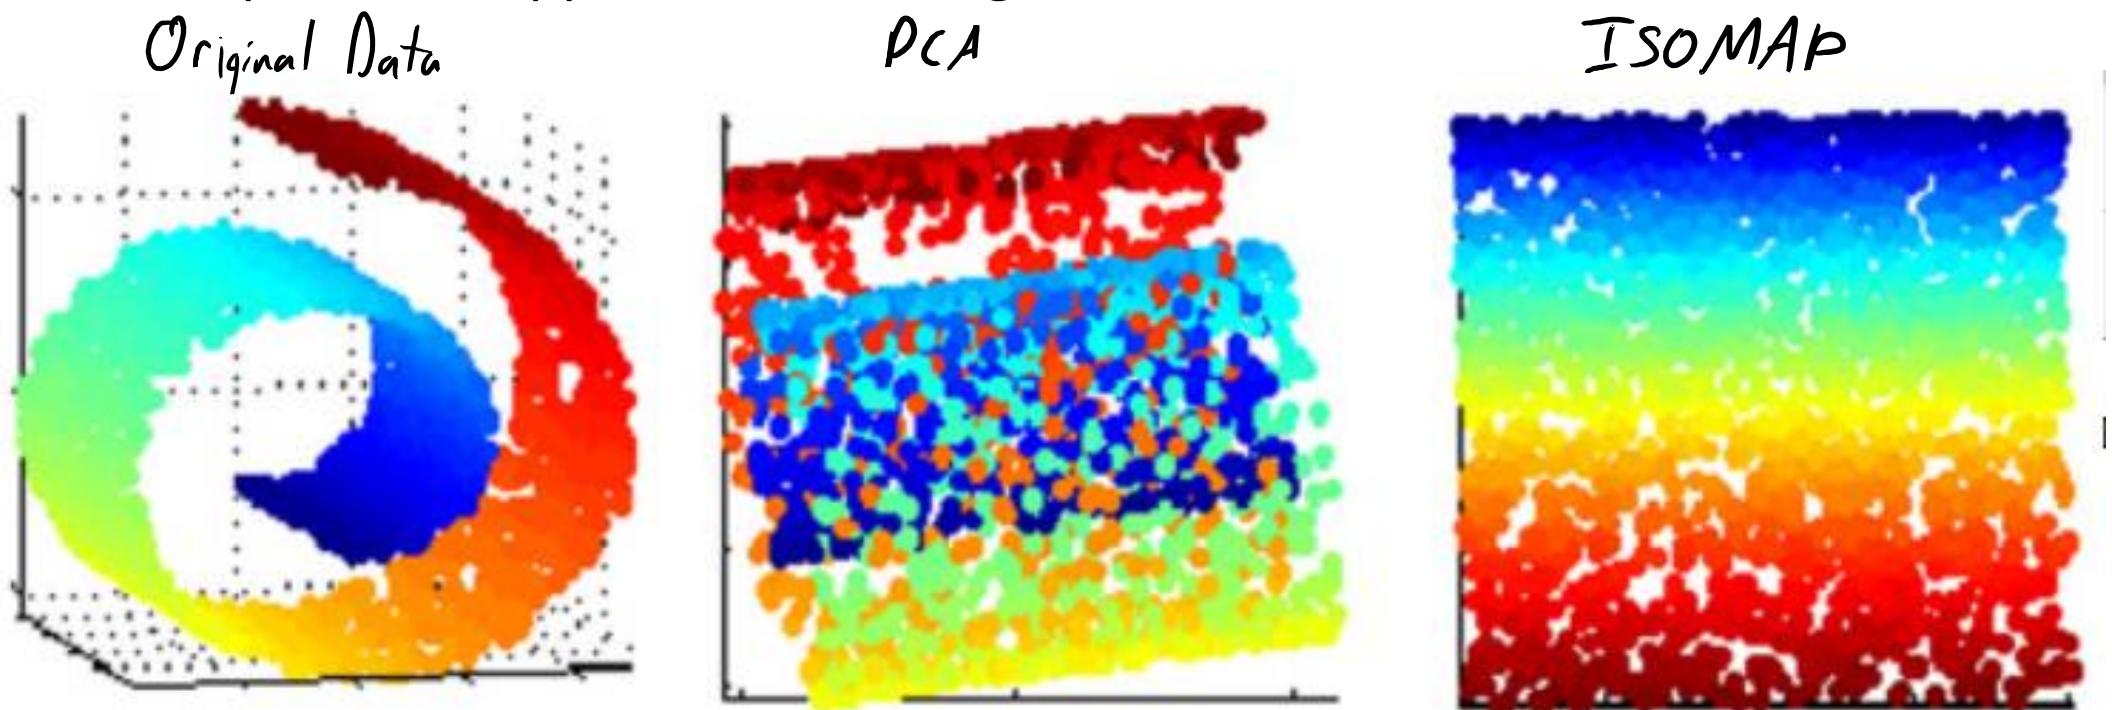
$$D = \begin{bmatrix} 0 & 1 & 2 & 3 & \cdots \\ 1 & 0 & 1 & 2 & \cdots \\ 2 & 1 & 0 & 1 & \cdots \\ 3 & 2 & 1 & 0 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$


ISOMAP z_i values in 1D or 2D

Run MDS with these approximate geodesic distances.

ISOMAP

- ISOMAP can “unwrap” the roll:
 - Shortest paths are approximations to geodesic distances.



- Sensitive to having the right graph:
 - Points off of manifold and gaps in manifold cause problems.

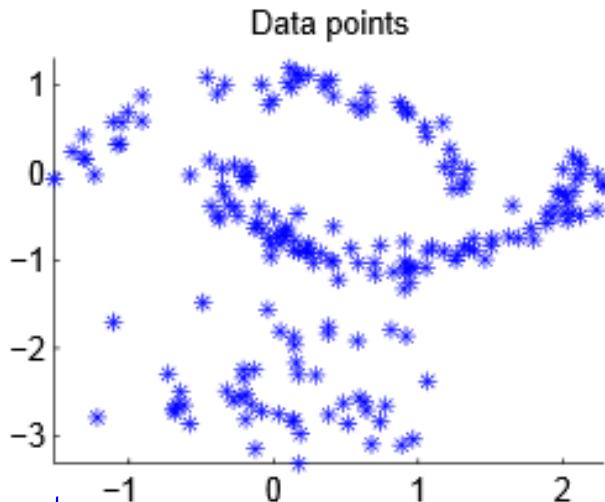
Constructing Neighbour Graphs

- Sometimes you can **define the graph/distance without features**:
 - Facebook friend graph.
 - Connect YouTube videos if one video tends to follow another.
- But we can also **convert from features x_i to a “neighbour” graph**:
 - Approach 1 (“**epsilon graph**”): connect x_i to all x_j within some threshold ε .
 - Like we did with density-based clustering.
 - Approach 2 (“**KNN graph**”): connect x_i to x_j if:
 - x_j is a KNN of x_i **OR** x_i is a KNN of x_j .
 - Approach 2 (“**mutual KNN graph**”): connect x_i to x_j if:
 - x_j is a KNN of x_i **AND** x_i is a KNN of x_j .

Converting from Features to Graph

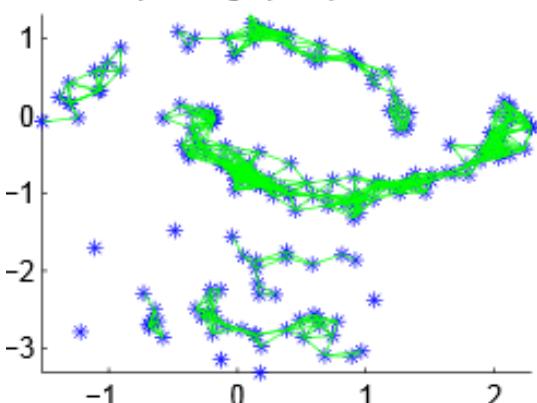
add edge if $\|x_i - x_j\| \leq 0.3$

or
 i is NN of j or j is NN of i

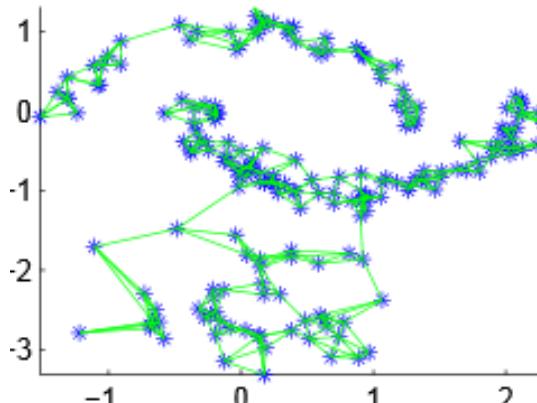


add edge if i and j are kNNs of each other.

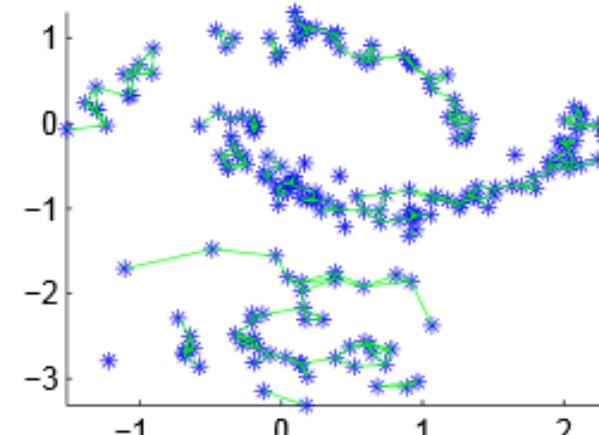
epsilon-graph, epsilon=0.3



kNN graph, k = 5

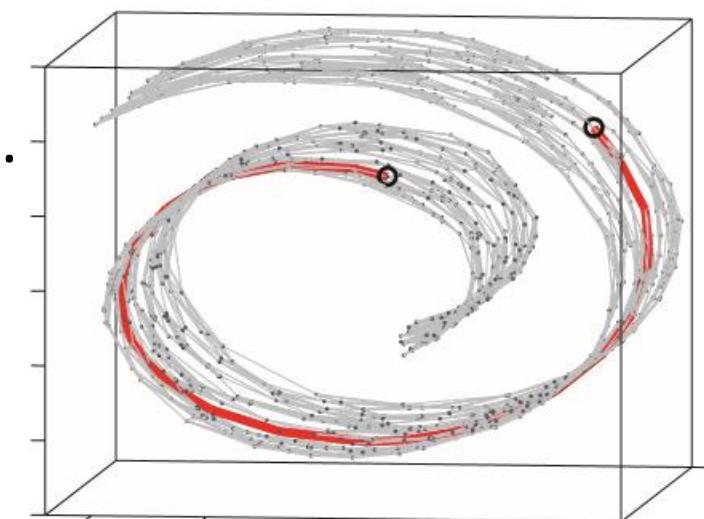


Mutual kNN graph, k = 5

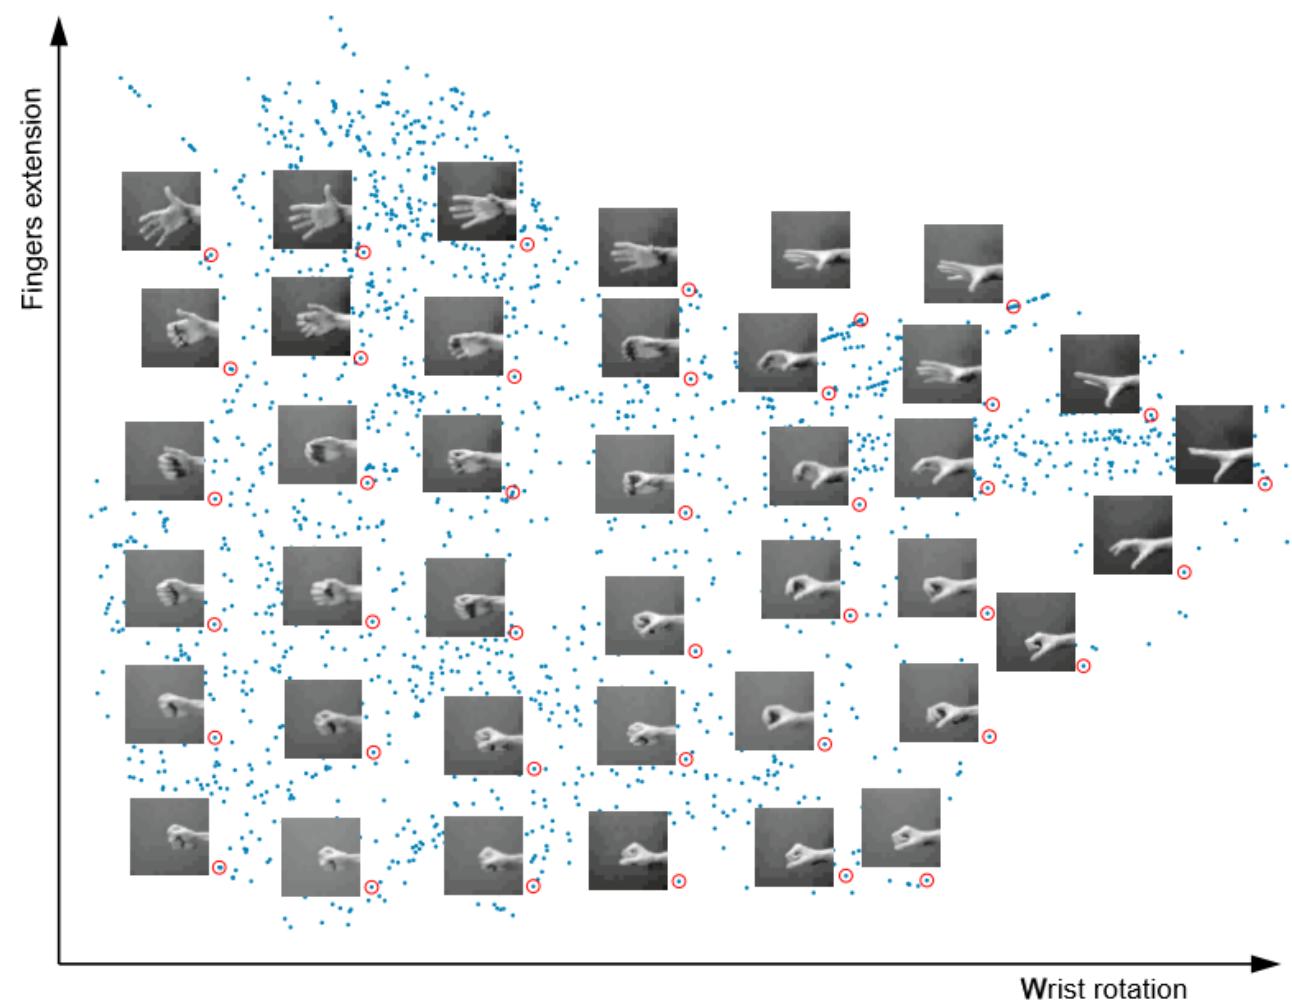


ISOMAP

- ISOMAP is latent-factor model for visualizing data on manifolds:
 1. Find the neighbours of each point.
 - Usually “k-nearest neighbours graph”, or “epsilon graph”.
 2. Compute edge weights:
 - Usually distance between neighbours.
 3. Compute weighted shortest path between all points.
 - Dijkstra or other shortest path algorithm.
 4. Run MDS using these distances.



ISOMAP on Hand Images



- Related method is “local linear embedding”.

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

Sammon's Map vs. ISOMAP vs. PCA

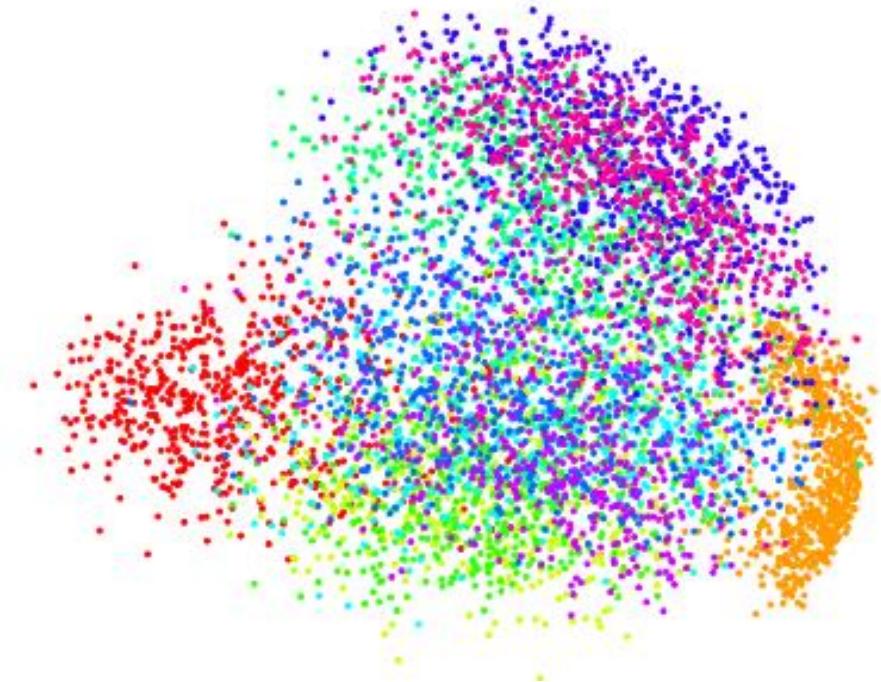
Sammon Map



ISOMAP



PCA



- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

Sammon's Map vs. ISOMAP vs. t-SNE

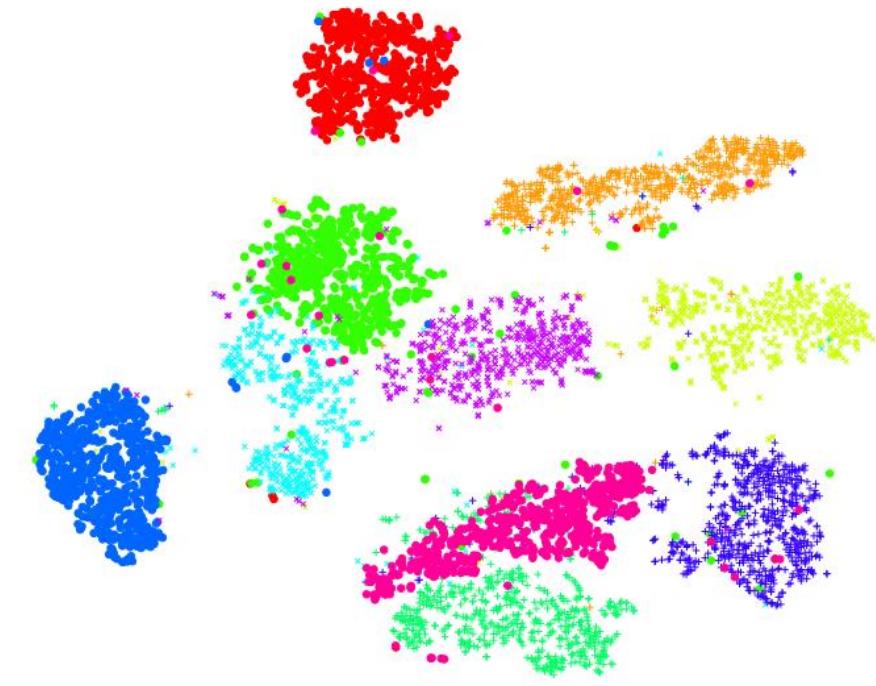
Sammon Map



ISOMAP



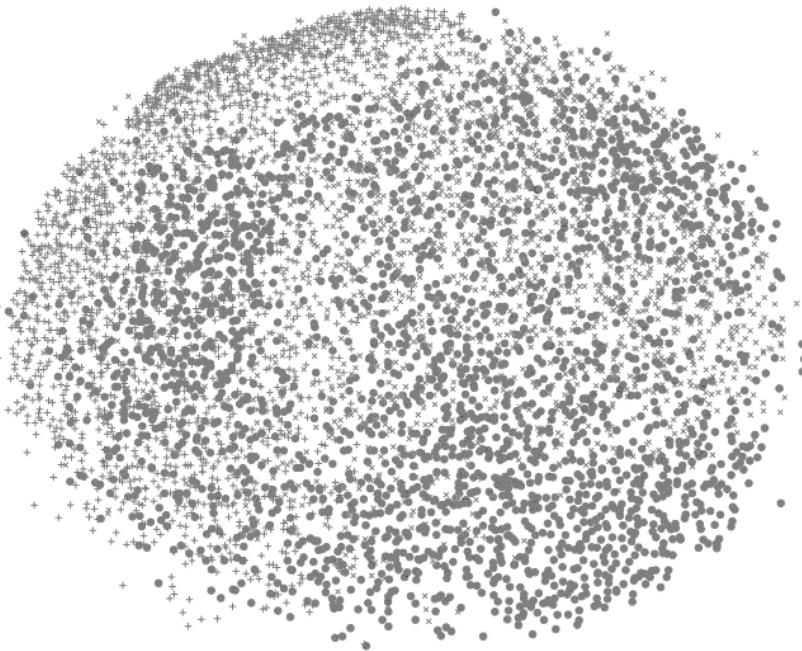
t-SNE



- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

Sammon's Map vs. ISOMAP vs. t-SNE

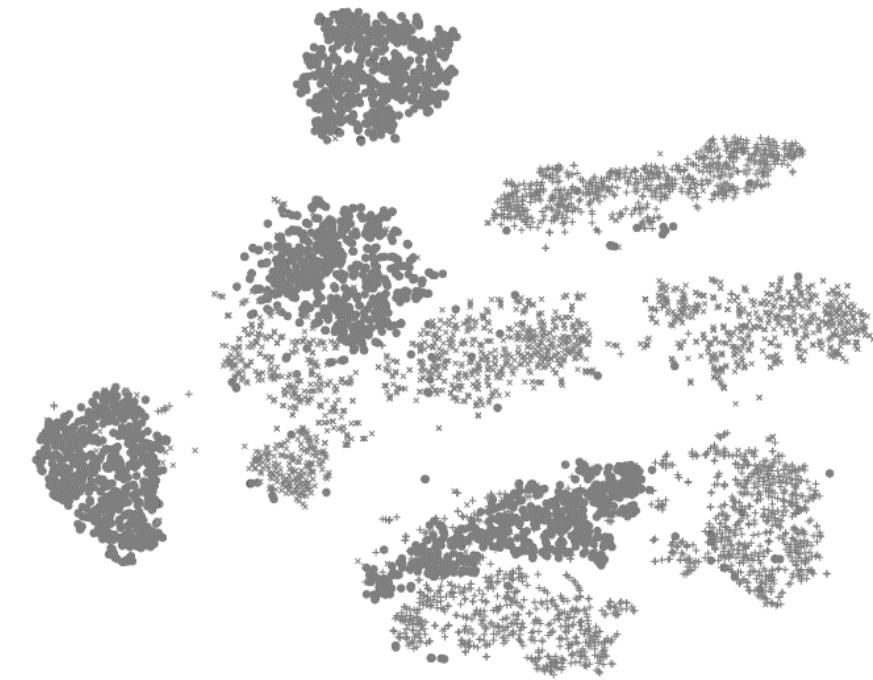
Sammon Map



ISOMAP



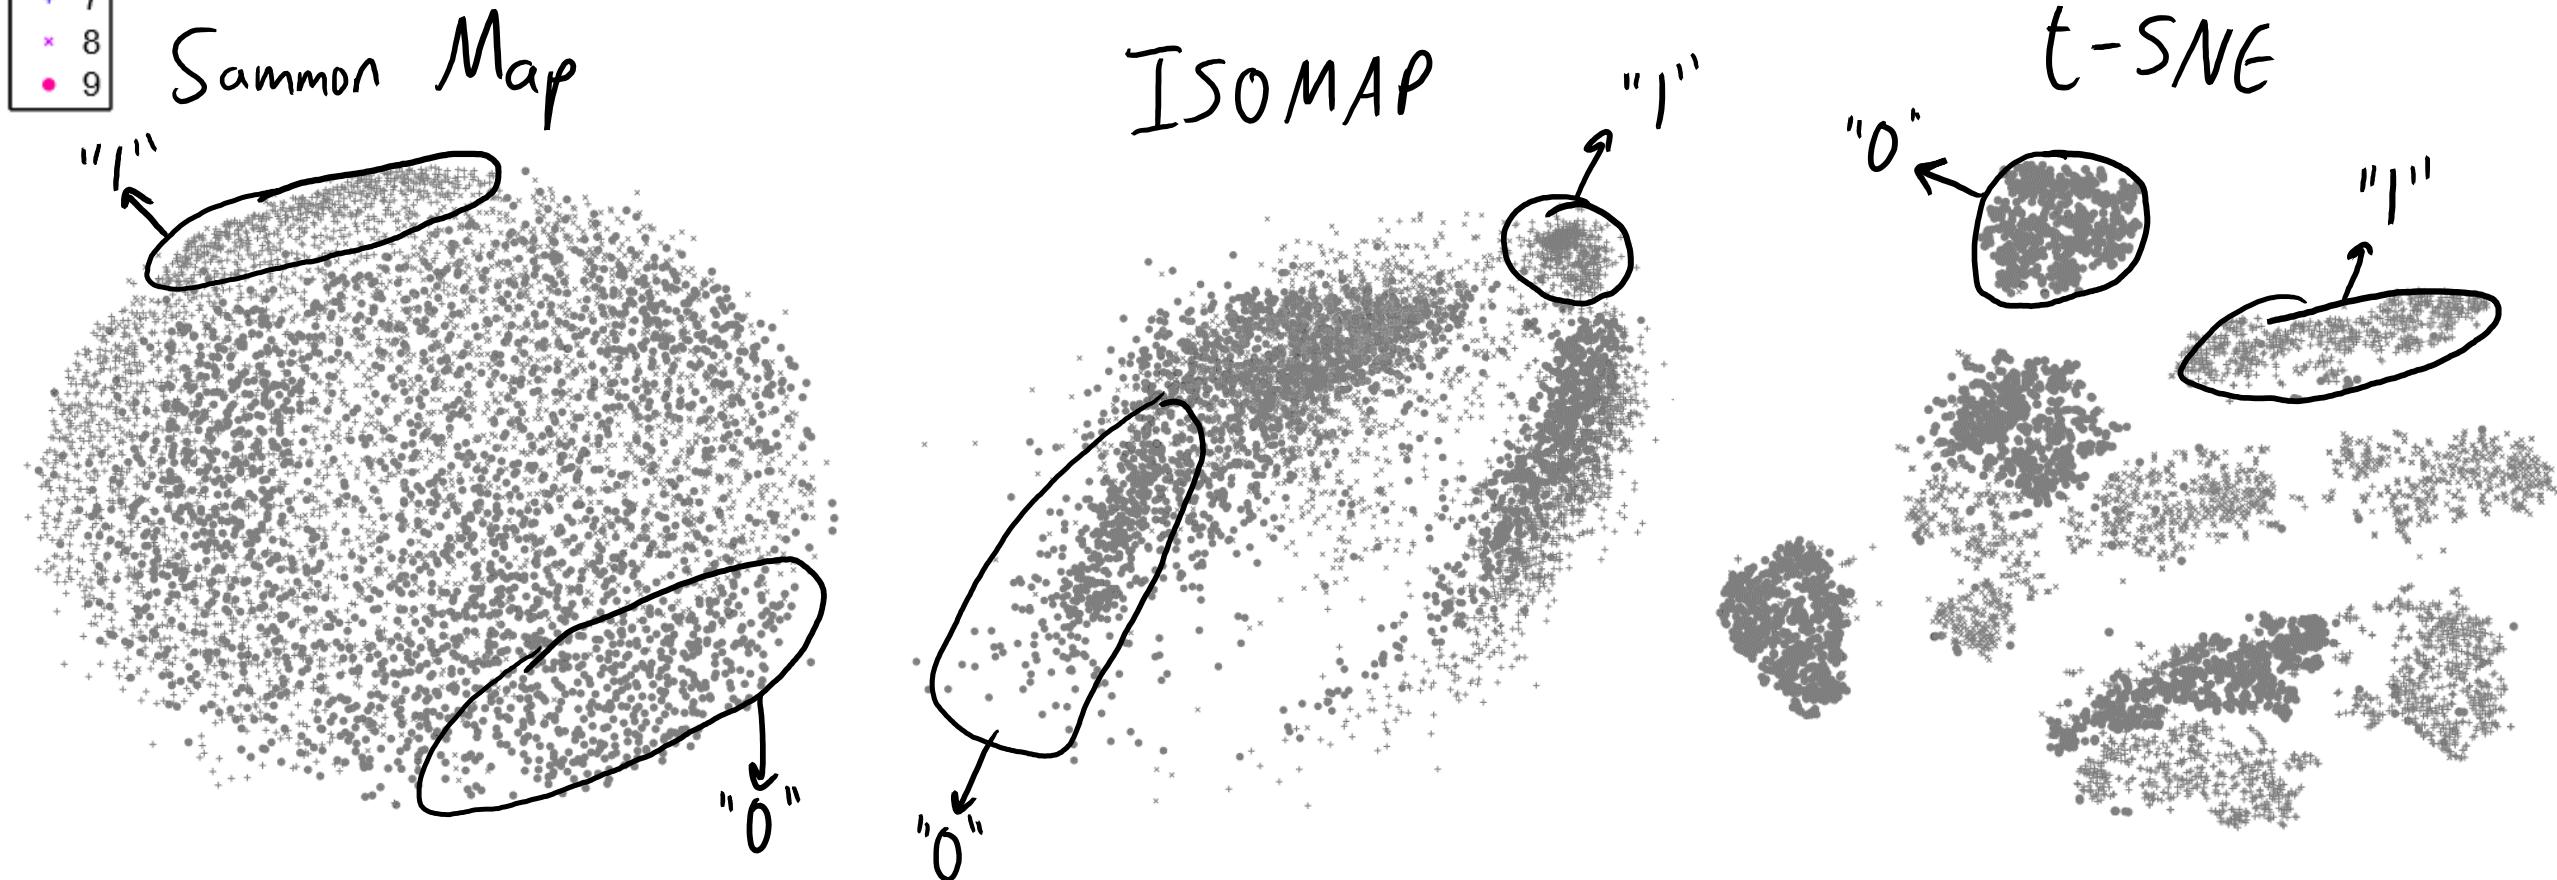
t-SNE



Remember this is unsupervised, algorithms do not know the labels.

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

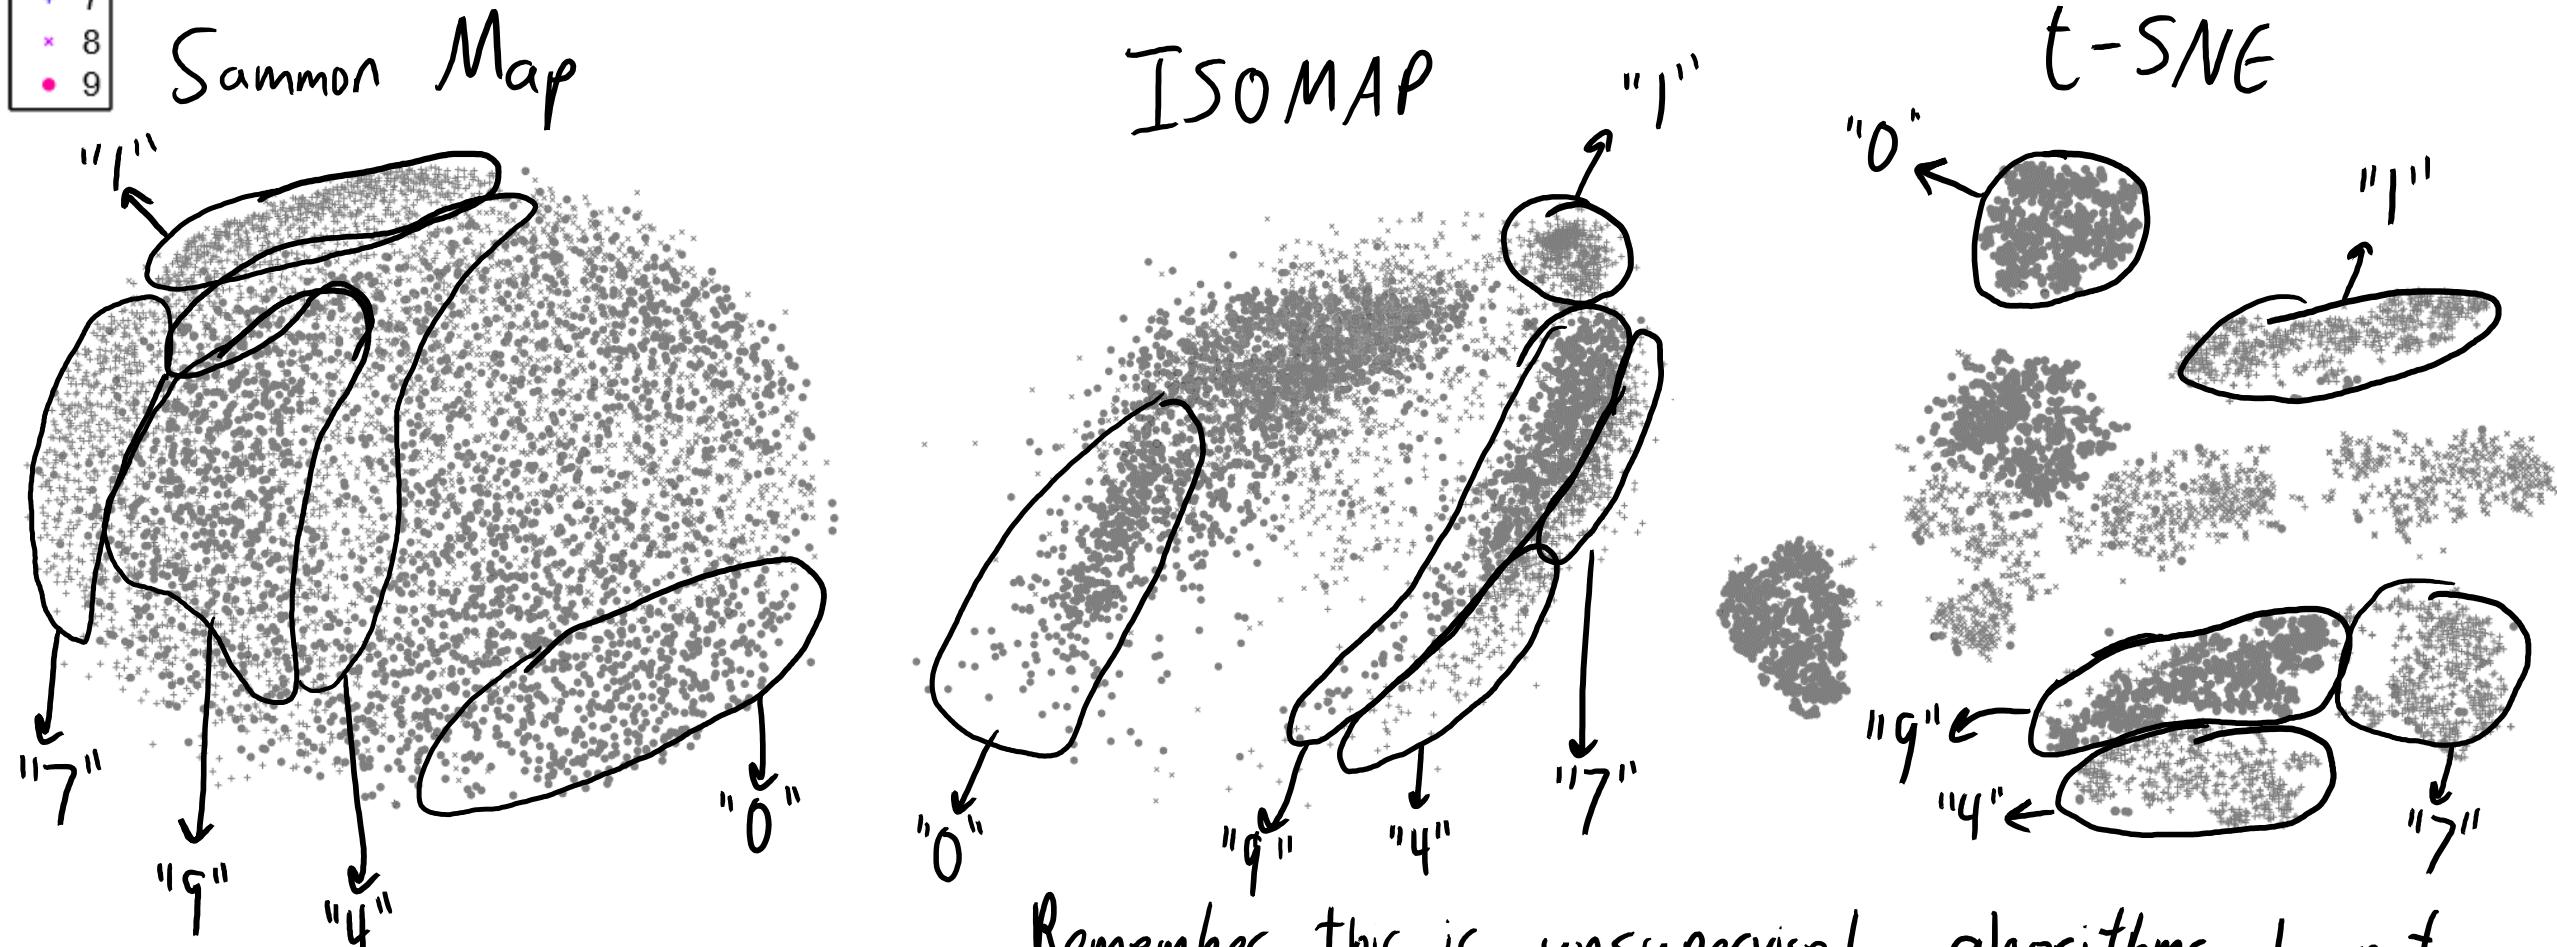
Sammon's Map vs. ISOMAP vs. t-SNE



Remember this is unsupervised, algorithms do not know the labels.

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

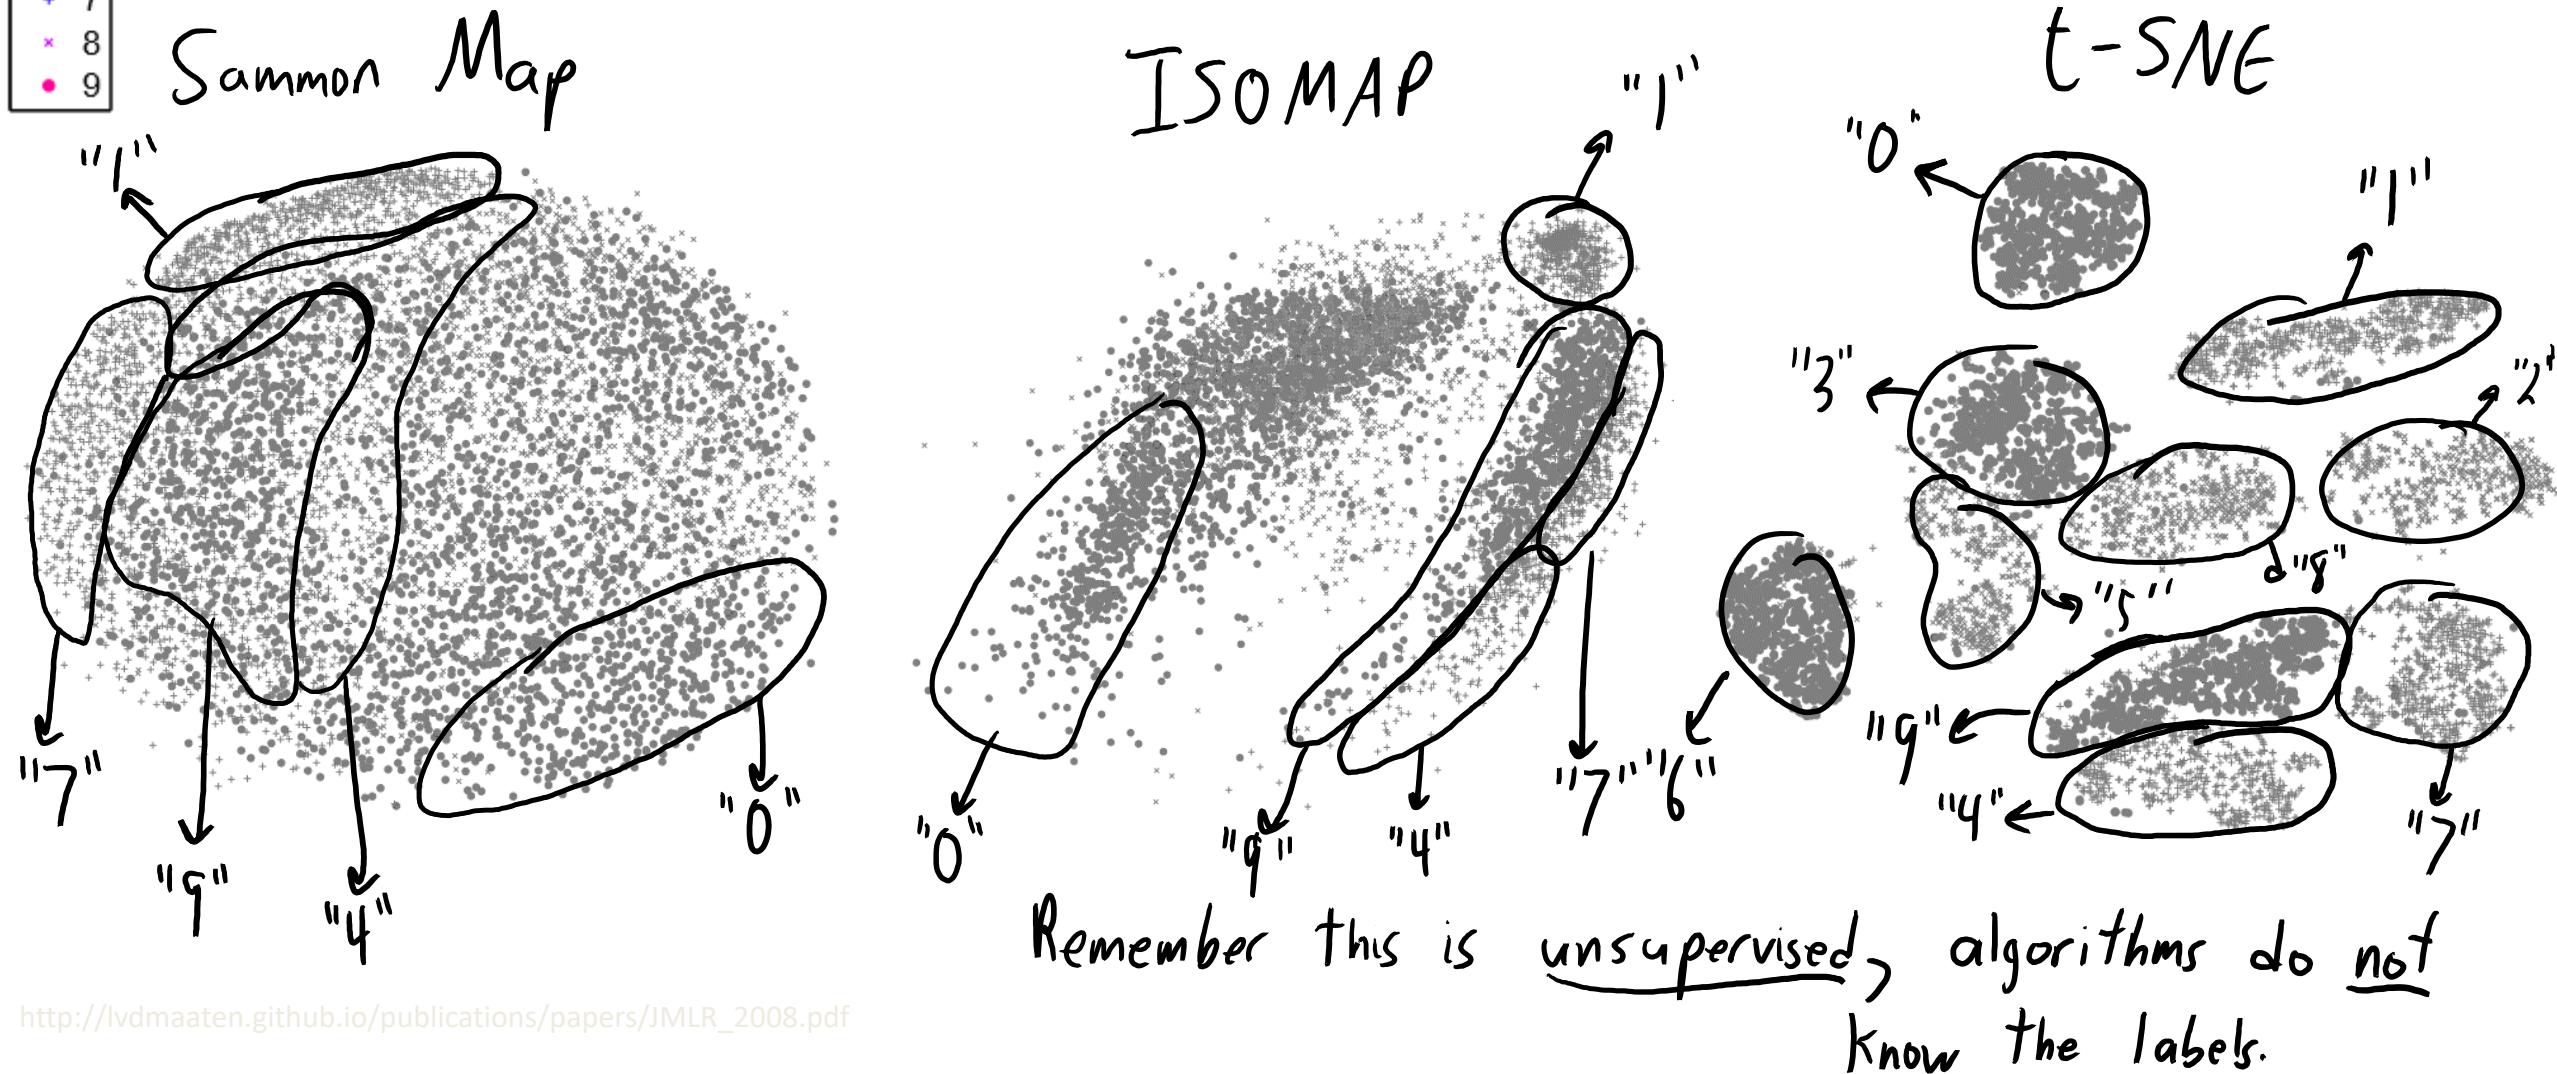
Sammon's Map vs. ISOMAP vs. t-SNE



Remember this is unsupervised, algorithms do not know the labels.

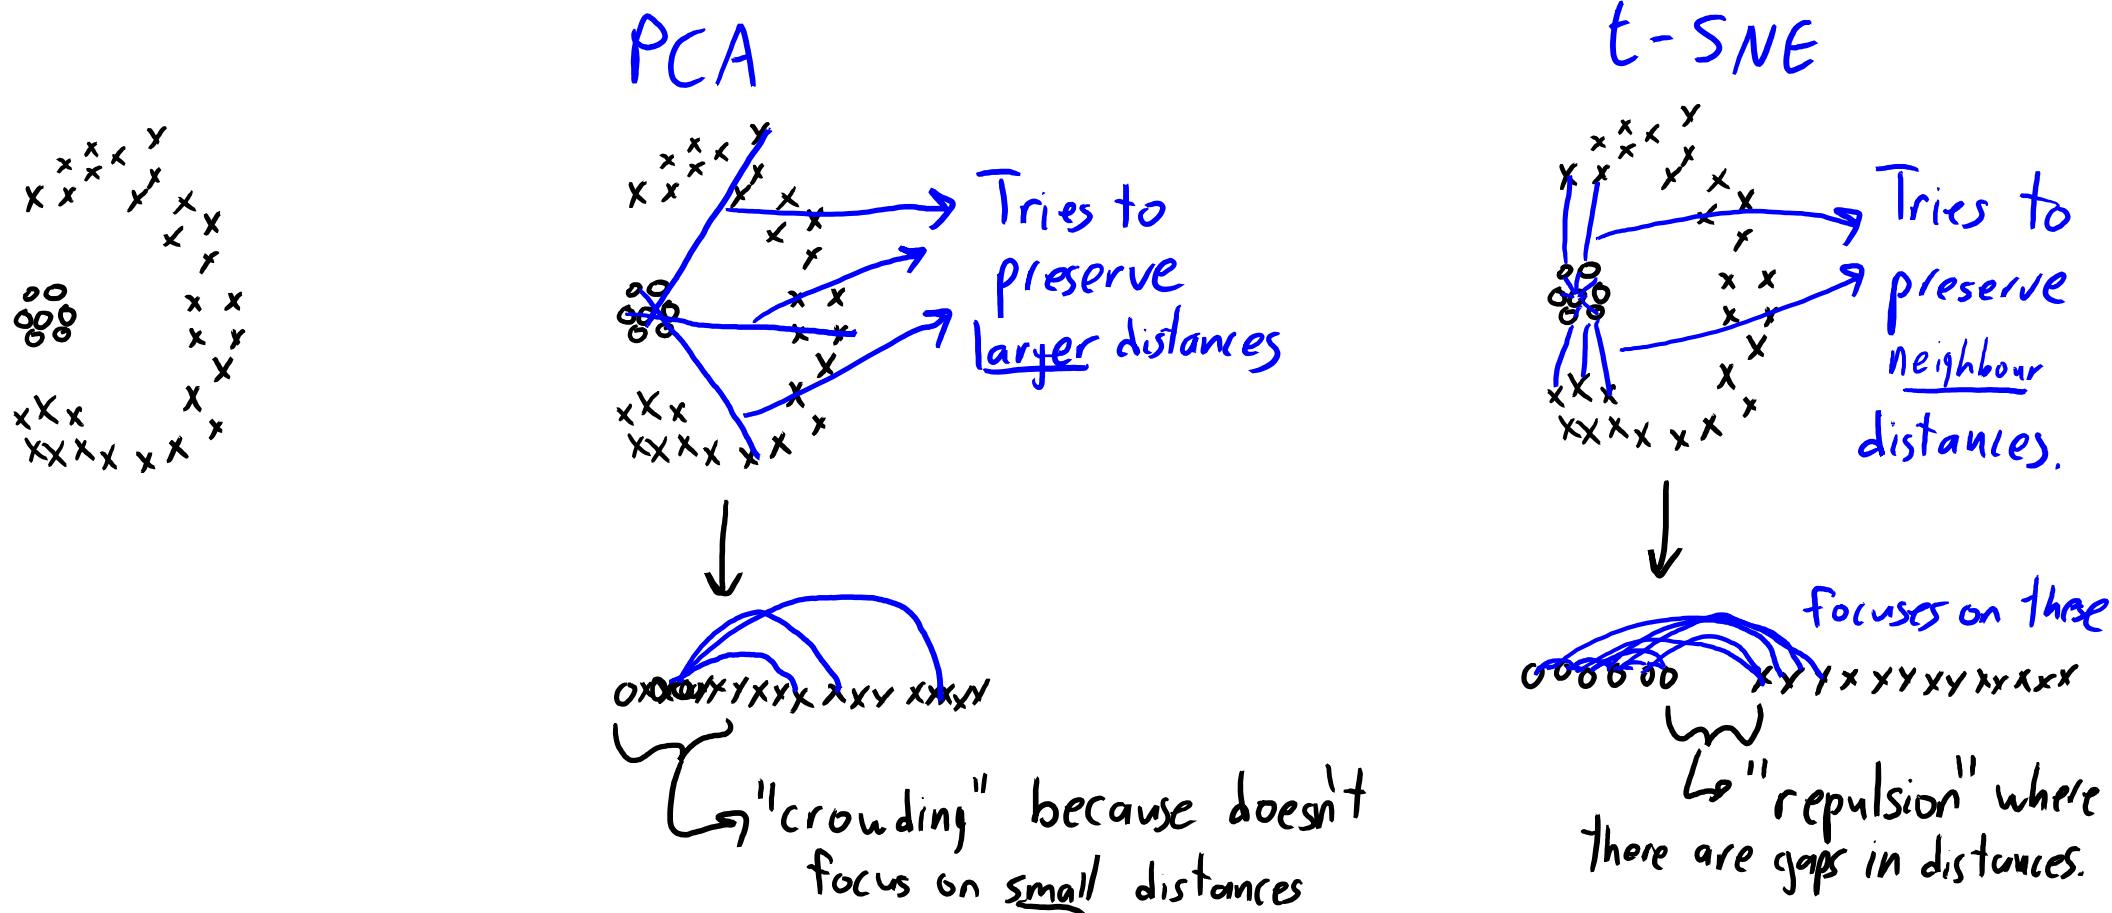
- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

Sammon's Map vs. ISOMAP vs. t-SNE



t-Distributed Stochastic Neighbour Embedding

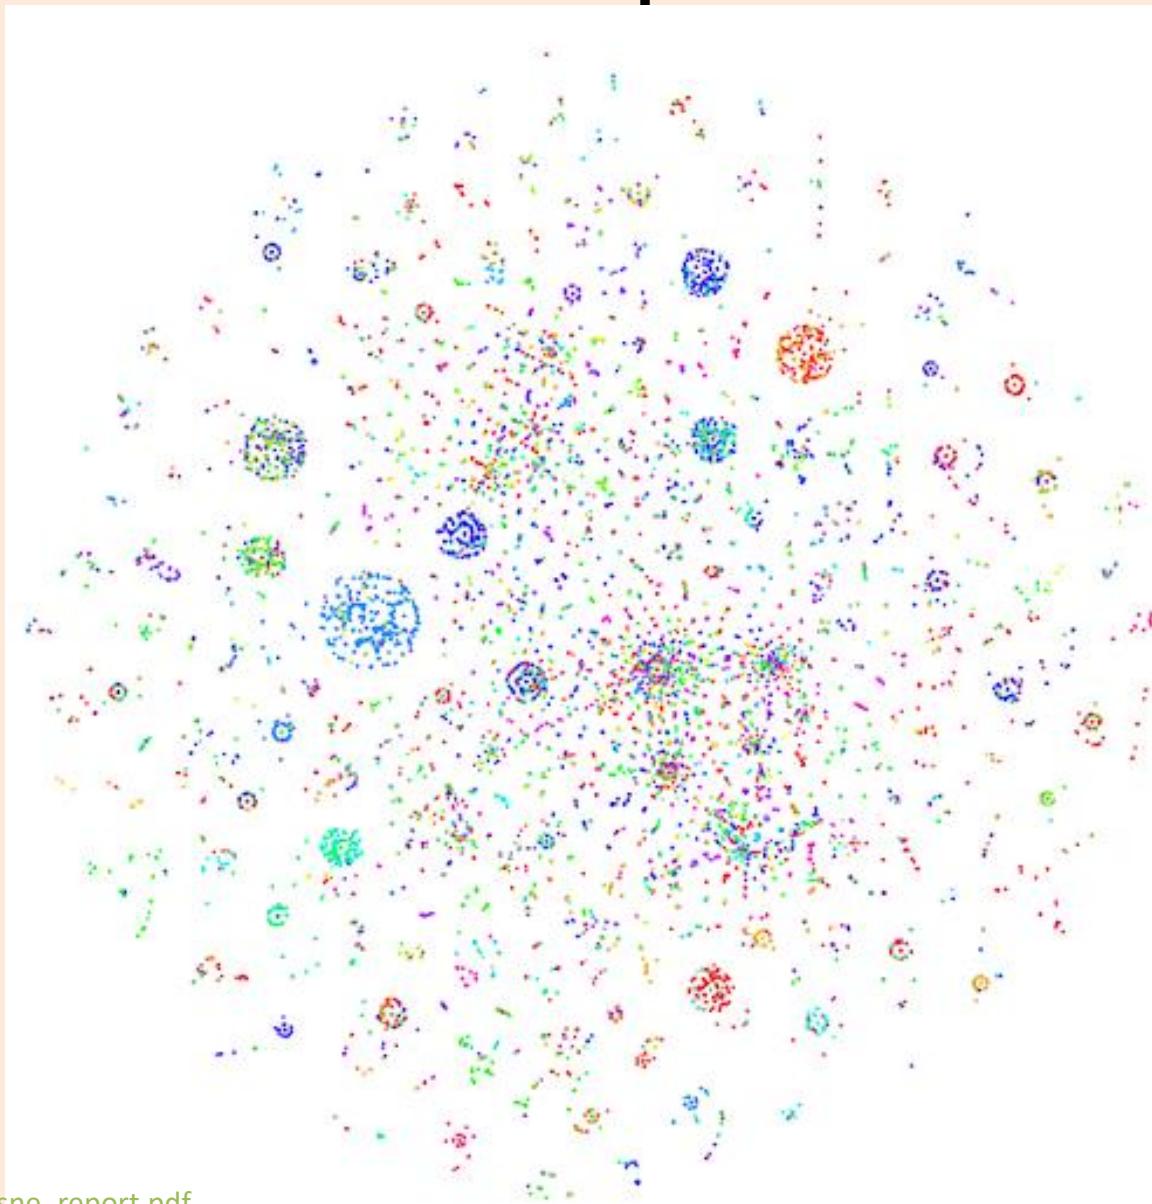
- One key idea in t-SNE:
 - Focus on distance to “neighbours”(allow large variance in other distances)



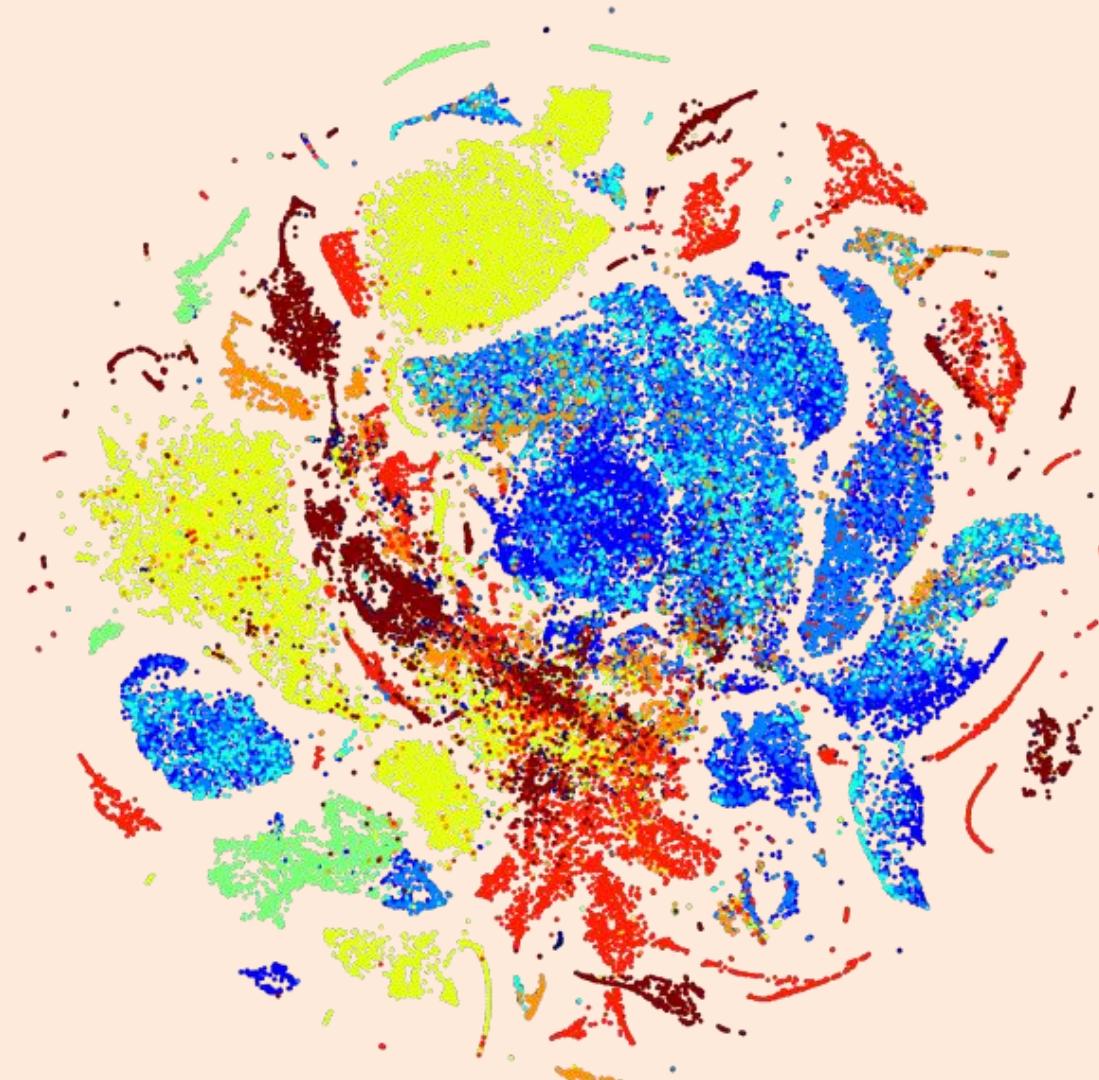
t-Distributed Stochastic Neighbour Embedding

- t-SNE is a special case of MDS (specific d_1 , d_2 , and d_3 choices):
 - d_1 : for each x_i , compute probability that each x_j is a ‘neighbour’.
 - Computation is similar to k-means++, but most weight to close points (Gaussian).
 - Doesn’t require explicit graph.
 - d_2 : for each z_i , compute probability that each z_j is a ‘neighbour’.
 - Similar to above, but uses student’s t (grows really slowly with distance).
 - Avoids ‘crowding’, because you have a huge range that large distances can fill.
 - d_3 : Compare x_i and z_i using an entropy-like measure:
 - How much ‘randomness’ is in probabilities of x_i if you know the z_i (and vice versa)?
- Interactive demo: <https://distill.pub/2016/misread-tsne>

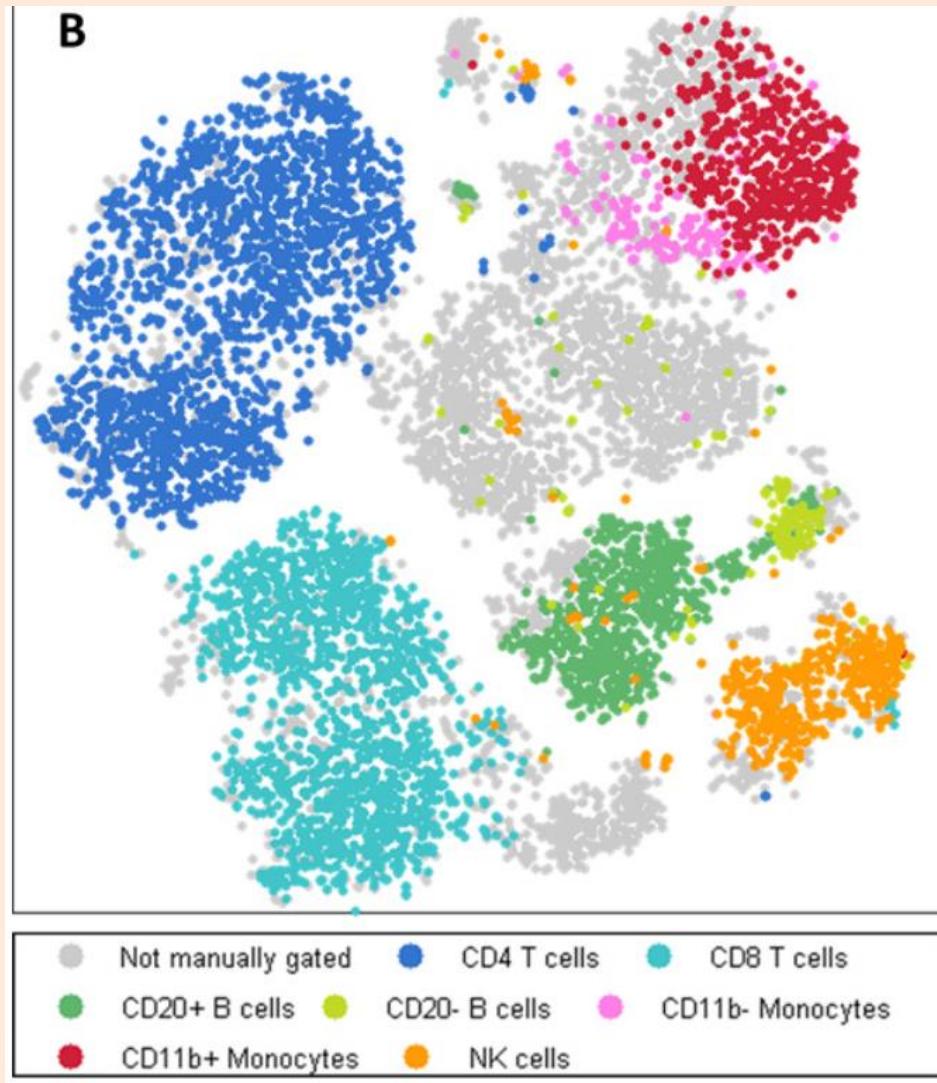
t-SNE on Wikipedia Articles



t-SNE on Product Features



t-SNE on Leukemia Heterogeneity



(pause)

Latent-Factor Representation of Words

- For natural language, we often **represent words by an index**.
 - E.g., “cat” is word 124056 among a “bag of words”.
- But this may be inefficient:
 - Should “cat” and “kitten” **share parameters** in some way?
- We want a **latent-factor representation** of individual words:
 - Closeness in latent space should indicate similarity.
 - Distances could represent meaning?
- Recent alternative to PCA/NMF is **word2vec...**

Using Context

- Consider these phrases:
 - “the cat purred”
 - “the kitten purred”
 - “black cat ran”
 - “black kitten ran”
- Words that occur in the same context likely have similar meanings.
- Word2vec uses this insight to design an MDS distance function.

Word2Vec

- Two common **word2vec** approaches:
 1. Try to predict word from surrounding words (**continuous bag of words**).
 2. Try to predict surrounding words from word (**skip-gram**).

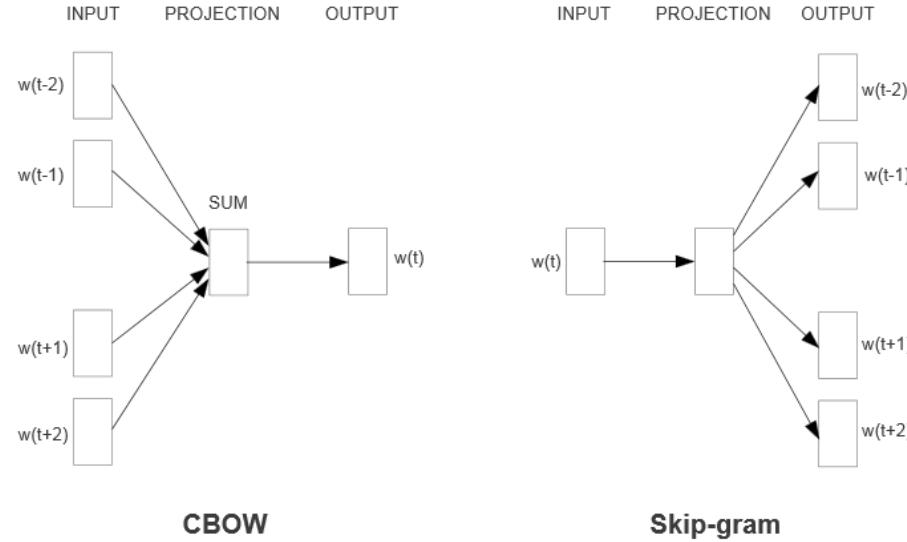


Figure 1: New model architectures. The CBOW architecture predicts the current word based on the context, and the Skip-gram predicts surrounding words given the current word.

- Train latent-factors to solve one of these supervised learning tasks.

Word2Vec

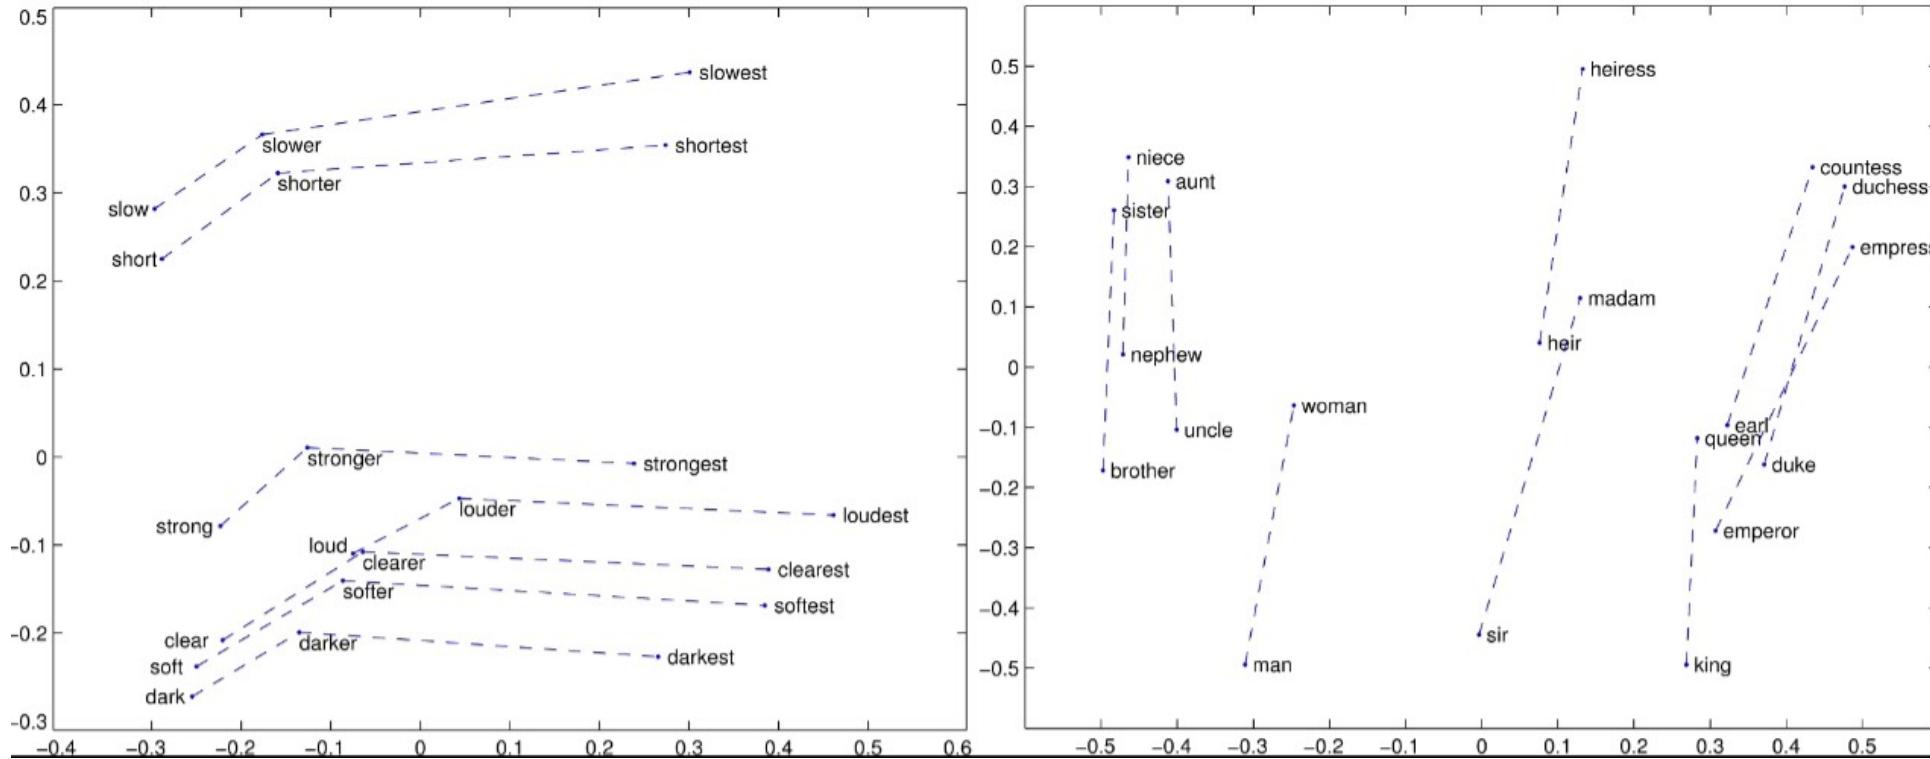
- In both cases, each word ‘i’ is represented by a vector z_i .
- In continuous bag of words (CBOW), we optimize the following likelihood:

$$\begin{aligned} p(x_i | x_{\text{surround}}) &= \prod_{j \in \text{surround}} p(x_i | x_j) && (\text{independence assumption}) \\ &= \prod_{j \in \text{surround}} \frac{\exp(z_i^T z_j)}{\sum_{c=1}^k \exp(z_c^T z_j)} && (\text{softmax over all words}) \end{aligned}$$

- Apply gradient descent to logarithm:
 - Encourages $z_i^T z_j$ to be big for words in same context (making z_i close to z_j).
 - Encourages $z_i^T z_j$ to be small for words not appearing in same context (makes z_i and z_j far).
- For CBOW, denominator sums over all words.
- For skip-gram it will be over all possible surrounding words.
 - Common trick to speed things up: sample terms in denominator (“negative sampling”).

Word2Vec Example

- MDS visualization of a set of related words:



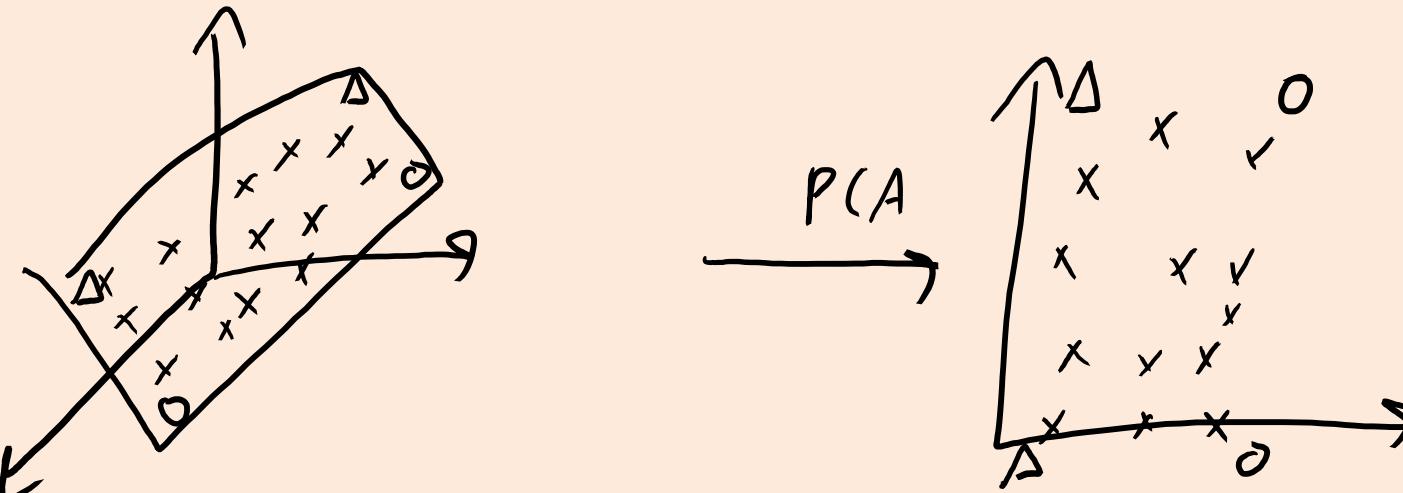
- Distances between vectors might represent semantics.

Summary

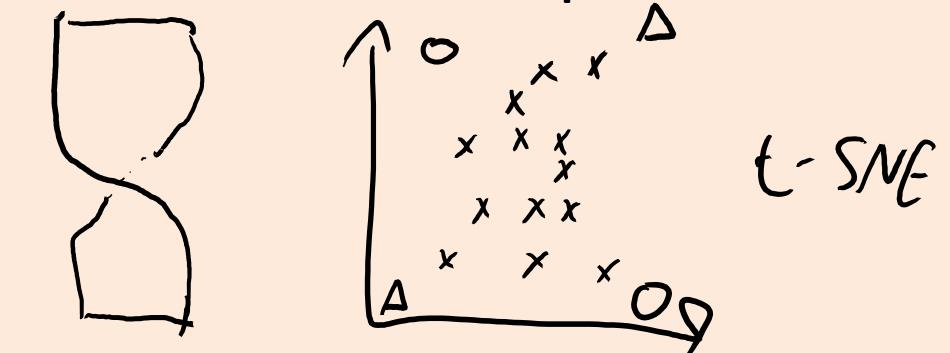
- Different MDS distances/losses/weights usually gives better results.
- Manifold learning focuses on low-dimensional curved structures.
- ISOMAP is most common approach:
 - Approximates geodesic distance by shortest path in weighted graph.
- t-SNE is promising new data MDS method.
- Word2vec:
 - Latent-factor (continuous) representation of words.
 - Based on predicting word from its context (or context from word).
- Next time: deep learning.

Does t-SNE always outperform PCA?

- Consider 3D data living on a 2D hyper-plane:

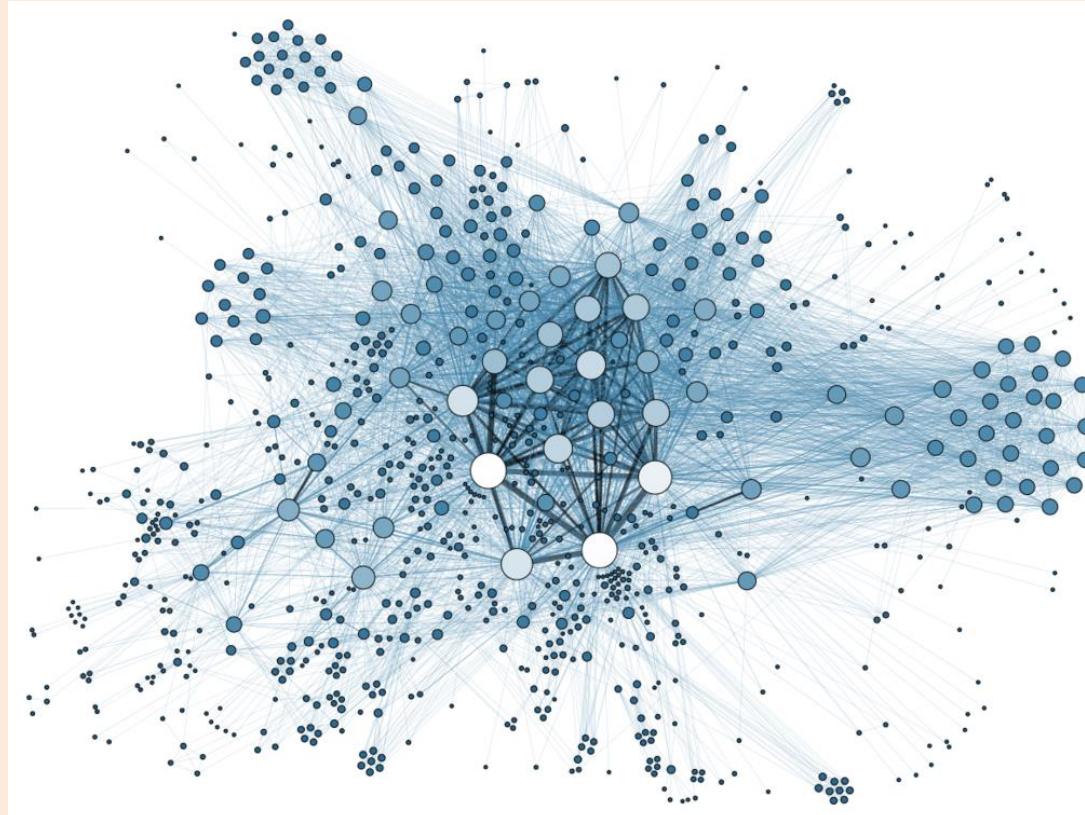


- PCA can perfectly capture the low-dimensional structure.
- T-SNE can capture the local structure, but can “twist” the plane.
 - It doesn't try to get long distances correct.



Graph Drawing

- A closely-related topic to MDS is **graph drawing**:
 - Given a graph, how should we display it?
 - Lots of interesting methods: https://en.wikipedia.org/wiki/Graph_drawing



Bonus Slide: Multivariate Chain Rule

- Recall the univariate chain rule:

$$\frac{d}{dw} [f(g(w))] = f'(g(w)) g'(w)$$

- The multivariate chain rule:

$$\nabla [f(g(w))] = \underbrace{f'(g(w))}_{\text{1x1}} \underbrace{\nabla g(w)}_{\text{dx1}}$$

- Example:

$$\nabla \left[\frac{1}{2} (w^T x_i - y_i)^2 \right]$$

$$= \nabla [f(g(w))]$$

$$\text{with } g(w) = w^T x_i - y_i$$

$$\text{and } f(r_i) = \frac{1}{2} r_i^2$$

$$\nabla g(w) = x_i$$

$$f'(r_i) = r_i$$

$$\nabla [f(g(w))] = r_i x_i$$

$$= (w^T x_i - y_i) x_i$$

Bonus Slide: Multivariate Chain Rule for MDS

- General MDS formulation:

$$\underset{Z \in \mathbb{R}^{n \times k}}{\operatorname{argmin}} \sum_{i=1}^n \sum_{j=i+1}^n g(d_1(x_i, x_j), d_2(z_i, z_j))$$

- Using multivariate chain rule we have:

$$\nabla_{z_i} g(d_1(x_i, x_j), d_2(z_i, z_j)) = g'(d_1(x_i, x_j), d_2(z_i, z_j)) \nabla_{z_i} d_2(z_i, z_j)$$

- Example: If $d_1(x_i, x_j) = \|x_i - x_j\|$ and $d_2(z_i, z_j) = \|z_i - z_j\|$ and $g(d_1, d_2) = \frac{1}{2}(d_1 - d_2)^2$

$$\nabla_{z_i} g(d_1(x_i, x_j), d_2(z_i, z_j)) = - (d_1(x_i, x_j) - d_2(z_i, z_j)) \left[- \frac{(z_i - z_j)}{2\|z_i - z_j\|} \right] \nabla_{z_i} d_2(z_i, z_j)$$

Assuming $z_i \neq z_j$

(move distances closer) (how distance changes in z -space)