

CPSC 340: Machine Learning and Data Mining

Convolutional Neural Networks

Fall 2019

Admin – Lectures this week



- Planned bus strike Wednesday-Friday.
- I'm planning to finish the “testable content” of the course today.
 - I might go a bit over time.
- Wednesday will be about “fun with deep learning”.
- Friday, I might cover different topics in the different sections:
 - Possible topics include semi-supervised learning, Google's PageRank, or proofs:
 - How many gradient descent iterations do we need?
 - What does the approximation error depend on?

Last Lectures: Deep Learning

- We've been discussing **neural network / deep learning** models:

$$\hat{y}_i = v^T h(W^{(m)} h(W^{(m-1)} h(\dots h(W^{(2)} h(W^{(1)} x_i) \dots)))$$

- We discussed **unprecedented vision/speech performance**.
- We discussed **methods to make SGD work better**:
 - **Parameter initialization** and **data transformations**.
 - Setting the **step size(s)** in stochastic gradient and using **momentum**.
 - Alternative non-linear functions like **ReLU**.



“Residual” Networks (ResNets)

- Impactful recent idea is residual networks ([ResNets](#)):




Figure 2. Residual learning: a building block.

- You can **take previous (non-transformed) layer as input** to current layer.
 - Also called “skip connections” or “highway networks”.
- **Non-linear part of the network only needs to model residuals.**
 - Non-linear parts are just “pushing up or down” a linear model in various places.
- This was a key idea behind first methods that used 100+ layers.
 - Evidence that biological networks have skip connections like this.

DenseNet

- More recent variation is “DenseNets”:
 - Each layer can see all the values from many previous layers.
 - Gets rid of vanishing gradients.
 - May get same performance with fewer parameters/layers.




Figure 1: A 5-layer dense block with a growth rate of $k = 4$. Each layer takes all preceding feature-maps as input.

Deep Learning and the Fundamental Trade-Off



- Neural networks are subject to the fundamental trade-off:
 - With increasing depth, training error of global optima decreases.
 - With increasing depth, training error may poorly approximate test error.
- We want deep networks to model highly non-linear data.
 - But increasing the depth can lead to overfitting.
- How could GoogLeNet use 22 layers?
 - Many forms of regularization and keeping model complexity under control.
 - Unlike linear models, typically use multiple types of regularization.

Standard Regularization

- Traditionally, we've added our usual **L2-regularizers**:

$$f(v, W^{(3)}, W^{(2)}, W^{(1)}) = \frac{1}{2} \sum_{i=1}^n (v^T h(W^{(3)} h(W^{(2)} h(W^{(1)} x_i))) - y_i)^2 + \frac{\lambda_4}{2} \|v\|^2 + \frac{\lambda_3}{2} \|W^{(3)}\|_F^2 + \frac{\lambda_2}{2} \|W^{(2)}\|_F^2 + \frac{\lambda_1}{2} \|W^{(1)}\|_F^2$$

- L2-regularization often called “**weight decay**” in this context.
 - Could also use L1-regularization: gives **sparse network**.



Standard Regularization


- Traditionally, we've added our usual **L2-regularizers**:

$$f(v, W^{(3)}, W^{(2)}, W^{(1)}) = \frac{1}{2} \sum_{i=1}^n (v^T h(W^{(3)} h(W^{(2)} h(W^{(1)} x_i))) - y_i)^2 + \frac{\lambda_4}{2} \|v\|^2 + \frac{\lambda_3}{2} \|W^{(3)}\|_F^2 + \frac{\lambda_2}{2} \|W^{(2)}\|_F^2 + \frac{\lambda_1}{2} \|W^{(1)}\|_F^2$$

- L2-regularization often called “**weight decay**” in this context.
 - Could also use L1-regularization: gives **sparse network**.
- “**Hyper-parameter**” optimization:
 - Try to optimize validation error in terms of $\lambda_1, \lambda_2, \lambda_3, \lambda_4$.
- Recent result:
 - Adding a regularizer in this way **creates bad local optima**.

Early Stopping

- Another common type of regularization is “early stopping”:
 - Monitor the validation error as we run stochastic gradient.
 - Stop the algorithm if validation error starts increasing.




Unfortunately it might look more like

↑ hopefully you don't stop here.

Dropout


- **Dropout** is a more recent form of explicit regularization:
 - On each iteration, **randomly set some x_i and z_i to zero** (often use 50%).



- Encourages **distributed representation** rather than relying on specific z_i .
 - Alternately, you are adding **invariance to missing inputs or latent factors**.
- After a lot of success, dropout may already be going out of fashion.

“Hidden” Regularization in Neural Networks


- Fitting single-layer neural network with SGD and no regularization:



- Training goes to 0 with enough units: we’re finding a global min.
- What should happen to training and test error for larger #hidden?

“Hidden” Regularization in Neural Networks

- Fitting single-layer neural network with SGD and no regularization:




- Test error continues to go down!?! Where is fundamental trade-off??
- There exist global mins with large #hidden units have test error = 1.
 - But among the global minima, SGD is somehow converging to “good” ones.

Implicit Regularization of SGD

- There is growing evidence that **using SGD regularizes parameters**.
 - We call this the “**implicit regularization**” of the optimization algorithm.
- Beyond empirical evidence, we know this happens in simpler cases.
- Example of implicit regularization:
 - Consider a **least squares** problem where there **exists a ‘w’ where $Xw=y$** .
 - Residuals are all zero, we fit the data exactly.
 - You run [stochastic] gradient descent starting from $w=0$.
 - Converges to **solution $Xw=y$ that has the minimum L2-norm**.
 - So **using SGD is equivalent to L2-regularization** here, but regularization is “implicit”.

Implicit Regularization of SGD

- Example of implicit regularization:
 - Consider a **logistic regression** problem where **data is linearly separable**.
 - We can fit the data exactly.
 - You run gradient descent from any starting point.
 - Converges to **max-margin solution** of the problem.
 - So **using gradient descent is equivalent to encouraging large margin**.



- Similar result known for **boosting**.

(pause)

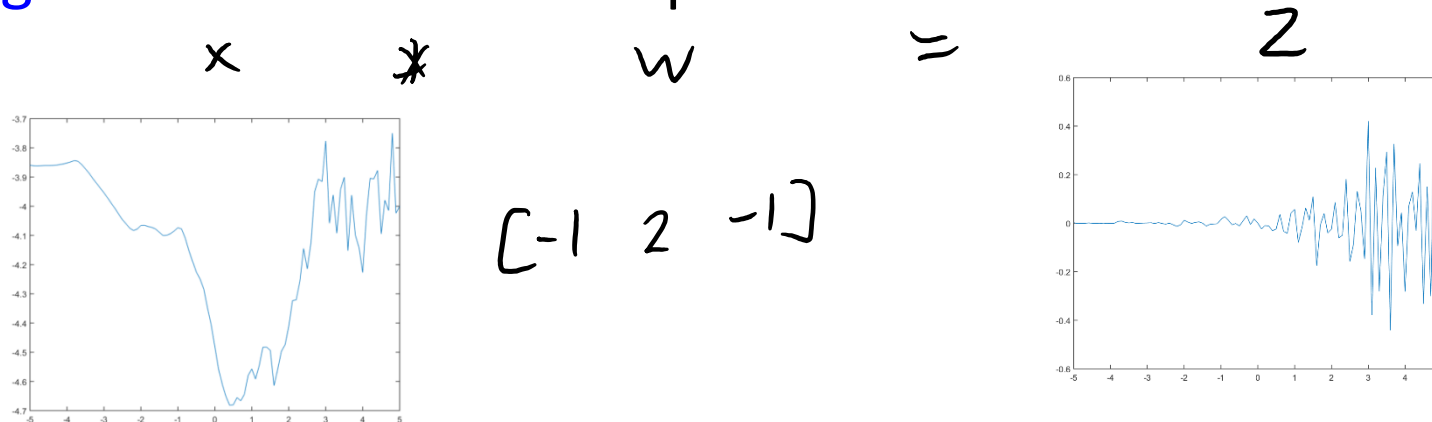
Deep Learning “Tricks of the Trade”

- We’ve discussed **heuristics to make deep learning work**:
 - Parameter initialization and data transformations.
 - Setting the **step size(s)** in stochastic gradient and using **momentum**.
 - **ResNets** and alternative non-linear functions like **ReLU**.
 - Different forms of regularization:
 - L2-regularization, early stopping, dropout, implicit regularization from SGD.
- These are often **still not enough** to get deep models working.
- Deep computer vision models are all **convolutional neural networks**:
 - The $W^{(m)}$ are **very sparse and have repeated parameters** (“tied weights”).
 - Drastically reduces number of parameters (speeds training, reduces overfitting).

1D Convolution as Matrix Multiplication

- 1D convolution:

- Takes signal 'x' and filter 'w' to produces vector 'z':



- Can be written as a matrix multiplication:

$$W_x = \begin{bmatrix} 1 & -2 & 1 & 0 & 0 & 0 & \dots & 0 & 0 & 0 & 0 \\ 0 & 1 & -2 & 1 & 0 & 0 & \dots & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & -2 & 1 & 0 & \dots & 0 & 0 & 0 & 0 \\ \vdots & & & \vdots & & & & & & & \\ 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0 & 1 & -2 & 1 \end{bmatrix} x = z$$

1D Convolution as Matrix Multiplication

- Each element of a convolution is an **inner product**:

$$\begin{aligned}
 z_i &= \sum_{j=-m}^m w_j x_{i+j} \\
 &= w^T x_{(i-m:i+m)} \\
 &= \tilde{w}^T x \quad \text{where } \tilde{w} = [0 \ 0 \ 0 \ \underbrace{\quad w \quad}_{\text{positions } i-m \text{ through } i+m} \ 0 \ 0]
 \end{aligned}$$

- So **convolution is a matrix multiplication** (I'm ignoring boundaries):

$$z = \tilde{W} x \quad \text{where } \tilde{W} = \begin{bmatrix} \overbrace{\quad w \quad} & \overbrace{\quad 0 \quad} & \overbrace{\quad 0 \quad} & \overbrace{\quad 0 \quad} \\ 0 & \overbrace{\quad w \quad} & \overbrace{\quad 0 \quad} & \overbrace{\quad 0 \quad} \\ 0 & 0 & \overbrace{\quad w \quad} & \overbrace{\quad 0 \quad} \\ 0 & 0 & 0 & \overbrace{\quad w \quad} \end{bmatrix}$$



matrix can be very sparse and only has $2m+1$ variables.

- The shorter 'w' is, the more sparse the matrix is.

2D Convolution as Matrix Multiplication

- 2D convolution:

- Signal 'x', filter 'w', and output 'z' are now all **images/matrices**:



- Vectorized 'z' can be written as a **matrix multiplication** with vectorized 'x':

$$W = \begin{bmatrix} -2 & -1 & 0 & 0 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & 0 & 0 & \dots & 0 & 0 & 1 & 2 & 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & -2 & -1 & 0 & 0 & 0 & \dots & 0 & 0 & -1 & 0 & 1 & 0 & 0 & \dots & 0 & 0 & 0 & 1 & 2 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0 & 0 & 0 & 0 & -1 & 0 & 1 & \dots & 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0 & 2 & -1 & 0 & 0 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0 & 0 & 2 & -1 & 0 & 0 & 0 & \dots & 0 & 0 & -1 & 0 & 1 & 0 & 0 & \dots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \end{bmatrix}$$

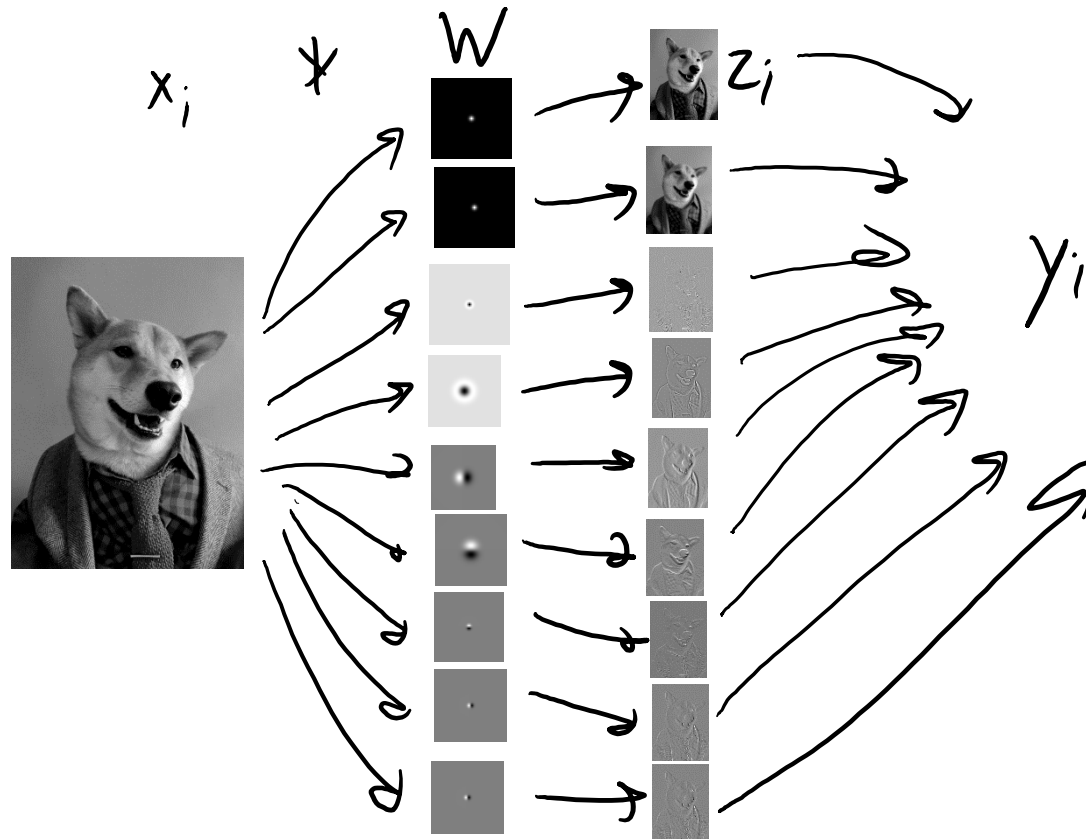
Motivation for Convolutional Neural Networks

- Consider training neural networks on 256 by 256 images.
 - This is 256 by 256 by 3 \approx 200,000 inputs.
- If first layer has $k=10,000$, then it has **about 2 billion parameters**.
 - We want to avoid this huge number (due to storage and overfitting).
- Key idea: make Wx_i act like several convolutions (to make it sparse):
 1. Each row of W only applies to part of x_i .
 2. Use the same parameters between rows.
- Forces most weights to be zero, reduces number of parameters.

$$w_1 = [0 \quad 0 \quad 0 \quad \text{---} \quad w \quad \text{---} \quad 0 \quad 0 \quad 0]$$
$$w_2 = [0 \quad \text{---} \quad w \quad \text{---} \quad 0 \quad 0 \quad 0 \quad 0 \quad 0]$$


Motivation for Convolutional Neural Networks

- Classic vision methods uses **fixed convolutions** as features:
 - Usually have **different types/variances/orientations**.
 - Can do subsampling or take **maxes across locations/orientations/scales**.





Motivation for Convolutional Neural Networks

- Convolutional neural networks learn the convolutions:
 - Learning 'W' and 'v' automatically chooses types/variances/orientations.
 - Don't pick from fixed convolutions, but learn the elements of the filters.



Motivation for Convolutional Neural Networks

- Convolutional neural networks learn the convolutions:
 - Learning 'W' and 'v' automatically chooses types/variances/orientations.
 - Can do multiple layers of convolution to get deep hierarchical features.



Convolutional Neural Networks


- Convolutional Neural Networks classically have 3 layer “types”:
 - Fully connected layer: usual neural network layer with unrestricted W.

$$W^{(m)} = \begin{bmatrix} \text{---} w_1^{(m)} \text{---} \\ \text{---} w_2^{(m)} \text{---} \\ \vdots \\ \text{---} w_k^{(m)} \text{---} \end{bmatrix}$$

Convolutional Neural Networks


- **Convolutional Neural Networks** classically have 3 layer “types”:
 - **Fully connected layer**: usual neural network layer with unrestricted W .
 - **Convolutional layer**: restrict W to act like several convolutions.

1D example

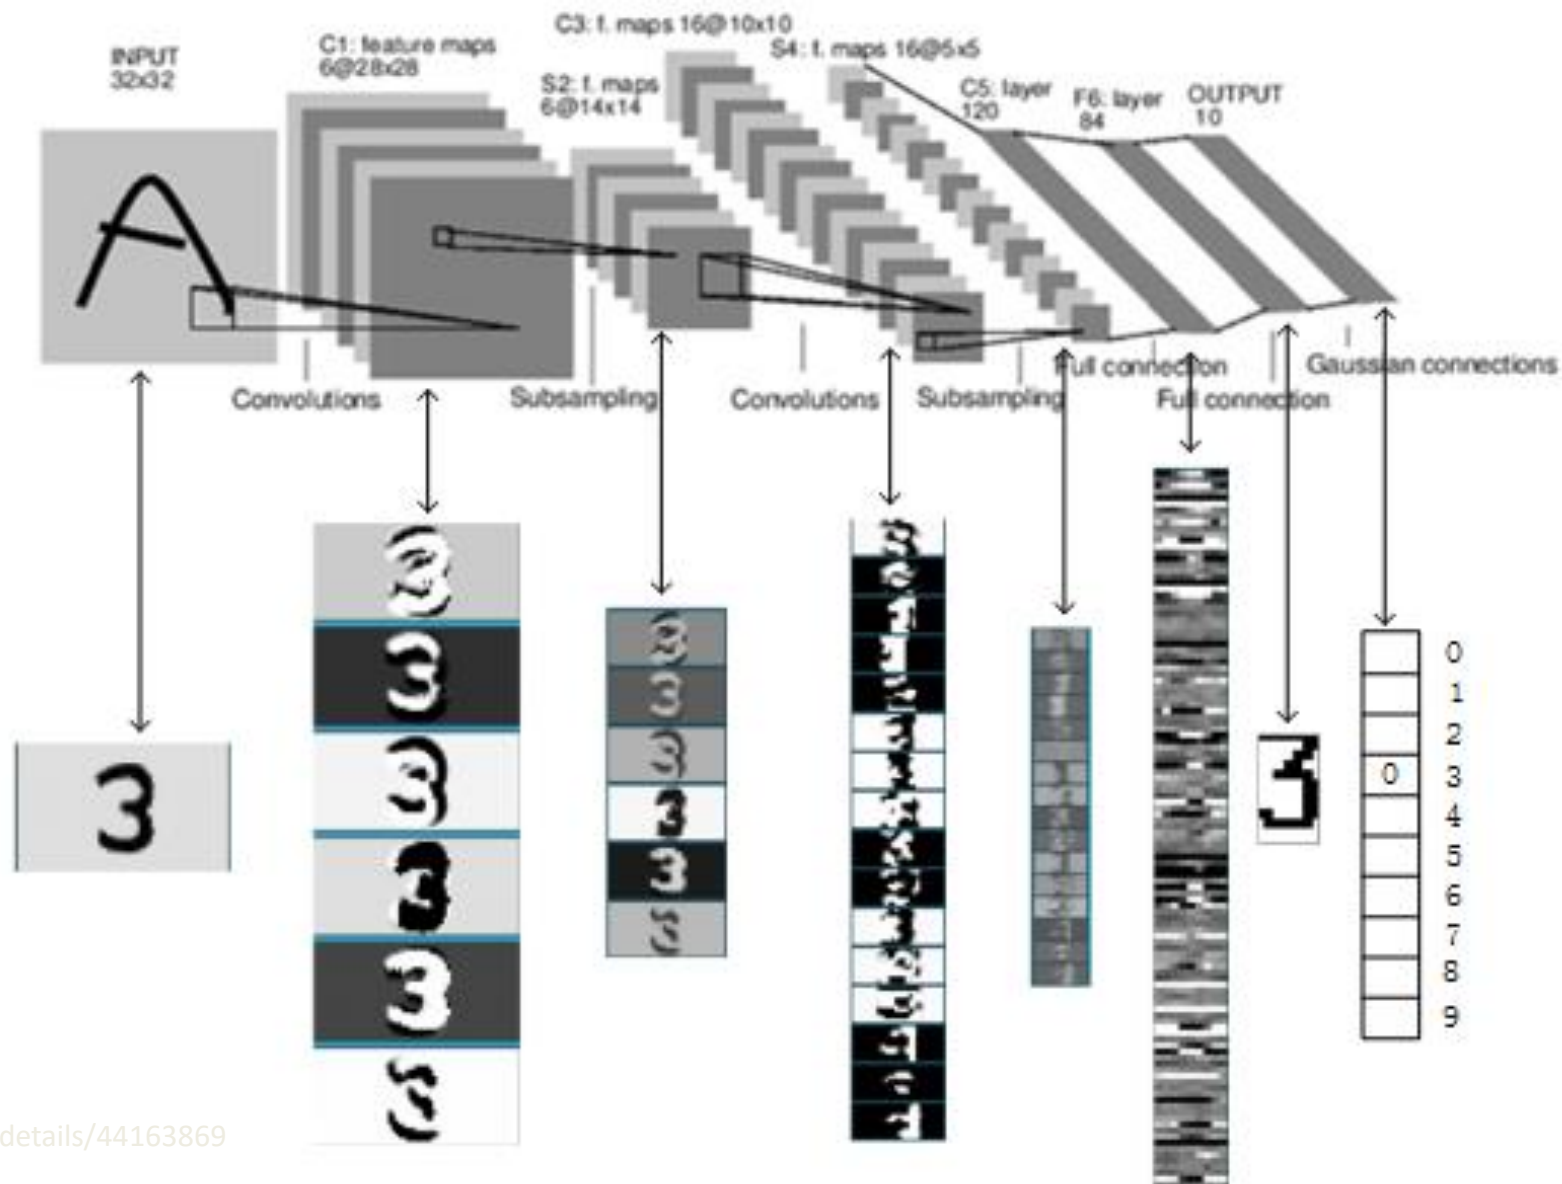


Convolutional Neural Networks




- Convolutional Neural Networks classically have 3 layer “types”:
 - Fully connected layer: usual neural network layer with unrestricted W .
 - Convolutional layer: restrict W to act like several convolutions.
 - Pooling layer: combine results of convolutions.
 - Can add some invariance or just make the number of parameters smaller.
 - Usual choice is ‘max pooling’:





LeNet for Optical Character Recognition



Deep Hierarchies in the Visual System




Deep Hierarchies in Optics




Last Time: Convolutional Neural Networks

- Classic **convolutional neural network** (LeNet):



- Visualizing the “activations” of the layers:

- <http://scs.ryerson.ca/~aharley/vis/conv>
- <http://cs231n.stanford.edu>



→ softmax
} 2 "fully-connected"
} max pooling
} 3D convolutions
} max pooling
} 2D convolutions

Partial Summary

- **ResNets** include untransformed previous layers.
 - Network focuses non-linearity on residual, allows huge number of layers.
- **Regularization** is crucial to neural net performance:
 - L2-regularization, early stopping, dropout, implicit regularization of SGD.
- **Convolutional neural networks:**
 - Restrict $W^{(m)}$ matrices to represent sets of convolutions.
 - Often combined with max (pooling).
- Next time: modern convolutional neural networks and applications.
 - Image segmentation, depth estimation, image colorization, artistic style.

(End of testable content for final exam)

Topics from Previous Years

- Slides for other topics that were covered in previous years:
 - [Ranking](#): finding “highest ranked” training examples (Google PageRank).
 - [Semi-supervised](#): using unlabeled data to help supervised learning.
 - [Sequence mining](#): approximate matching of patterns in large sequences.
- In previous years we did a [course review](#) on the last day:
 - Overview of topics covered in 340, and topics coming in 540.
 - [Slides here](#): this could help with studying for the final.

CPSC 330 vs. 340

- CPSC 330 starts next semester: “Applied Machine Learning”.
 - Not intended as a sequel to 340 (or even a prequel).
- There is some overlap in content, but focus is different:
 - More emphasis on the other steps of the data processing pipeline:
 - Data cleaning, feature extraction, reproducible workflows, communicating results.
 - More emphasis of “how to use packages”, less on “how stuff works”.
- If you found 340 too hard to keep up with, 330 might make sense.
 - In this situation, tell your friends about 330.

CPSC 330 vs. 540

- Next semester I'm teaching **CPSC 540**.
 - Intended as a **direct sequel to 340**.
 - We're basically starting with CNNs and going from there.
- Main focuses:
 - What if y_i is a sentence or an image or a protein?
 - Giving you the **background to understand the latest advances**.
- **Prerequisites**:
 - I expect you to know everything in this course and CPSC 320.
- Longer term, I expect this course to also be listed as **CPSC 440**.
 - 540 next semester is a “trial run” for CPSC 440.
 - I removed topics related to optimization research from the course.

Other ML-Related Courses

- CPSC 406:
 - Numerical optimization algorithms (like gradient descent).
- CPSC 422:
 - Includes topics like time series and reinforcement learning.
- CPSC 532R/533R:
 - Deep learning for vision, sound, and language.
- CPSC 533V:
 - Deep learning for computer graphics.
- EECE 592:
 - Deep learning and reinforcement learning.
- STAT 406:
 - Similar/complementary topics.
- STAT 460/461:
 - Advanced statistical issues (what happens when 'n' goes to ∞ ?)

Final Slide

- “Calling Bullshit in the Age of Big Data”:
 - <https://www.youtube.com/playlist?list=PLPnZfvKID1Sje5jWxt-4CSZD7bUI4gSPS>
 - Every “data scientist” should watch all these lectures.
 - You should be able to recognize non-sense, and not accidentally produce non-sense!
- Thank you for your patience.
 - This was the first time someone has taught multiple sections of this class.
- Good luck with finals/projects and the next steps!